



**KESHAV MEMORIAL INSTITUTE OF TECHNOLOGY
(AN AUTONOMOUS INSTITUTION)**



Accredited by NBA & NAAC, Approved by AICTE, Affiliated to JNTUH,
Narayanguda, Hyderabad, Telangana – 500029



DEPARTMENT OF INFORMATION TECHNOLOGY

LAB RECORD

SOFTWARE ENGINEERING LAB

B. Tech. III YEAR I SEM (KR23)

ACADEMIC YEAR 2025-26



KESHAV MEMORIAL INSTITUTE OF TECHNOLOGY
(AN AUTONOMOUS INSTITUTE)



NBA
NATIONAL BOARD
FOR ACCREDITATION

Accredited by NBA & NAAC, Approved by AICTE, Affiliated to JNTUH, Hyderabad
Narayanguda, Hyderabad, Telangana – 500029

Certificate

This is to certify that following is a Bonafide Record of the workbook task done by

_____ bearing Roll No _____ of _____

Branch of _____ year B. Tech. Course in the _____

Subject during the Academic year _____ & _____ under our supervision.

Number of week tasks completed: _____

Signature of Staff Member Incharge

Signature of Head of the Dept.

Signature of Internal Examiner

Signature of External Examiner



**KESHAV MEMORIAL INSTITUTE OF TECHNOLOGY
(AN AUTONOMOUS INSTITUTE)**

**Accredited by NBA & NAAC, Approved by AICTE, Affiliated to JNTUH, Hyderabad
Narayanguda, Hyderabad, Telangana – 500029**



NBA

Daily Laboratory Assessment Sheet

Name of the Lab:
Branch & Section:

Student Name:
HT. No:

Sr. No.	Name of the Experiment	Date	Observation Marks (5M)	Record Marks (5M)	Viva Voice Marks(10M)	Total Marks (20M)	Signature of Faculty
	Total						

Faculty Incharge

INDEX

Sr. NO.	CONTENTS	PAGE NO.
1.	<p>Software Installation & SRS Document</p> <ul style="list-style-type: none"> a. Abstract b. Functional Requirements (FR) c. Non-Functional Requirements (NFR) d. User Identification e. Workflow of Each User f. Use Cases 	
2.	<p>Exploring git local and remote commands on the multi-folder project</p> <ul style="list-style-type: none"> a. Pushing multi-folder project into private repository (by student). b. Students must explore all listed git commands on the multi-folder project in local and remote repository. c. Students must explore all git commands on given scenario-based question 	
3.	<p>Collaborative coding using git</p> <ul style="list-style-type: none"> a. To work on collaborative coding by: b. Creating Organization. c. Coordinating with others through a shared repository d. To resolve conflicts when collaborating on same part of code. e. To create and apply patch. 	
4.	<p>Build and package Java and Web applications using Maven</p> <ul style="list-style-type: none"> a. Understand the structure and lifecycle of a Maven project. b. Build and package Java and Web applications using Maven. c. Add dependencies using pom.xml, compile and test using plugins. d. Resolve errors and conflicts arising from dependency mismatches. e. Work with parent and multi-module Maven projects. f. Generate executable JARs and deployable WARs using Maven. 	
5.	<p>Docker CLI commands</p> <ul style="list-style-type: none"> a. Learn how to pull, run, stop, start, remove, and inspect containers and images. b. Gain the ability to create, monitor, and troubleshoot running containers. c. Configure and manage networks for container communication. d. Create and manage persistent storage for containers. e. Learn how to list, remove, and manage images efficiently. 	
6.	<p>Docker</p> <ul style="list-style-type: none"> a. Learn how to define and run multiple interdependent services (e.g., web server, database) in a single configuration file. 	

	<ul style="list-style-type: none"> b. Gain skills in writing and interpreting docker-compose.yml files for service setup. c. Deploy the same setup across different machines without manual configuration. d. Configure container networking and persistent storage within Compose. e. Reduce setup time and enable faster iteration during application development. 	
7.	<p><u>Creating a Multi-Module Maven Project</u></p> <ul style="list-style-type: none"> a. Build and package Java and Web applications using Maven. b. Add dependencies using pom.xml, compile and test using plugins. c. Resolve errors and conflicts arising from dependency mismatches. d. Work with parent and multi-module Maven projects. e. Generate executable JARs and deployable WARs using Maven 	
8.	<p><u>Jenkins Automation</u></p> <ul style="list-style-type: none"> a. Hands-on practice on manual creation of Jenkins pipeline using Maven projects from Github b. Create the job and build the pipeline for maven-java and maven-web project. 	
9.	<p><u>Pipeline Creation using script</u></p> <ul style="list-style-type: none"> a. Evaluation of Jenkins pipeline. b. WORKING ON BUILD TRIGGERS FOR LAST JENKINS PIPILINE c. Hands-on practice on creation of scripted Jenkins pipeline. d. Take the screenshots for above task 	
10.	<p><u>Working with minikube and Nagios</u></p> <ul style="list-style-type: none"> a. Hands-on practice of creating, running and scaling pods in minikube. b. Running Nginx server on specified port number by explaining the Nginx monitoring tool c. Running Nagios server and Understanding the Monitoring tool using Docker. d. AWS-free Trier account Creation steps 	
11.	<p><u>Jenkins-CI/CD</u></p> <ul style="list-style-type: none"> a. CI-Continuous Integration using Webhooks. b. Sending E-mail Notification on Build Failure or success 	
12.	<p><u>Creation of virtual machine for Ubuntu OS and Deploying the web application</u></p> <ol style="list-style-type: none"> 1. Creation of virtual machine 2. Deploying the web application 3. Accessing it publicly 	

LAB LESSON PLAN FOR WEEK 1

Lab Topics:

- Installation of Eclipse and other software (Installation PDF is available and will be shared before lab.)
- Downloading Maven and Configuring
- Student has to prepare on III/I project topic the following:
 - ✓ Abstract
 - ✓ Functional requirement
 - ✓ Non- Functional requirement
 - ✓ Identify user of application
 - ✓ Modules in application
 - ✓ Use case diagram

• **Use case diagram**

Unified Modeling Language (UML) diagrams are a way to visualize systems and software. It helps the software engineers understand designs and proposed implementation of complex software systems. UML (Unified Modeling Language) diagrams are categorized into two main types: **Structural diagrams** and **Behavioral diagrams**. These diagrams serve to model different aspects of a system, from its architecture and components to its behaviour and interactions.

Software Requirements Specification (SRS):

Project Title: Gen AI in Education - Personalized Quiz and Content Generator Using LLM.

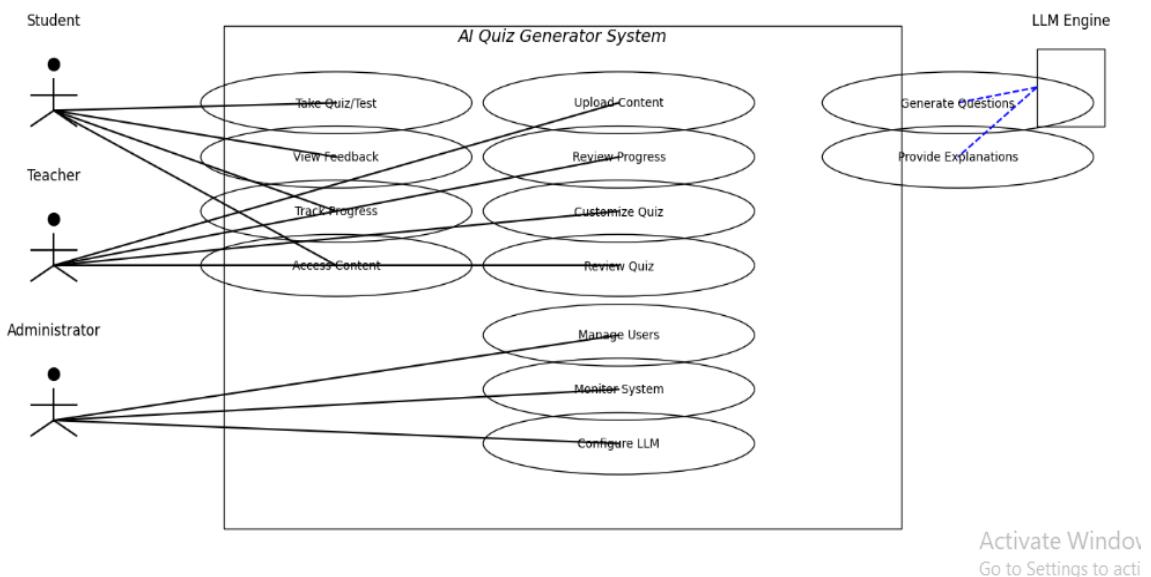
- **Abstract:** This project develops an AI-based system that uses Large Language Models (LLMs) to generate personalized quizzes and tests from educational content. It adapts questions to each student's learning level and performance, providing instant feedback and explanations. The system helps educators track progress and tailor instruction, enhancing learning outcomes through intelligent assessment.
- **Functional Requirements:**
 - a. **User Registration & Authentication:** Students, teachers, and admins can register and log in securely.
 - b. **Content Upload:** Teachers can upload educational materials (PDFs, notes, transcripts).
 - c. **Content Processing:** The system extracts key concepts from uploaded content using LLM.
 - d. **Personalized Quiz Generation:** LLM generates quizzes tailored to each student's learning profile and content.
 - e. **Quiz Attempt and Submission:** Students can take quizzes and submit answers through the platform.
 - f. **Feedback Generation:** The system provides instant feedback and explanations for each question.

- g. **Performance Tracking:** Student scores, accuracy, and progress are stored and visualized.
 - h. **Teacher Dashboard:** Teachers can view student analytics and quiz content.
 - i. **Admin Controls:** Admins can manage users, content repositories, and LLM configuration.
 - j. **LLM API Integration:** The system connects to external LLM services (e.g., OpenAI API) for generational tasks.
- **Non-Functional Requirements:**
 1. **Usability:** The interface should be user-friendly and intuitive for all users.
 2. **Performance:** Quiz generation time should be under 5 seconds.
 3. **Scalability:** The system should support 1,000+ concurrent users.
 4. **Security:** All data should be encrypted; authentication must be secure.
 5. **Availability:** The system should have at least 99.9% uptime.
 6. **Maintainability:** Codebase should be modular and well-documented for future updates.
 7. **Compatibility:** Must work on desktops, tablets, and smartphones.
 - **Use Cases:**
 1. **User Registration and Login:** Users (students, teachers, admins) create accounts and securely log in.
 2. **Content Upload:** Teachers upload educational materials such as PDFs, notes, and transcripts.
 3. **Content Processing:** The system processes uploaded content to extract key concepts using LLM.
 4. **Personalized Quiz Generation:** The system generates quizzes tailored to each student's learning progress and content.
 5. **Quiz Attempt:** Students take personalized quizzes within the platform.
 6. **Instant Feedback and Explanation:** After quiz submission, students receive AI-generated feedback and detailed explanations.
 7. **Performance Tracking:** The system records student quiz scores and tracks progress over time.
 8. **Teacher Analytics Dashboard:** Teachers view student performance reports

and quiz results.

9. **Quiz Review and Approval:** Teachers can review and edit AI-generated quizzes before assigning them.
10. **User Management:** Admins manage users, roles, and access permissions.
11. **System and LLM Configuration:** Admins configure LLM settings and integration parameters.

- **Workflow Diagram:**



Task : Students must explore all git commands on given scenario-based question. (Note student write command in observation book and paste the answer (git command) below each question for uploading).

Scenario based questions on basic git commands

1. You made changes to one file in the above project but haven't staged them yet. You realize they were a mistake. What Git command will you use to discard the local changes?

`git checkout -- <filename>`

2. You accidentally ran `git add file1.txt` (note instead of `file.txt` consider one file from above project), but you're not ready to commit yet. How do you remove it from the staging area without losing the changes?

`git reset file1.txt`

3. You made a commit but typed the wrong commit message. You haven't pushed it yet. How do you fix it?

`git commit --amend`

4. You want to view the commit history of the current branch in a readable format. What Git command should you use?

`git log --oneline --graph`

5. After cloning a repo, how can you set your name and email globally for all Git repositories?

`git config --global user.name "Name"`

`git config --global user.email "Mail Id"`

6. You've made some edits to your files but haven't staged them. How can you view the changes?

`git diff`

7. You're on the main branch but need to switch to feature/login. What command do you

`git checkout -b feature -branch`

8. You deleted the feature-ui branch by mistake. You haven't pushed the deletion. How

`git branch -d feature-branch`

9. You've made some commits locally and now want to upload them to the remote repository. What do you run?

`git push`

10. How can you fetch the latest changes from the remote repository without merging them automatically?

`git fetch`

11. You're on the main branch and want to start working on a new feature called search-filter. What command do you use?

`git checkout -b search-filter`

12. You committed a file containing an API key and want to completely remove it from the repo's history. What should you do?

`git filter-branch --force`

13. You want to see all the branches that exist both locally and on the remote. What
`git branch -a`

14. You're on main and want to merge the completed feature/signup branch into it. What command do you use?

`git merge feature/signup`

15. You tried to merge two branches and Git reported a conflict in app.js. What are the general steps to resolve it?

`git add app.js`

`git commit`

16. You don't want Git to track changes to .log files or node_modules/. How do you achieve this?

`echo "*.log" >>.gitignore`

`echo "node_modules/">.gitignore`

17. You're investigating a bug and want to know who last changed line 25 in script.py. What command do you use?

`git blame -L 25,25 script.py`

18. You're in the middle of working on a file but need to switch branches quickly. What do you do to save your work?

`git stash`

19. You previously ran git stash and now want to restore those changes. What command do you use?

`git stash pop`

20. The feature/test branch is no longer needed. How do you delete it locally?

`git branch -d feature/test`

21. You just merged the feature-ui branch into main. You want to clean up your local branches. How do you safely delete feature-ui?

`git reset file1.txt`

22. You created a branch feature-experiment, made some changes, but now realize the code is no longer needed. You want to delete it, even though it hasn't been merged. What command will you use?

`git branch -D featureexperiment`

23. You're currently on the feature-login branch and want to delete feature-ui. What must you ensure before running the delete command?

`git checkout main`

`git branch -d feature-ui`

24. You want to delete bugfix-footer, but you're unsure if it's been merged. What should you do before deleting it?

`git branch --merged | grep`

25. You finished working on feature-a, feature-b, and feature-c. All have been merged into main. How do you delete them in one command?

Scenario based questions on working with remote repository using Git commands

1. Specify the git command when you're starting work on a project and need to clone a remote repository to your local machine.

`git clone <repo url>`

2. Specify the git command if you want to see the remotes that are connected to your local repository.

`git remote -v`

3. Specify the git command if you want to add a new remote repository to your local repository.

`git remote add <name> <url>`

4. Specify the git command to remove a remote repository from your local configuration:

`git remote remove <name>`

5. Specify the git command if you want to rename an existing remote:

`git remote rename <old-name><new-name>`

6. Specify the git command to fetch updates from the remote repository but not merge them into your local branch.

`git fetch`

7. Specify the git command to pull the latest changes from the remote repository and merge them into your local branch.

`git pull`

8. Specify the git command to push your local commits to a remote repository.

`git push`

9. Specify the git command to when pushing for the first time and want to set the remote branch you are pushing to.

`git push -u origin main`

10. Specify the git command to change the URL of a remote (e.g., after changing the remote repository address).

`git remote set -url origin https://github.com/url`

11. Specify the git command if you want to list all branches on the remote repository.

`git branch -r`

12. Specify the git command if a branch was deleted on the remote but still shows up locally.

`git fetch -p`

13. Specify the git command to fetch a specific branch from a remote.

`git fetch <remotename><branchname>`

14. Specify the git command if you want to see detailed information about a remote repository.

```
git remote show <remotename>
```

15. Specify the git command if you want to rebase your local branch onto a remote branch (this can be useful to keep your history linear)

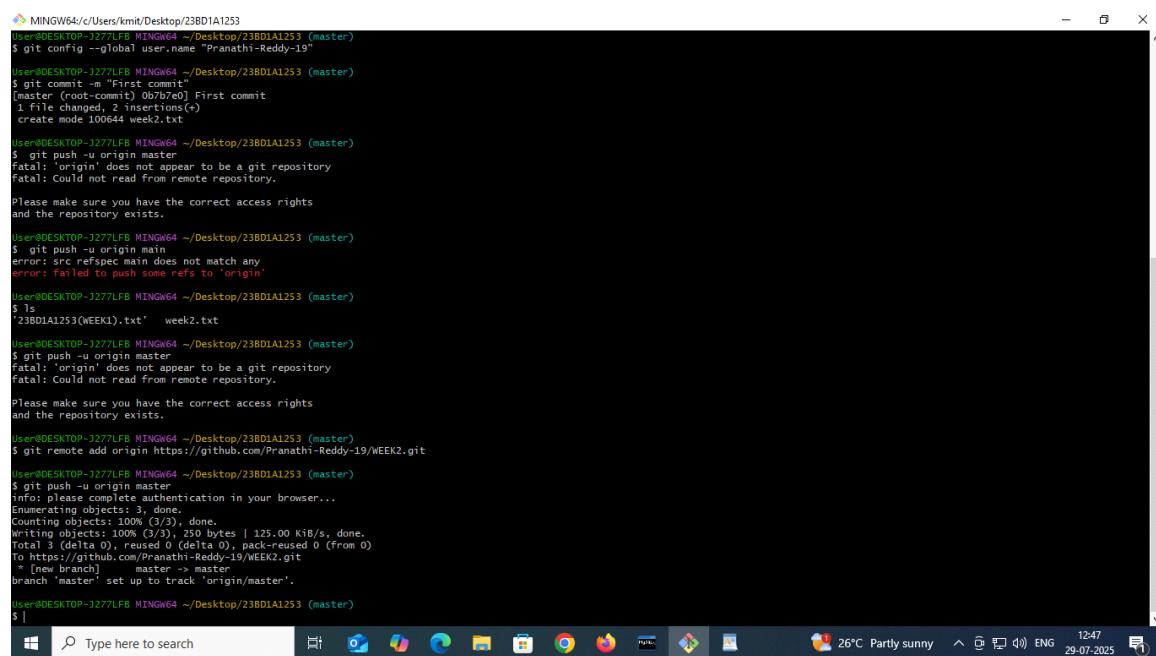
```
git pull --rebase origin <branchname>
```

Conclusion:

Successfully uploading a multi-folder project to both public and private GitHub repositories demonstrates student's practical understanding of **core Git commands** and **remote repository management**. By handling diverse file types (e.g., images, Java files, web pages, and pom.xml) and pushing them to different remote repositories, students show proficiency in:

- Initializing and configuring Git repositories
- Staging, committing, and managing version history
- Connecting to and working with multiple remote repositories
- Handling Gitignore and tracking large files appropriately
- Applying best practices for public vs. private repository usage

Completing these tasks equips students with essential version control skills and the ability to collaborate securely and professionally. It also reflects their readiness to contribute effectively to real-world software projects where source code management, team coordination, and secure code sharing are vital.



```
MINGW64/c/Users/kmit/Desktop/23BD1A1253
User@DESKTOP-J277LFB MINGW64 ~/Desktop/23BD1A1253 (master)
$ git config --global user.name "Pranathi-Reddy-19"
User@DESKTOP-J277LFB MINGW64 ~/Desktop/23BD1A1253 (master)
$ git commit -m "First Commit"
[master (root-commit) 0b87b7e] First commit
 1 file changed, 2 insertions(+)
 create mode 100644 week2.txt

User@DESKTOP-J277LFB MINGW64 ~/Desktop/23BD1A1253 (master)
$ git push -u origin main
fatal: 'origin' does not appear to be a git repository
fatal: Could not read from remote repository.

Please make sure you have the correct access rights
and the repository exists.

User@DESKTOP-J277LFB MINGW64 ~/Desktop/23BD1A1253 (master)
$ git push -u origin main
error: src refspec main does not match any
error: failed to push some refs to 'origin'

User@DESKTOP-J277LFB MINGW64 ~/Desktop/23BD1A1253 (master)
$ ls
'23BD1A1253(WEEK1).txt'  week2.txt

User@DESKTOP-J277LFB MINGW64 ~/Desktop/23BD1A1253 (master)
$ git push -u origin master
fatal: 'origin' does not appear to be a git repository
fatal: Could not read from remote repository.

Please make sure you have the correct access rights
and the repository exists.

User@DESKTOP-J277LFB MINGW64 ~/Desktop/23BD1A1253 (master)
$ git remote add origin https://github.com/Pranathi-Reddy-19/WEEK2.git
User@DESKTOP-J277LFB MINGW64 ~/Desktop/23BD1A1253 (master)
$ git push -u origin master
info: please complete authentication in your browser...
Enumerating objects: 3, done.
Counting objects: 100% (3/3), done.
Writing objects: 100% (3/3), 100 bytes | 125.00 KiB/s, done.
Total 3 (delta 0), reused 0 (delta 0), pack-reused 0 (from 0)
To https://github.com/Pranathi-Reddy-19/WEEK2.git
 * [new branch]    master --> master
branch 'master' set up to track 'origin/master'.

User@DESKTOP-J277LFB MINGW64 ~/Desktop/23BD1A1253 (master)
$ |
```

LAB ACTION PLAN FOR WEEK 3

Objective:

1. To work on collaborative coding by:
 - a. Creating Organization.
 - b. Coordinating with others through a shared repository
2. To resolve conflicts when collaborating on same part of code.
3. To create and apply patch
4. Students need to make note in observation book and also upload the screenshots of above executed exercise and scenario-based question.

1. Collaborative coding.

Code collaboration refers to the process of multiple people working together on the same codebase.

There are **two ways** of collaborative coding:

- a. Creating Organization
- b. Coordinating with others through a shared repository

a. Facilitating Collaborative Work by creating Organization

Organization - Organizations are shared accounts where businesses and open-source projects can collaborate across many projects at once. Owners and administrators can manage member access to the organization's data and projects with sophisticated security and administrative features. Steps to facilitate collaborative work by creating organization are:

Step 1: Click on New organization in Git hub

Click on Create a free organization.

Step 2: Set up your organization by entering the details like name, associated email, account verification, etc.

Click on Next

Step 3: Complete the setup by

- ✓ adding members to the group.
- ✓ Now, the invitation goes to other collaborator and they must accept.
- ✓ Next click on complete setup
- ✓ Goto Github and click on

Settings → Member privileges → Base permissions → Select write → change base permission to “Write”

Next under,

Projects base permissions → Select Write → change base permission to “Write”

Step 4: Create the remote repository for storing the project related files in Organization. This repository is accessible to every member of the team as per the permissions given.

Note: The repository can be made private so that it is accessible only to the group members rather than being in a public domain.

Now all the members of the team can contribute to the development of the project and the different files with all the versions and modification notices will be available in the respective repositories and is accessible to all the members

b. To collaborate effectively on GitHub shared repository by sharing code and coordinating with others through a shared repository.

→**Steps for Collaborator 1 are:** (collaborator-1 has repository that he wants to share with collaborator-2)

- ✓ Go to Settings > collaborators
- ✓ Now in Manage Access click on Add people
- ✓ Now you will be able to see ONE collaborator added
- ✓ After this invitation goes to collaborator 2 and they need to accept it.

→**Steps for Collaborator 2 are:**

✓ **1. Set Up Git and GitHub**

In git bash

```
git config --global user.name "Your Name"  
git config --global user.email you@example.com
```

Note : run \$ git remote -v

```
If origin git@github.com:OrgYOG/commonREPO.git (fetch)  
    origin git@github.com:OrgYOG/commonREPO.git (push)  
then
```

```
$ git remote -v // Note now, not connected to any remote
```

✓ **2. Goto Git hub and copy url of shared repository by collaborator 1**

- Copy url (option 1)
- If you're contributing to an existing project: (option 2)
Visit the repository page and click Fork (if you don't have write access)

✓ **3. Clone the Repository**

Get a local copy of the repo:

In git bash

```
git clone https://github.com/username/repository-name.git  
cd repository-name
```

✓ **4. Create a Branch for Your Work**

Always work on a separate branch:

In git bash

```
git checkout -b feature/your-feature-name
```

✓ 5. Make Changes and Commit

Edit files, then stage and commit your changes:

In git bash to existing file / new file

```
git add .
```

```
git commit -m "Add a clear and descriptive commit message"
```

✓ 6. Push Your Branch to GitHub

→ In git bash

→ git push origin feature/your-feature-name

Note: Following steps needed if you forked from any public repository to contribute (option 2)

✓ 7. Open a Pull Request (PR)

- Go to the GitHub repo in your browser
- Click Compare & pull request
- Add a meaningful title and description
- Submit the pull request for review

✓ 8. Review and Discuss Changes

- Collaborators can comment, request changes, or approve the PR(Pull Request)
- Make edits if needed and push again; they will update the PR automatically

✓ 9. Merge the Pull Request (has to be done by both collaborators)

Once approved:

- The repo owner can click Merge pull request
- Or you can squash and merge, depending on the team's workflow

✓ 10. Keep Your Branch Updated (has to be done by owner collaborator 1)

Before making new changes:

In git bash

```
git checkout main
```

```
git pull origin main
```

```
git checkout feature/your-new-feature
```

```
git merge main
```

2. To resolve conflicts when collaborating on same part of code.

Resolving conflicts when collaborating with GitHub (or Git in general) is a normal part of team development. Conflicts happen when two people make changes to the same part of the code, and Git doesn't know which version to keep.

✓ Steps to Resolve a Conflict

1. See Which Files Are in Conflict (Say f1.txt)

In git bash

```
git status
```

Conflicted files will be marked like:

In git bash
both modified: (Say f1.txt)

3. Open the File and Look for Conflict Markers by clicking on

- ✓ Click on **Resolve Conflict**

Git will insert conflict markers in the file like this:

```
def greet():
<<<<<< HEAD
    print("Hello from First collaborator ")
=====
    print("Hello from second collaborator ")
>>>>> feature/second collaborator branch
```

This shows the two conflicting changes:

- HEAD is your version
- The section below ===== is the other branch's version or second collaborator branch version

3. Edit the File to Fix the Conflict

- ✓ Choose one version or combine them:

```
def greet():
    print("Hello from First collaborator ")
    print("Hello from second collaborator ")
```

- ✓ Then **delete all the conflict markers** (<<<<<<, =====, >>>>>).

4. Mark the Conflict as Resolved

Once you've edited and saved the file:

- ✓ git add f1.txt

5. Complete the Merge or Pull

- ✓ git commit # Only needed if you're merging

Or if you're rebasing:

- ✓ git rebase --continue

⚠ If You Want to Abort the Merge

If it gets messy and you want to cancel:

```
git merge --abort
```

Here below screen shot shows conflict raised when collaborator 1 wanted to merge changes to same line that already collaborator 2 added the line of code at that line in same file but in their own branch.

Example :

Merge Conflict – Changes made to a file by two different users leads to merge conflict as below.

Say I am collaborator 2, now

- ➔ Accept invitation from collaborator 1
- ➔ Next goes to Shared repository by collaborator 1 (say abcrepo)
- ➔ Now open Git bash and clone the shared repo as below steps

Step 1: first connect to this “abcrepo” to local repository by clone

- ✓ Now change directory to “abcrepo”
- ✓ Switch to branch feature-abc
- ✓ Now add line in README.md file
- ✓ Now git add, commit
- ✓ Now see status and push to remote branch feature-abc
- ✓ Now go to Github and refresh or click on “abcrepo”, you can see README.md

Step 2: Now collaborator 2 adding line in README.md, save, git add, commit, push. No conflicts

Step 3: Also collaborator 1 adding line in README.md at same line collaborator 2 added, then git add, commit, push gives **merge conflict**.

Now open file, Edit the File to Fix the Conflict by keeping either or both collaborator changes and remove conflict markers. Add file, commit, and push.

3. To create and apply patch

A patch in Git is a text file that represents changes between two sets of files, or commits. It's essentially the output of the git diff command, packaged in a format that can be applied to another set of files.

Steps:

1. Goto github and take public project url
2. Clone into git bash
3. Now open file and edit it, add, commit
4. Run git log
5. Create patch with commit hash code as:

```
git format-patch -1 commit-hash-code
```

This creates 0001-commit-from-cloned-person.patch file.

Note: git format-patch -1 <commit-hash>

- -1 means "create a patch from 1 commit"
- This creates a .patch file for a single commit.
- Example:

```
git format-patch -1 abc1234
```

- This creates a file like 0001-Your-commit-message.patch

To create patches for the **last 3 commits**: git format-patch -3

or

Create patches from multiple commits

git format-patch <base_commit>..HEAD

6. Get the path of above patch file and email or what's up to the remote owner of file.
7. Once the remote owner gets the patch file 0001-commit-from-cloned-person.patch next the owner need to apply patch file in his git bash using git command below:

Syntax:

git apply my-changes.patch

Or, if it's a patch made by git format-patch, use:

git am 0001-Your-commit-message.patch

✓ **git apply 0001-Your-commit-message.patch**

4. Scenario based questions on working with remote repository

1. You've cloned a repository and made some changes to a local branch. Now you want to push these changes to the remote repository, but you're getting an error saying "rejected - non-fast-forward." How would you resolve this?

- Answer:

Run:

git pull --rebase origin branch-name

Resolve any conflicts, then:

git push origin branch-name

2. You've been working on a feature branch, and now you need to push it to the remote repository. However, the remote repository already has a main branch. How do you push your feature branch without affecting the main branch?

- Answer:

Run:

git push origin feature-branch-name

3. You cloned a remote repository, but after a while, the repository's structure changed and new branches were added. How would you keep your local repository updated with the latest changes from the remote repository?

- Answer:

Run:

```
git fetch --all  
git checkout new-branch-name
```

4. A colleague has pushed some changes to the main branch, but you have local changes in the same branch. You want to pull their changes, but you want to avoid merge conflicts. What steps would you take?

- Answer:

Run:

```
git add .  
git commit -m "your changes"  
git pull --rebase origin main
```

5. You accidentally pushed a sensitive file (e.g., API keys) to the remote repository. How would you fix this situation?

- Answer:

Run:

```
git rm --cached sensitive_file  
Add to .gitignore, commit.  
Use git filter-repo or BFG to clean history.  
git push --force
```

6. You're working on a feature branch, and your manager requests that you integrate the latest changes from main into your feature branch. What steps would you take?

- Answer:

Run:

```
git checkout feature-branch  
git pull origin main
```

7. You cloned a remote repository, but later you find that you need to push your changes to a different remote repository. How do you configure your local repository to push to this new remote?

- Answer:

Run:

```
git remote set-url origin new-url  
git push -u origin branch-name
```

8. After running git pull, you notice that your local branch is behind the remote branch. How would you proceed to bring your local branch up to date without losing your local changes?

- Answer:

Stash or commit local changes first, then:

```
git pull --rebase origin branch-name
```

9. You're working on a project with multiple collaborators, and you notice that your local changes conflict with changes that have been pushed by others. How would you resolve the conflicts?

- Answer:

Run:

```
git pull --rebase origin branch-name
```

Resolve conflicts, then:

```
git add .
```

```
git rebase --continue
```

10. You've pushed a feature branch to a remote repository, but now you need to delete the branch from the remote. How would you do that?

- Answer:

Run:

```
git push origin --delete feature-branch-name.
```

Scenario:

You are working on a collaborative project hosted on GitHub with a team of four developers. The main branch is main, and each developer works on their own feature branches.

You are assigned to implement a new UI component in a branch called feature/ui-update. Meanwhile, another team member has made a critical bug fix and pushed it directly to the main branch. When you try to push your changes, your push fails due to upstream changes.

Additionally, a teammate has submitted a patch file with a small CSS fix and shared it via email. You need to apply the patch, test it, and merge everything into main without breaking any functionality.

11. What command will you use to bring your local main branch up to date with the remote repository before merging it into your feature branch?

- Answer:

Run:

```
git checkout main
```

```
git pull origin main
```

12. How will you update your feature/ui-update branch to reflect the latest changes from main?

- Answer:

Run:

```
git checkout feature/ui-update  
git merge main
```

13. Which command should you use to attempt pushing your local feature branch again, and what should you do if it fails due to conflicts?

- Answer:

Run:

```
git push origin feature/ui-update  
If conflicts occur, resolve them, commit, and push again.
```

14. How do you apply a .patch file provided by your teammate and include it in your commit history?

- Answer:

Run:

```
git am 0001-Your-commit-message.patch
```

15. After successful testing, describe the steps (with commands) to merge your feature branch into main and push it to GitHub.

- Answer:

Run:

```
git checkout main  
git pull origin main  
git merge feature/ui-update  
git push origin main
```

Conclusion:

The effective use of Git in student collaborative projects promotes **teamwork, organization, and clear communication**, allowing students to manage code efficiently and track contributions accurately. By following structured Git workflows and best practices—such as regular commits, feature branching, and resolving conflicts early—students gain valuable experience in real-world development processes. This not only enhances their technical skills but also prepares them for **professional teamwork in software engineering**, where version control and collaboration are essential. Ultimately, mastering Git empowers students to deliver cleaner, more reliable code and succeed in collaborative environments.

LAB ACTION PLAN FOR WEEK 4

Objective:

To enable students to:

- Understand the structure and lifecycle of a Maven project.
- Build and package Java and Web applications using Maven.
- Add dependencies using **pom.xml**, compile and test using plugins.
- Resolve errors and conflicts arising from dependency mismatches.
- Work with parent and multi-module Maven projects.
- Generate executable JARs and deployable WARs using Maven.

Students must **document observations**, include **screenshots of executions**, and answer **scenario-based questions** after completing the tasks.

• Understanding Maven Project Structure

Maven standardizes the project structure for both Java and web-based applications. `src/`

```
└── main/
    ├── java/      → Application source code
    └── resources/ → Configuration files like config.properties
└── test/
    ├── java/      → Unit test source code
    └── resources/ → Test resources
```

The compiled files and reports are generated in the `target/` directory after a successful build.

• Steps to Perform Maven Build and Testing

---mvn clean install

- **clean**: Deletes the previous build (`target/` folder)
- **install**: Builds, tests, and installs the package to local `.m2` repository

After execution:

- `target/` folder contains compiled `.class` files and the final `.jar` or `.war`
- Artifact is stored in:
`~/ .m2/repository/groupId/artifactId/version/`

Add a Dependency (e.g., Gson):

In `pom.xml`:

```
<dependency>
<groupId>com.google.code.gson</groupId>
<artifactId>gson</artifactId>
<version>2.10</version>
```

```
</dependency>
```

After running:

```
mvn clean install
```

Check:

- Gson JAR is downloaded to `.m2/repository/com/google/code/gson/gson/`
- If version is wrong or not found → **BUILD FAILURE** with dependency resolution error

- **Configure Java Version via Compiler Plugin**

In `pom.xml`:

```
<plugin>
<groupId>org.apache.maven.plugins</groupId>
<artifactId>maven-compiler-plugin</artifactId>
<version>3.11.0</version>
<configuration>
<source>17</source>
<target>17</target>
</configuration> </plugin>
```

Check:

- Run `mvn package`
- Inspect generated JAR: `jar tf target/myapp.jar`

- **JUnit Testing and Reports**

Place test files in `src/test/java`. Example:

```
public class AppTest {
    @Test
    public void testSum() {
        assertEquals(5, 2 + 3);
    }
}
```

Run:

```
mvn test
```

Check:

- `target/test-classes/` → compiled test `.class` files
- `target/surefire-reports/` → `.txt` or `.xml` test results
If test fails → corresponding log and failure trace shown

- **Handling Errors in pom.xml**

- Typos in version numbers or missing repositories cause **BUILD FAILURE**
- Maven shows precise error in console
- Fix the dependency tag → re-run `mvn clean install`

- **Adding Resource Files**

Put config.properties inside:

```
src/main/resources/config.properties
```

To read:

```
InputStream input =  
getClass().getClassLoader().getResourceAsStream("config.properties");
```

After build, check target/classes/ → file should exist there.

- **Multi-Module and Parent Projects**

Structure:

```
parent/  
|   └── pom.xml (packaging: pom)  
└── core/  
    |   └── pom.xml  
    └── web/  
        └── pom.xml
```

- Parent pom.xml defines <modules> and common dependencies
- Each submodule builds independently into its target/ directory

- **Executable JARs**

Add to pom.xml:

```
<build>  
  <plugins>  
    <plugin>  
      <groupId>org.apache.maven.plugins</groupId>  
      <artifactId>maven-jar-plugin</artifactId>  
      <configuration>  
        <archive>  
          <manifest>  
            <mainClass>com.example.Main</mainClass>  
          </manifest>  
        </archive>  
      </configuration>  
    </plugin>  
  </plugins>  
</build>
```

Run:

```
mvn package  
java -jar target/myapp.jar
```

- **Building a WAR File**

Create a Maven web project with structure:

```
src/main/webapp/  
└── WEB-INF/web.xml
```

Add:

```
<packaging>war</packaging>
```

Command:

```
mvn package
```

Generates target/mywebapp.war → deploy on Tomcat server.

- **Scenario-Based Questions:**

- If my error is about:

*Failed to execute goal org.apache.maven.plugins:maven-compiler-plugin:3.8.1:compile
[ERROR] -> [Help 1]*

[ERROR] To see the full stack trace of the errors, re-run Maven with the -X switch.

How can I resolve it using maven terminal?

- A dependency you added is not recognized by the compiler. What steps would you take to confirm it is available in .m2 and listed in dependency tree?
- A teammate sends a .patch file for a bug fix. How would you apply it and include it in your Maven build?
- You have multiple **JUnit** test failures and want to rerun only failed tests. How would you approach this?
- Your Maven build fails due to “Unsupported class version error.” What plugin and configuration would you review?
- You need to change your Java application from a **WAR** to a standalone **JAR**. What pom.xml changes are needed?
- You are required to change the default build output directory from **target/** to **build_output/**. How would you configure it?
- You want to skip tests during the Maven build. What command would you use?
- How would you generate a site report (with test coverage, dependency analysis) for a Maven project?

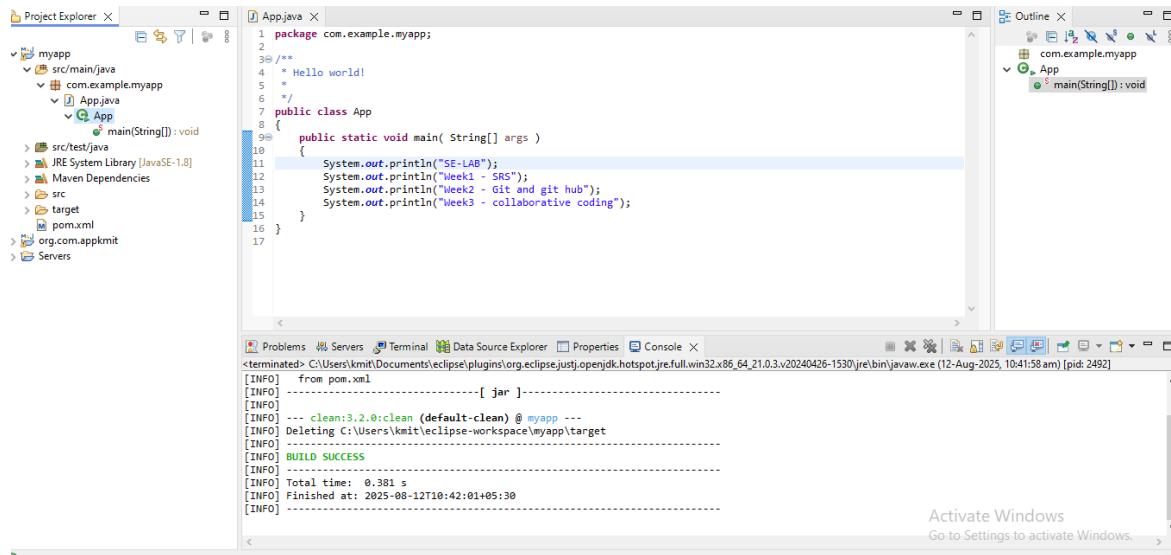
- How do you build a Java project using Maven, and what files are generated in the **target/** folder after running `mvn clean install`?
- How does Maven resolve dependency conflicts when two libraries use different versions of the same dependency, and how can you view and manage the **dependency tree**?
- How do you write and run a JUnit test in a Maven project, and where are the compiled test classes and reports stored after running `mvn test`?
- How can you create an executable **JAR** with a main method using Maven, and which **plugin** helps configure this behavior?
- How do you install and use a custom third-party JAR file in your Maven project, and how can you confirm it's included in the build and **classpath**?
- How do you create a Maven web project that packages into a WAR file, and what is the standard folder structure for such a project?
- What command do you use to build a WAR file in Maven, where is it generated, and how can you deploy it to a server like **Apache Tomcat**?
- How do you add JSTL and **servlet-api dependencies** in a Maven web project, and why should the servlet API use provided scope instead of compile?
- How do you set up a multi-module Maven web project with separate modules for core logic and web interface, and how are these modules built and connected?
- How do you configure a Maven web project, and how does its packaging and execution differ from a traditional WAR-based application?

Conclusion

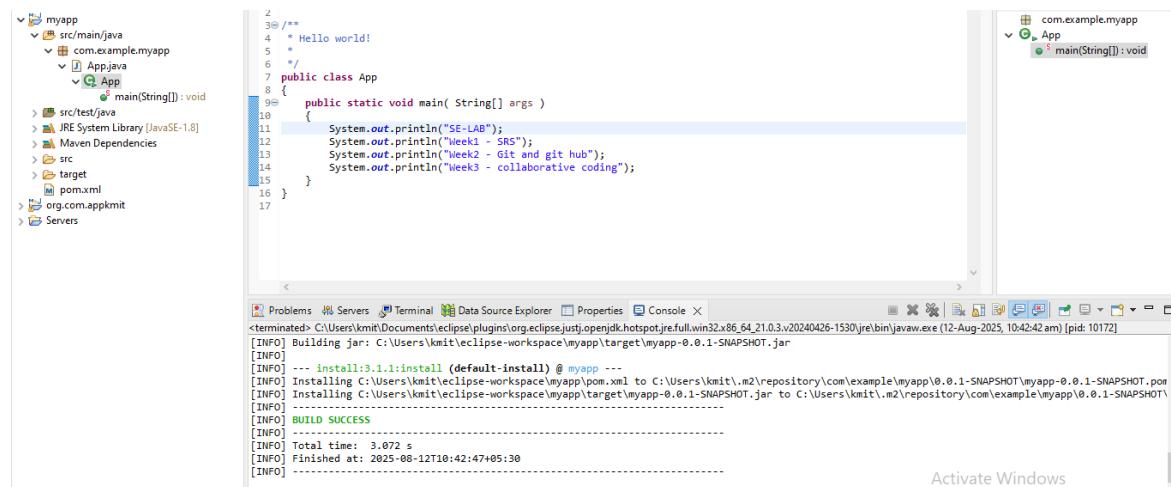
Mastering Maven empowers students to structure projects efficiently, manage complex dependencies, and automate testing and packaging. By practicing real-world scenarios—such as resolving build errors, handling resource files, applying patches, and deploying to servers—students gain valuable hands-on experience aligned with professional software development workflows. Maven not only streamlines builds but also enforces standardization and reusability across projects. Through Maven, students learn efficient project management, dependency control, multi-module integration, and reproducible builds. Understanding Maven's lifecycle, plugin system, and error handling prepares students for professional DevOps and CI/CD workflows. This lab equips students with hands-on experience in packaging, testing, resolving conflicts, and deploying real-world Java and Web applications.

Maven Project

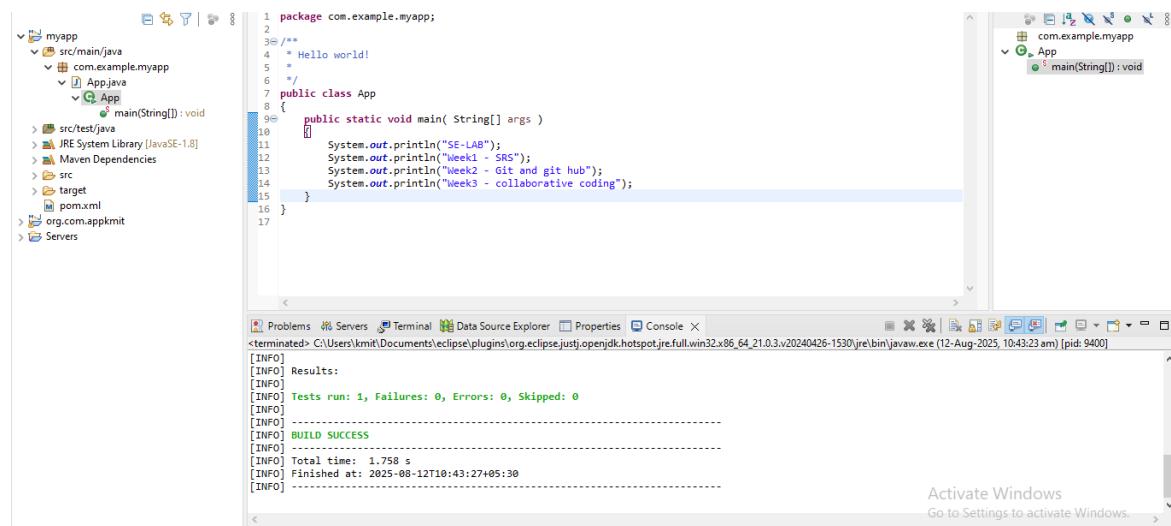
1. maven clean



2. maven install



3. maven test



4. maven build

The screenshot shows the Eclipse IDE interface. On the left is the Project Explorer with a single project named 'myapp'. Inside 'myapp' are 'src/main/java' and 'target'. Under 'src/main/java', there is a package 'com.example.myapp' containing an 'App.java' file. The code in 'App.java' is:

```

1 package com.example.myapp;
2 /**
3  * Hello world!
4  */
5 /**
6  */
7 public class App {
8     /**
9      * @param args
10     */
11    public static void main( String[] args )
12    {
13        System.out.println("SE-LAB");
14        System.out.println("Week1 - SRS");
15        System.out.println("Week2 - Git and git hub");
16        System.out.println("Week3 - collaborative coding");
17    }

```

On the right, the Java Editor shows the same code. Below the editor is the Console view, which displays the output of a Maven build:

```

<terminated> myapp () [Maven Build] C:\Users\kmt\Documents\plugins\org.eclipse.jdt\openjdk.hotspot.jre.full.win32.x86_64_21.0.3.v20240426-1530\jre\bin\javaw.exe (12-Aug-2025, 10:44:13 am)
[WARNING] target value 8 is obsolete and will be removed in a future release
[WARNING] To suppress warnings about obsolete options, use -Xlint:-options.
[INFO]
[INFO] --- surefire:3.2.5:test (default-test) @ myapp ---
[INFO] Skipping execution of surefire because it has already been run for this configuration
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 3.061 s
[INFO] finished at: 2025-08-12T10:44:18+05:30
[INFO] -----

```

5. run java application

The screenshot shows the Eclipse IDE interface. On the left is the Project Explorer with a single project named 'myapp'. Inside 'myapp' are 'src/main/java' and 'target'. Under 'src/main/java', there is a package 'com.example.myapp' containing an 'App.java' file. The code in 'App.java' is identical to the one in the previous screenshot.

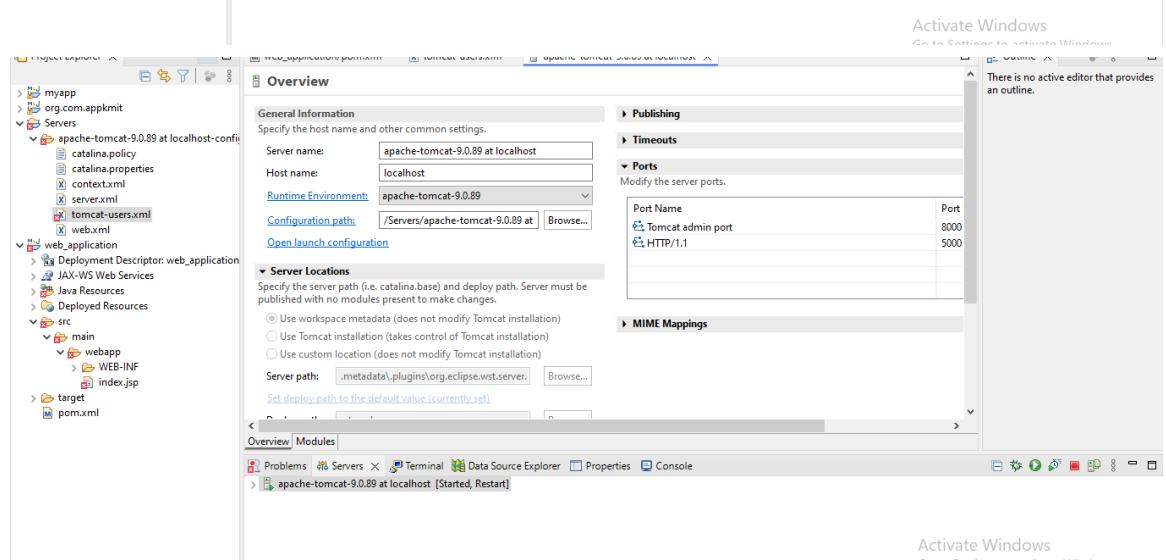
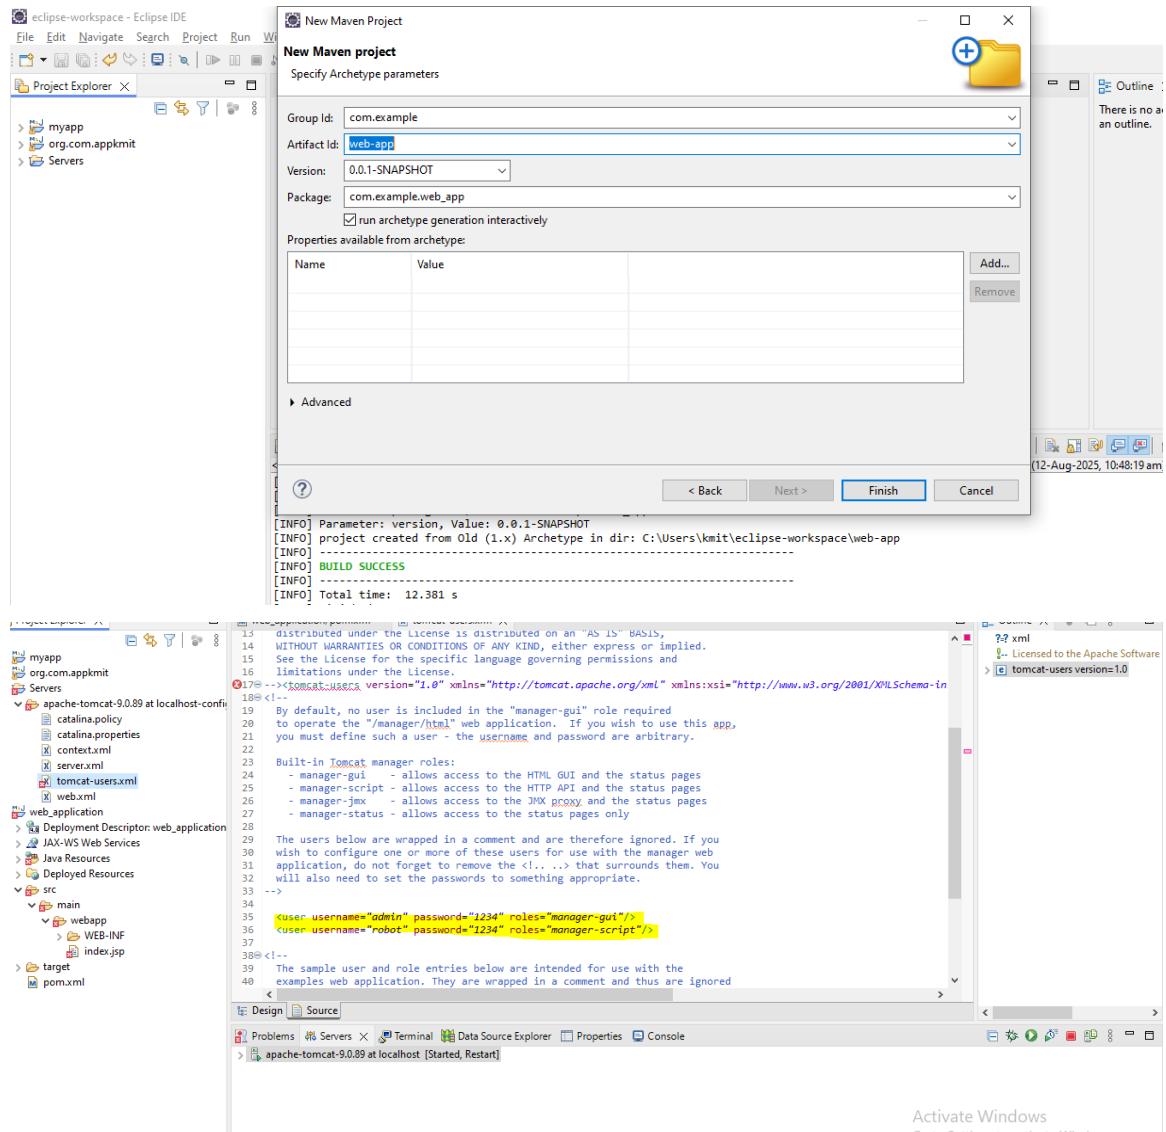
On the right, the Java Editor shows the same code. Below the editor is the Console view, which displays the output of running the application:

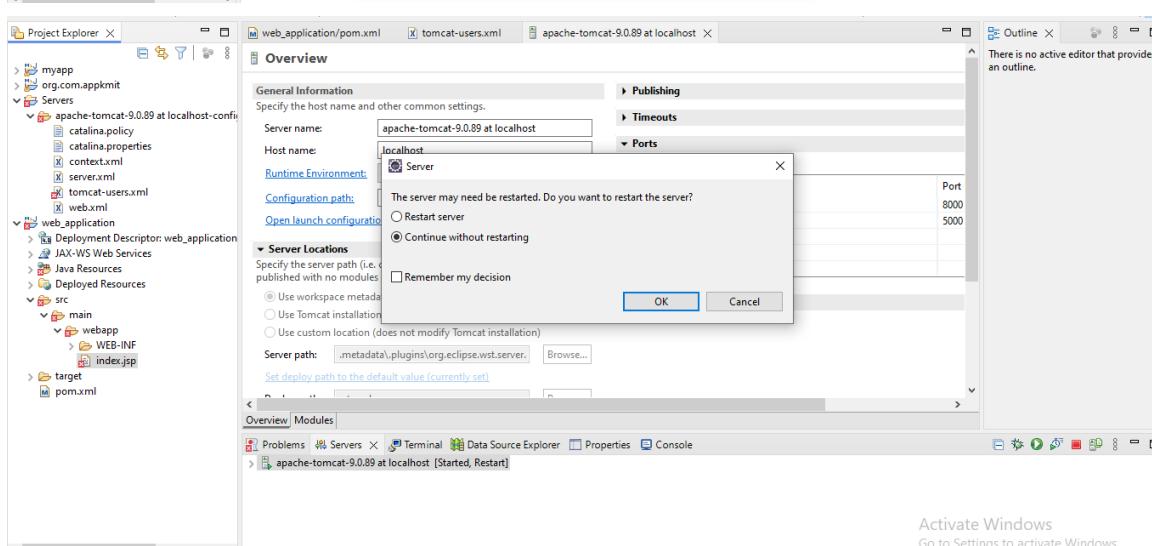
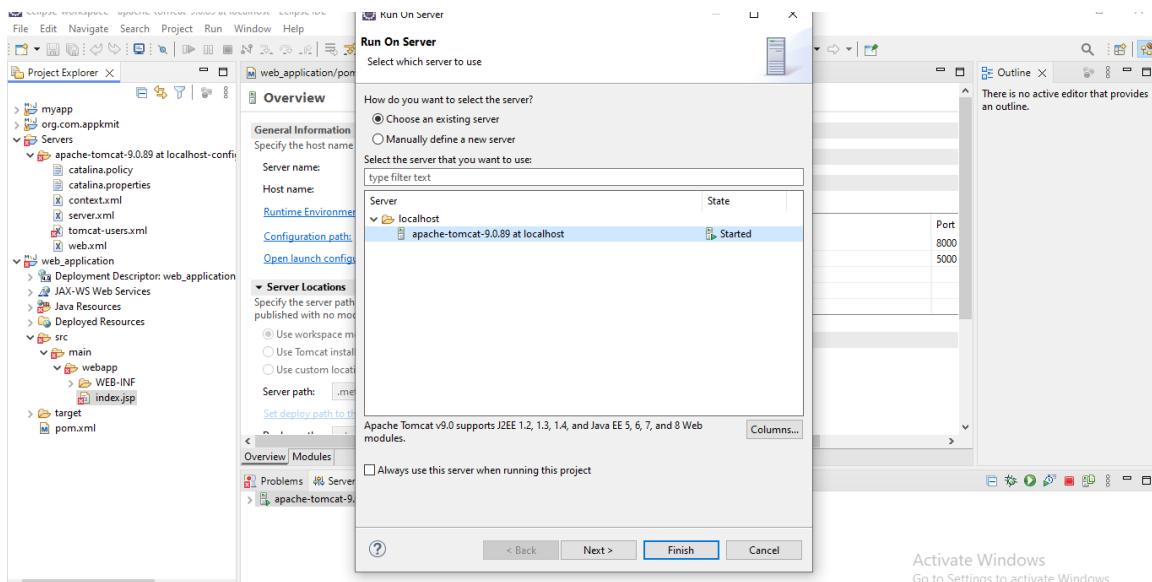
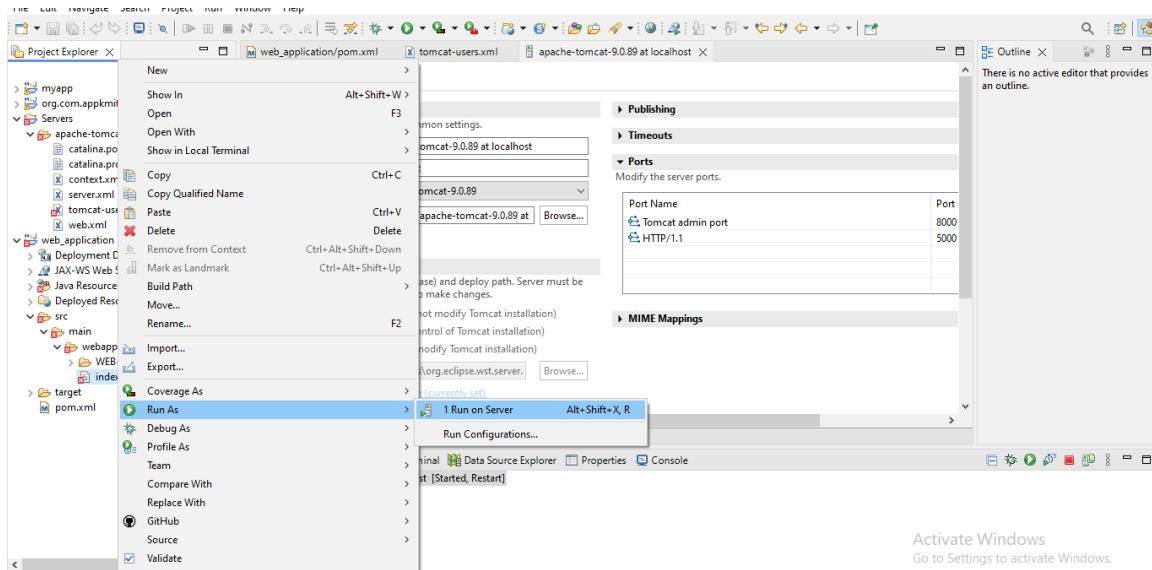
```

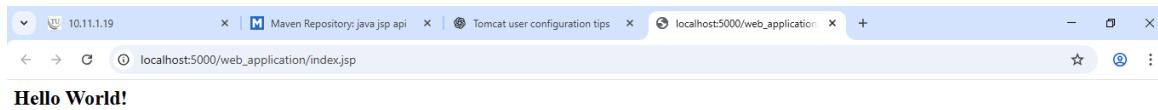
<terminated> App (1) [Java Application] C:\Users\kmt\Documents\plugins\org.eclipse.jdt\openjdk.hotspot.jre.full.win32.x86_64_21.0.3.v20240426-1530\jre\bin\javaw.exe (12-Aug-2025, 10:45:00 am)
SE-LAB
Week1 - SRS
Week2 - Git and git hub
Week3 - collaborative coding

```

WEB-APP







Hello World!

Activate Windows
Go to Settings to activate Windows.

LAB ACTION PLAN FOR WEEK 5

Objectives:

Students will able to:

- Learn how to pull, run, stop, start, remove, and inspect containers and images.
- Gain the ability to create, monitor, and troubleshoot running containers.
- Configure and manage networks for container communication.
- Create and manage persistent storage for containers.
- Learn how to list, remove, and manage images efficiently.

Docker:

Docker is an open-source platform that automates the process of building, packaging, and running applications inside lightweight, portable containers, ensuring they run consistently across different environments.

Docker Images:

A read-only template that contains the application code, runtime, libraries, and dependencies. It is like a blueprint for creating containers.

Docker Containers:

A running instance of a Docker image.

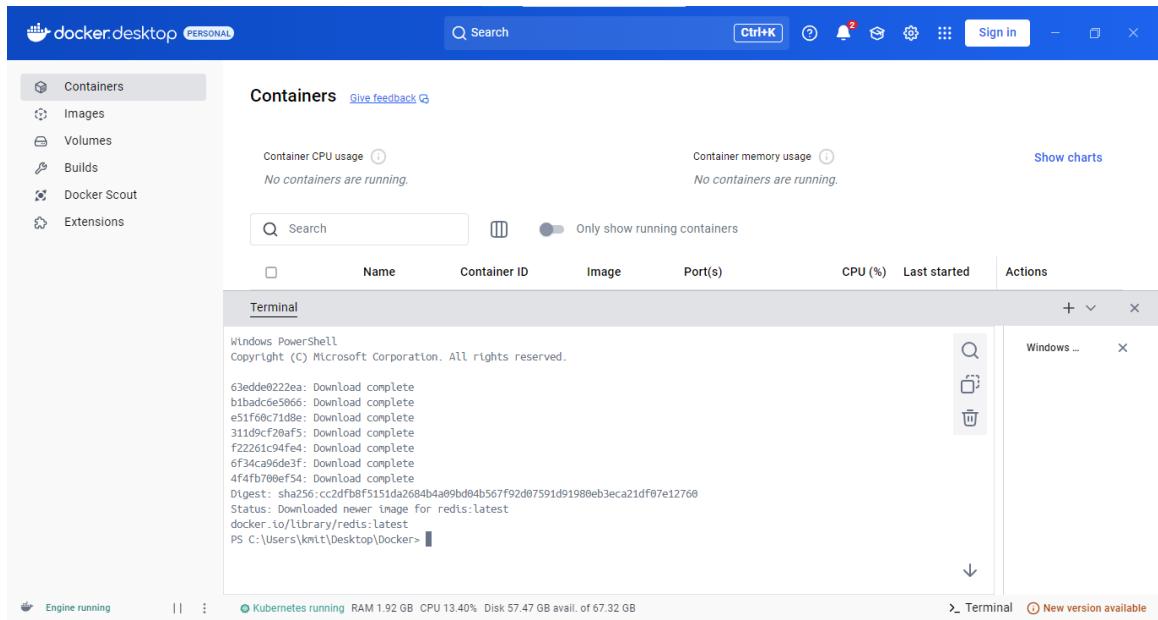
Working with Docker CLI Commands:

Docker CLI Commands with redis

Step 1: Pull the redis Image

Command:

```
docker pull redis
```



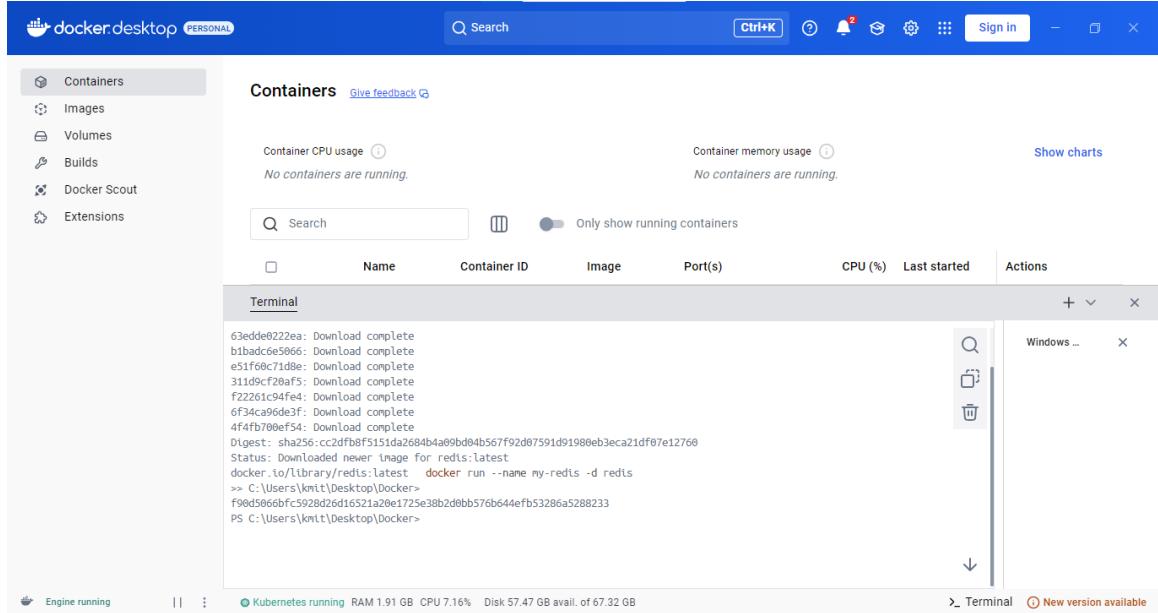
Step 2: Run a Redis Container

Command:

```
docker run --name my-redis -d redis
```

What It Does:

- Creates and starts a container named `my-redis` from the `redis` image.
- The `-d` flag runs the container in the background.



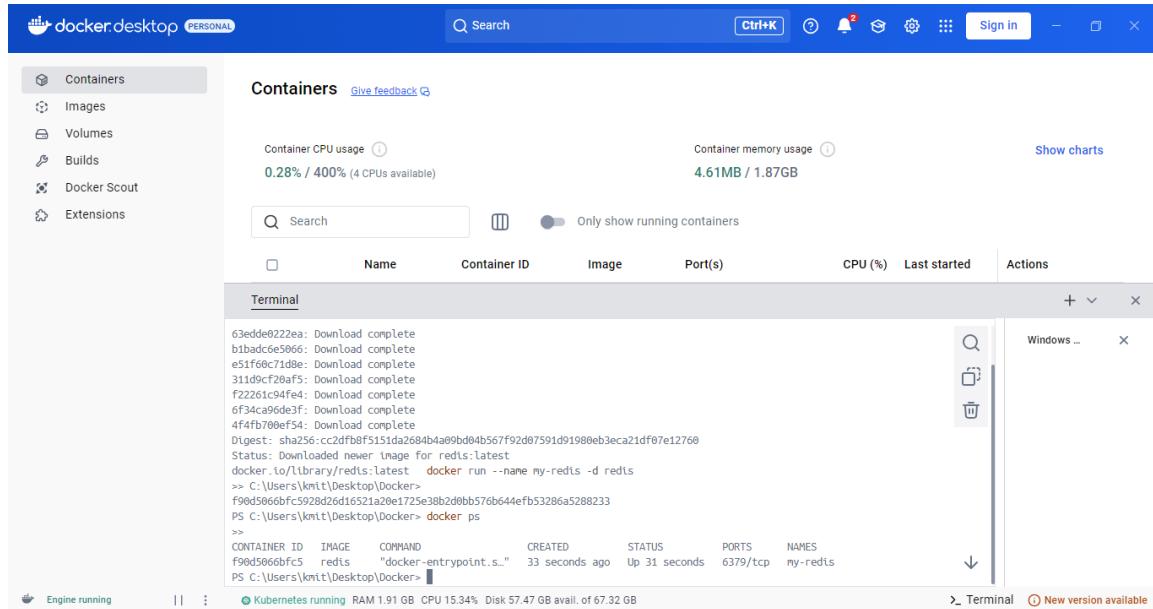
Step 3: Check Running Containers

Command:

`docker ps`

What It Does:

Lists all running containers.



Step 4: Access Redis

Command:

`docker exec -it my-redis redis-cli`

Opens the Redis command-line tool (redis-cli) inside the container.

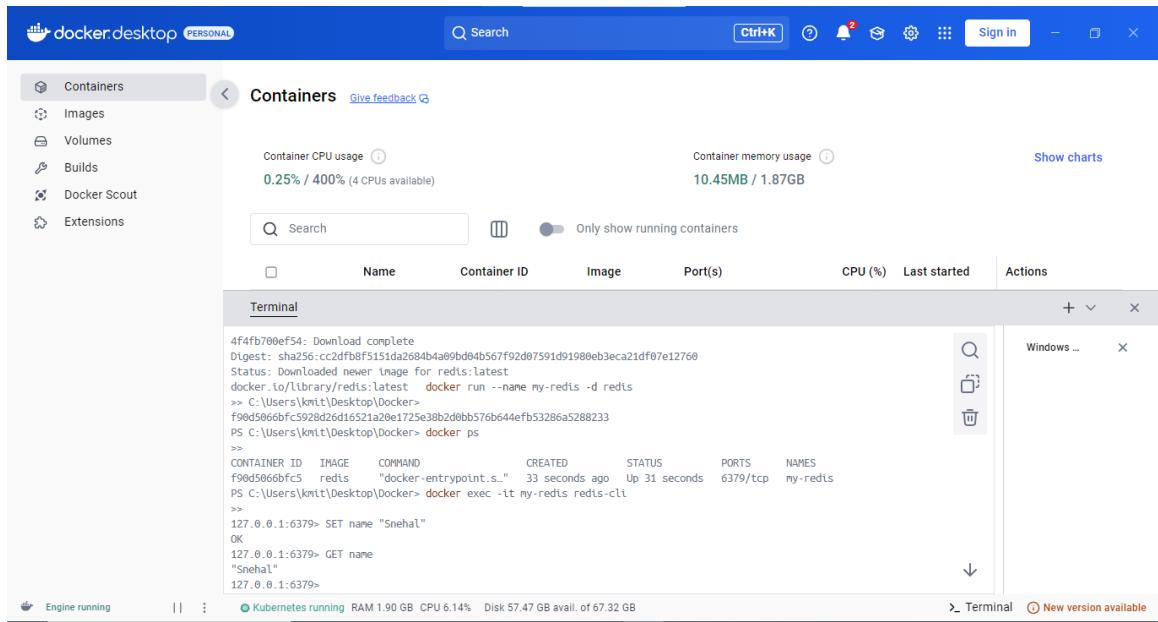
Example Redis Commands:

`127.0.0.1:6379> SET name "Alice"`

`OK`

`127.0.0.1:6379> GET name`

`"Alice"`



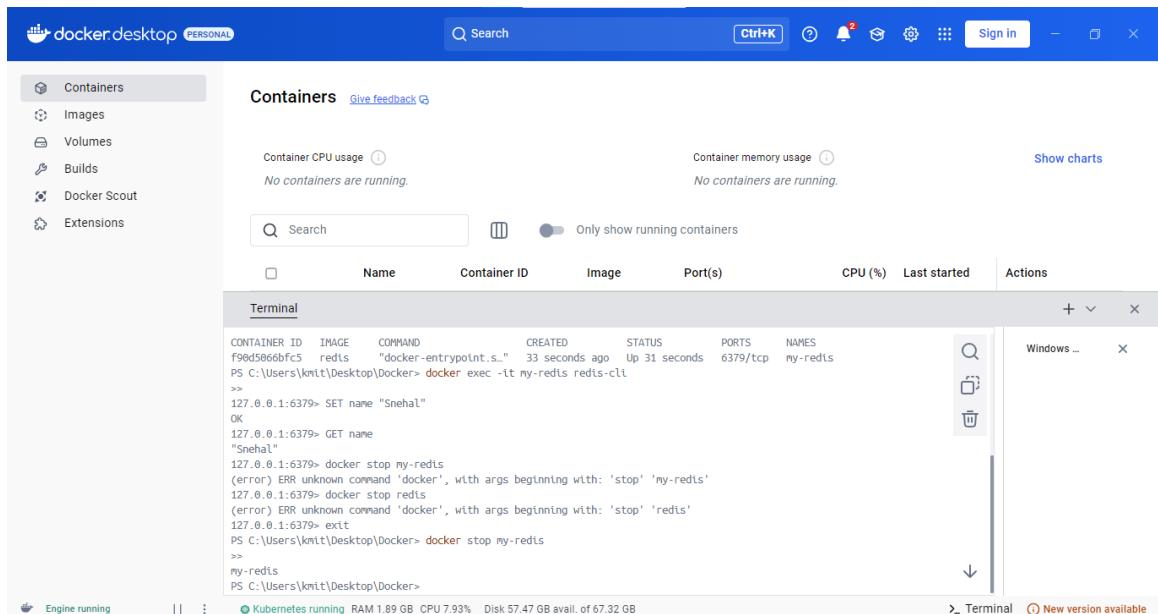
Step 5: Stop the Redis Container

Command:

```
docker stop my-redis
```

What It Does:

stops the Redis container but doesn't delete it.



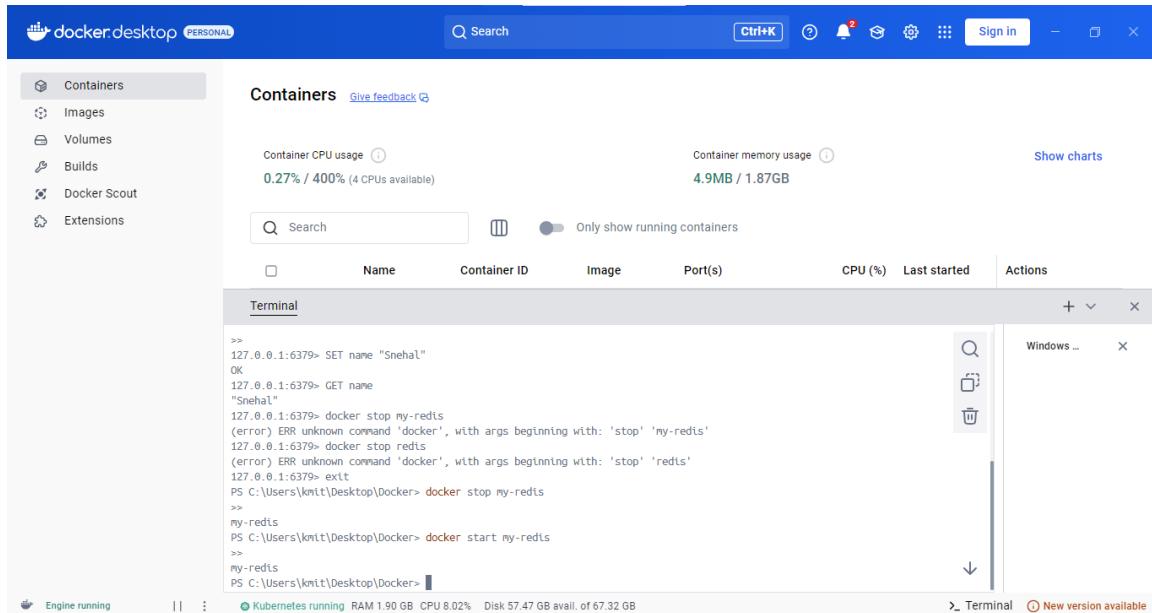
Step 6: Restart the Redis Container

Command:

`docker start my-redis`

What It Does:

Restarts the stopped container.



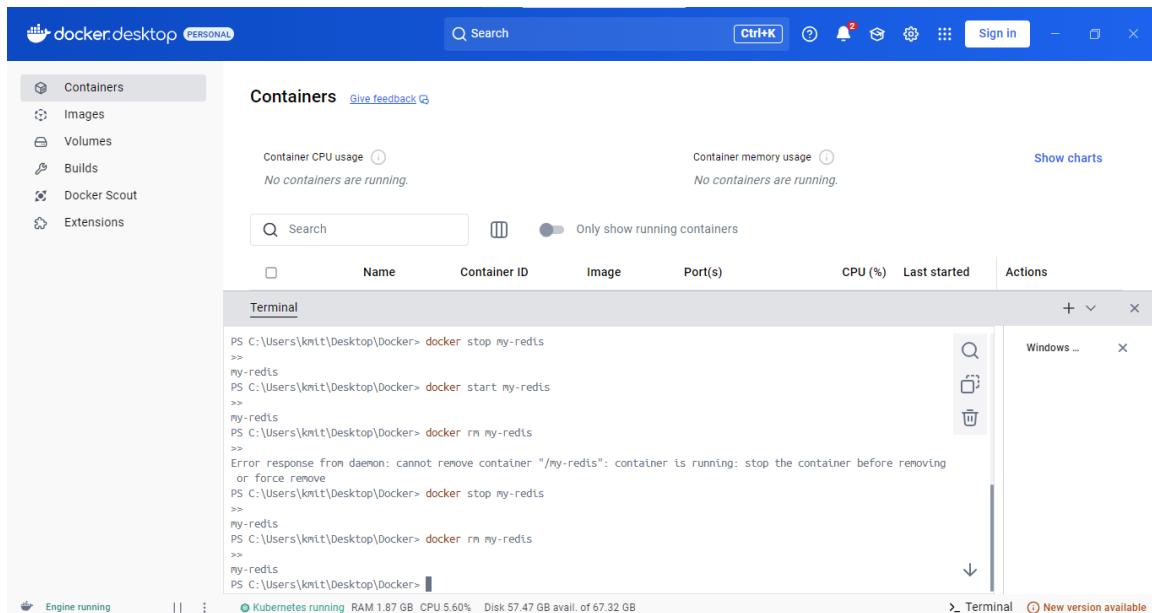
Step 7: Remove the Redis Container

Command:

`docker rm my-redis`

What It Does:

Deletes the container permanently.



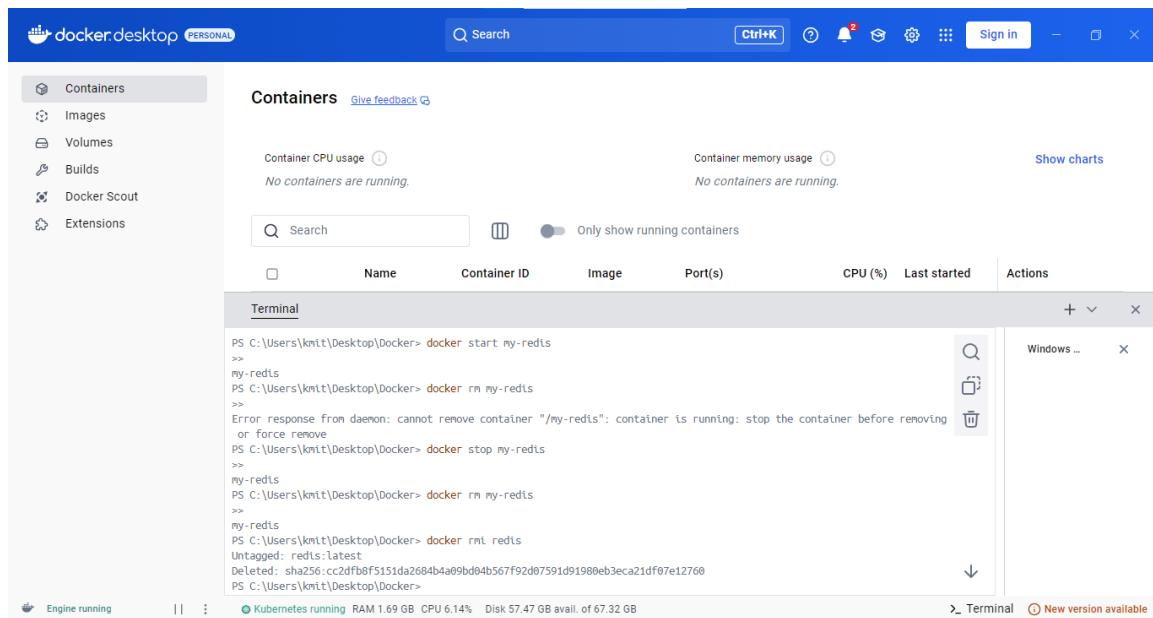
Step 8: Remove the Redis Image

Command:

```
docker rmi redis
```

What It Does:

- Deletes the Redis image from your local system.



2. Working with Docker file

A Docker file is a text file with instructions to create a custom Docker image.

Step 1: Set Up Your Folder

1. Windows:

- Create a folder like C:\DockerProjects\Redis.

- Open Git Bash and navigate to the folder:

```
cd /c/DockerProjects/Redis
```

2. Mac/Linux:

- Create a folder:

```
mkdir ~/DockerProjects/Redis
```

```
cd ~/DockerProjects/Redis
```

Step 2: Write the Dockerfile

1. Inside the folder, create a file named Dockerfile (no extension).

2. Add the following content:

```
FROM redis:latest
```

```
CMD ["redis-server"]
```

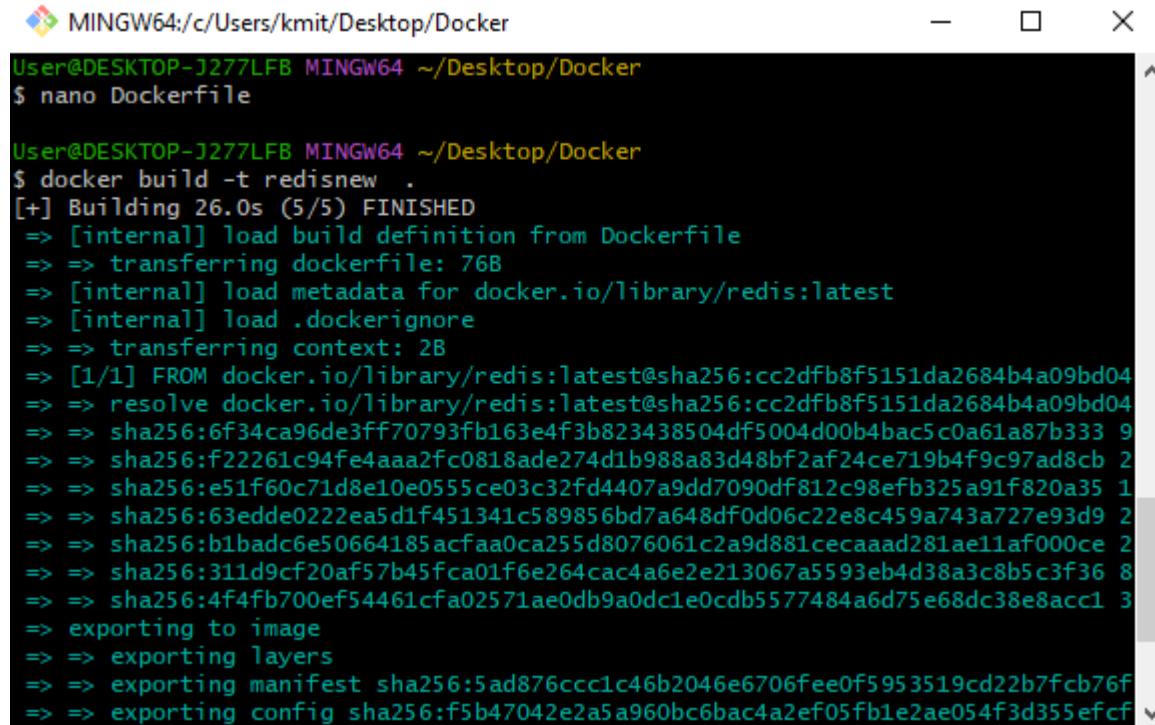
Docker Commands (Step-by-step):

1. **docker build -t redisnew .**

What it does:

① This creates (builds) a Docker image using the recipe (Dockerfile) in the current folder (.).

② -t redisnew: Gives the image a name/tag ("redisnew"), so you can find it easily.



```
MINGW64:/c/Users/kmit/Desktop/Docker
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ nano Dockerfile

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker build -t redisnew .
[+] Building 26.0s (5/5) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 76B
=> [internal] load metadata for docker.io/library/redis:latest
=> [internal] load .dockerrcignore
=> => transferring context: 2B
=> [1/1] FROM docker.io/library/redis@sha256:cc2dfb8f5151da2684b4a09bd04
=> => resolve docker.io/library/redis@sha256:cc2dfb8f5151da2684b4a09bd04
=> => sha256:6f34ca96de3ff70793fb163e4f3b823438504df5004d00b4bac5c0a61a87b333 9
=> => sha256:f22261c94fe4aaa2fc0818ade274d1b988a83d48bf2af24ce719b4f9c97ad8cb 2
=> => sha256:e51f60c71d8e10e0555ce03c32fd4407a9dd7090df812c98efb325a91f820a35 1
=> => sha256:63edde0222ea5d1f451341c589856bd7a648df0d06c22e8c459a743a727e93d9 2
=> => sha256:b1badc6e50664185acf0ca255d8076061c2a9d881cecaaad281ae11af000ce 2
=> => sha256:311d9cf20af57b45fc01f6e264cac4a6e2e213067a5593eb4d38a3c8b5c3f36 8
=> => sha256:4f4fb700ef54461cfa02571ae0db9a0dc1e0cdb5577484a6d75e68dc38e8acc1 3
=> exporting to image
=> => exporting layers
=> => exporting manifest sha256:5ad876ccc1c46b2046e6706fee0f5953519cd22b7fcb76f
=> => exporting config sha256:f5b47042e2a5a960bc6bac4a2ef05fb1e2ae054f3d355efcf
```

2. **docker run --name myredisnew -d redisnew**

What it does:

① Starts a new container (mini computer) from the redisnew image.

② --name myredisnew: Names the container "myredisnew" so it's easy to identify.

③ -d: Runs the container in the background.

```
MINGW64:/c/Users/kmit/Desktop/Docker
--> sha256:b1badc6e50664185acf0ca255d8076061c2a9d881cecaaad281ae11af000ce 2
--> sha256:311d9cf20af57b45fca01f6e264cac4a6e2e213067a5593eb4d38a3c8b5c3f36 8
--> sha256:4f4fb700ef54461cfa02571ae0db9a0dc1e0cd5577484a6d75e68dc38e8acc1 3
--> exporting to image
--> exporting layers
--> exporting manifest sha256:5ad876ccc1c46b2046e6706fee0f5953519cd22b7fc76f
--> exporting config sha256:f5b47042e2a5a960bc6bac4a2ef05fb1e2ae054f3d355efcf
--> exporting attestation manifest sha256:82942d1ea2122286606fb9ad836a37a0b63
--> exporting manifest list sha256:479c2a2f716056dcf47d41956210f52ecded3f0dc8
--> naming to docker.io/library/redisnew:latest
--> unpacking to docker.io/library/redisnew:latest

View build details: docker-desktop://dashboard/build/desktop-linux/desktop-linux

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ ^[[200~docker run --name myredisnew -d redisnew
bash: $'\E[200~docker': command not found

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker run --name myredisnew -d redisnew
856ae4657fa933973ccd4c1105089278b6644d3321a15a4181790d80f8a50b2c

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$
```

3. docker ps

What it does:

Shows a list of containers that are running right now.

```
MINGW64:/c/Users/kmit/Desktop/Docker
$ ^[[200~docker run --name myredisnew -d redisnew
bash: $'\E[200~docker': command not found

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker run --name myredisnew -d redisnew
856ae4657fa933973ccd4c1105089278b6644d3321a15a4181790d80f8a50b2c

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ ^[[200~. docker ps
bash: $'\E[200~.': command not found

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ . docker ps
/usr/bin/ps: /usr/bin/ps: cannot execute binary file

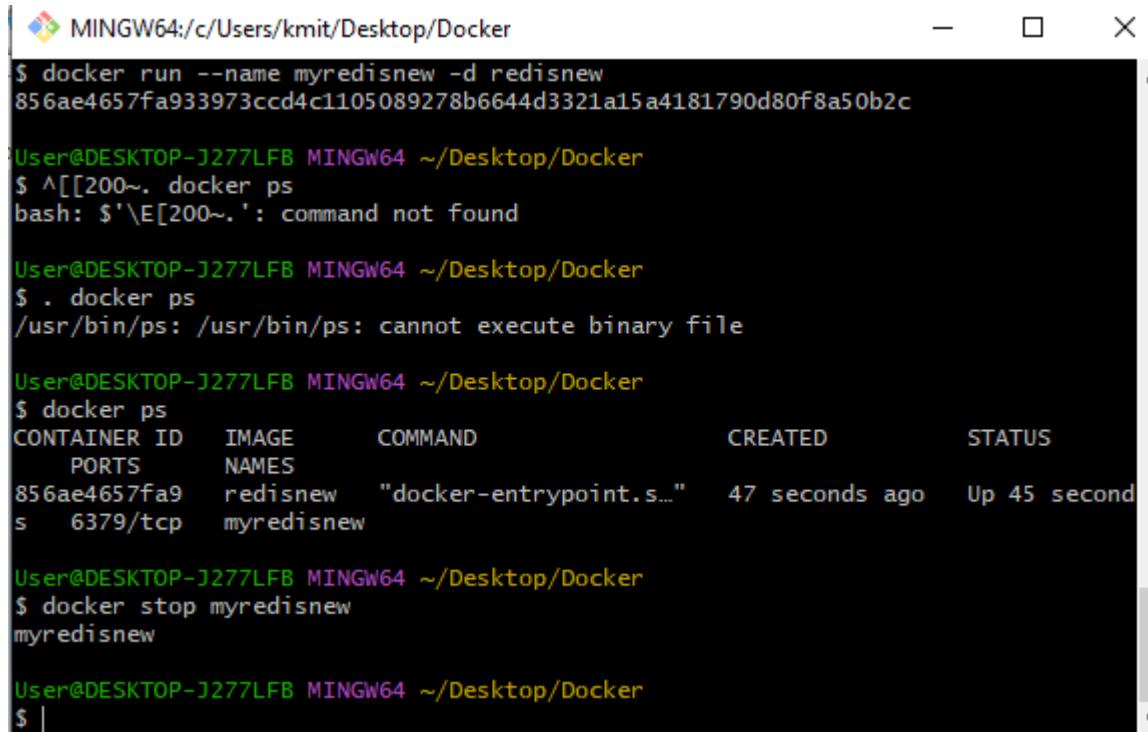
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker ps
CONTAINER ID IMAGE COMMAND CREATED STATUS
PORTS NAMES
856ae4657fa9 redisnew "docker-entrypoint.s..." 47 seconds ago Up 45 seconds
6379/tcp myredisnew

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$
```

4. docker stop myredisnew

What it does:

- stops the container named "myredisnew" (like turning off a computer).



The screenshot shows a terminal window titled "MINGW64:/c/Users/kmit/Desktop/Docker". The terminal output is as follows:

```
$ docker run --name myredisnew -d redisnew
856ae4657fa933973cccd4c1105089278b6644d3321a15a4181790d80f8a50b2c

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ ^[[200~. docker ps
bash: $'\E[200~.': command not found

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ . docker ps
/usr/bin/ps: /usr/bin/ps: cannot execute binary file

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker ps
CONTAINER ID IMAGE COMMAND CREATED STATUS
      PORTS NAMES
856ae4657fa9 redisnew "docker-entrypoint.s..." 47 seconds ago Up 45 seconds
s 6379/tcp myredisnew

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker stop myredisnew
myredisnew

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ |
```

5. docker login

What it does:

- Logs you into your Docker Hub account, so you can upload images.

```
MINGW64:/c/Users/kmit/Desktop/Docker
56ae4657fa9  redisnew  "docker-entrypoint.s..."  47 seconds ago  Up 45 seconds
s 6379/tcp  myredisnew

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker stop myredisnew
myredisnew

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker login

USING WEB-BASED LOGIN
To sign in with credentials on the command line, use 'docker login -u <username>'

Your one-time device confirmation code is: GNP0-FTFG
Press ENTER to open your browser or submit your device code here: https://login.docker.com/activate

Waiting for authentication in the browser...

Login Succeeded

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ |
```

6. docker ps -a

What it does:

- Shows a list of all containers, including stopped ones.

```
MINGW64:/c/Users/kmit/Desktop/Docker
Press ENTER to open your browser or submit your device code here: https://login.docker.com/activate

Waiting for authentication in the browser...

Login Succeeded

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker ps -a
CONTAINER ID   IMAGE          COMMAND       CR
EATED          STATUS         PORTS
NAMES
856ae4657fa9  redisnew      "docker-entrypoint.s..."  8
minutes ago    Exited (0) 7 minutes ago

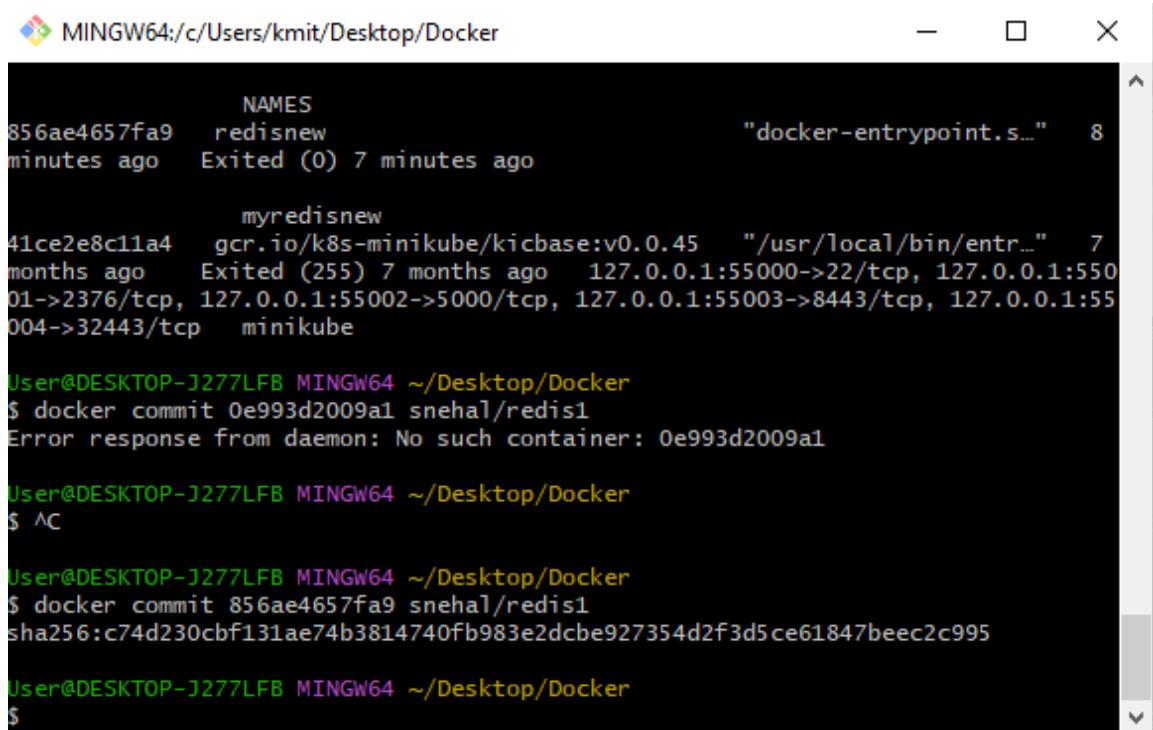
        myredisnew
41ce2e8c11a4  gcr.io/k8s-minikube/kicbase:v0.0.45  "/usr/local/bin/entr..."  7
months ago     Exited (255) 7 months ago  127.0.0.1:55000->22/tcp, 127.0.0.1:550
01->2376/tcp, 127.0.0.1:55002->5000/tcp, 127.0.0.1:55003->8443/tcp, 127.0.0.1:550
04->32443/tcp  minikube

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ |
```

7. docker commit 856ae4657fa9 snehal1729/redis1

What it does:

- ② Takes a snapshot (saves changes) of the container with ID 0e993d2009a1 and creates a new image called snehal1729/redis1.



The screenshot shows a terminal window titled "MINGW64:/c/Users/kmit/Desktop/Docker". The terminal output is as follows:

```
NAMES
856ae4657fa9  redisnew          "docker-entrypoint.s..."  8
    minutes ago   Exited (0) 7 minutes ago

        myredisnew
41ce2e8c11a4  gcr.io/k8s-minikube/kicbase:v0.0.45  "/usr/local/bin/entr..."  7
months ago      Exited (255) 7 months ago   127.0.0.1:55000->22/tcp, 127.0.0.1:550
01->2376/tcp, 127.0.0.1:55002->5000/tcp, 127.0.0.1:55003->8443/tcp, 127.0.0.1:55
004->32443/tcp  minikube

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker commit 0e993d2009a1 snehal1/redis1
Error response from daemon: No such container: 0e993d2009a1

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ ^C

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker commit 856ae4657fa9 snehal1/redis1
sha256:c74d230cbf131ae74b3814740fb983e2dcbe927354d2f3d5ce61847beec2c995

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$
```

8. docker images

What it does:

- ② Lists all images saved on your system.

```
MINGW64:/c/Users/kmit/Desktop/Docker
sha256:c74d230cbf131ae74b3814740fb983e2dcbe927354d2f3d5ce61847beec2c995

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker images
REPOSITORY          TAG      IMAGE ID      CREATED       SIZE
snehal/redis1        latest   c74d230cbf13  40 seconds ago  200MB
redisnew             latest   479c2a2f7160  13 hours ago   200MB
docker/desktop-kubernetes
-v1.29.0-cri-dockerd-v0.3.11-1-debian    kubernetes-v1.30.5-cni-v1.4.0-critools
registry.k8s.io/kube-apiserver           v1.30.5
                                             7a7b02256c8d  10 months ago  625MB
registry.k8s.io/kube-controller-manager   v1.30.5
                                             7746ea55ad74  11 months ago  153MB
registry.k8s.io/kube-scheduler            v1.30.5
                                             bbd15d267294  11 months ago  146MB
registry.k8s.io/kube-proxy                 v1.30.5
                                             62c91756a3c9  11 months ago  84.6MB
gcr.io/k8s-minikube/kicbase              <none>
                                             fa20f91153b9  11 months ago  118MB
gcr.io/k8s-minikube/kicbase              v0.0.45
                                             e7c9bc3bc515  11 months ago  1.81GB
                                             10. docker run -it --rm redisnew:latest
```

9. docker push snehal1729/redis1

What it does:

- Uploads the image 1729/redis1 to Docker Hub, so others can download it.

```
MINGW64:/c/Users/kmit/Desktop/Docker
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker push snehal1729/redis1:latest
The push refers to repository [docker.io/snehal1729/redis1]
tag does not exist: snehal1729/redis1:latest

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker tag redisnew snehal1729/redis1:latest

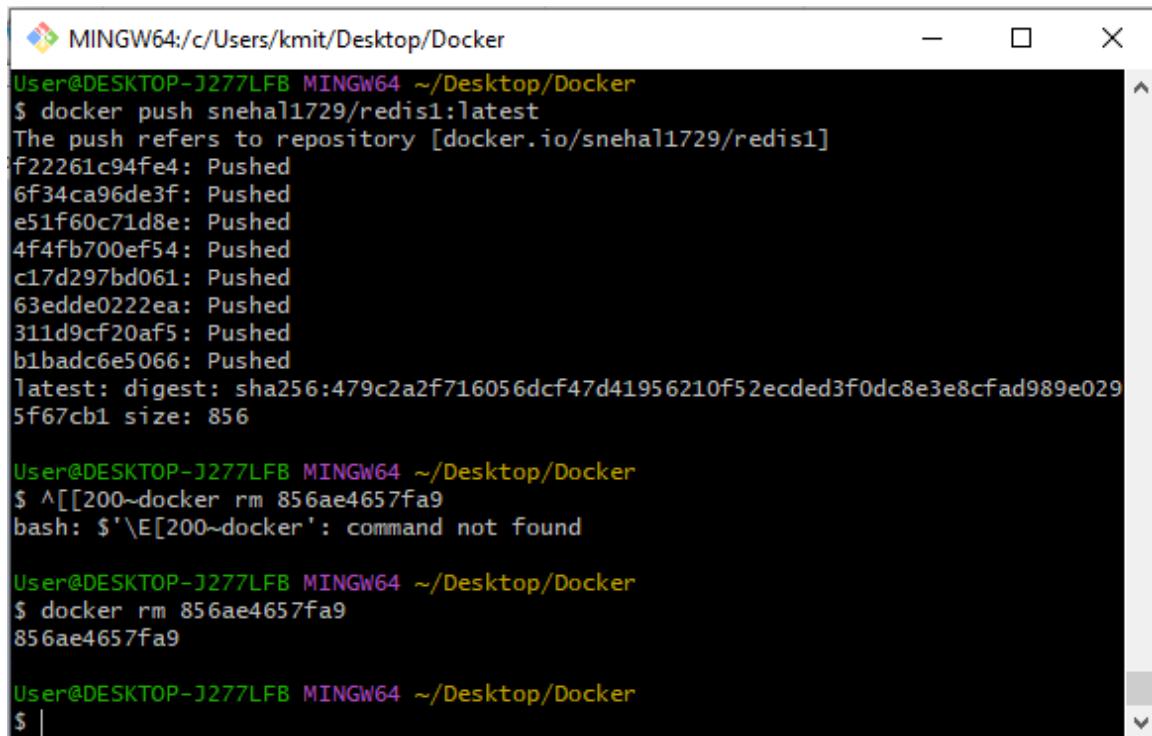
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker push snehal1729/redis1:latest
The push refers to repository [docker.io/snehal1729/redis1]
f22261c94fe4: Pushed
6f34ca96de3f: Pushed
e51f60c71d8e: Pushed
4f4fb700ef54: Pushed
c17d297bd061: Pushed
63edde0222ea: Pushed
311d9cf20af5: Pushed
b1badc6e5066: Pushed
latest: digest: sha256:479c2a2f716056dcf47d41956210f52ecded3f0dc8e3e8cfad989e029
5f67cb1 size: 856

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$
```

10. docker rm 856ae4657fa9

What it does:

- Deletes the container with ID 0e993d2009a1.



The screenshot shows a terminal window titled "MINGW64:/c/Users/kmit/Desktop/Docker". The user has run several Docker commands:

- \$ docker push snehal1729/redis1:latest
- The push refers to repository [docker.io/snehal1729/redis1]
- f22261c94fe4: Pushed
- 6f34ca96de3f: Pushed
- e51f60c71d8e: Pushed
- 4f4fb700ef54: Pushed
- c17d297bd061: Pushed
- 63edde0222ea: Pushed
- 311d9cf20af5: Pushed
- b1badc6e5066: Pushed
- latest: digest: sha256:479c2a2f716056dcf47d41956210f52ecded3f0dc8e3e8cfad989e0295f67cb1 size: 856

Then, the user attempts to delete a container:

- \$ docker rm 856ae4657fa9
- bash: \$'\E[200~docker': command not found

Finally, the user attempts to delete an image:

- \$ docker rmi budarajumadhurika/redis1

11. docker rmi budarajumadhurika/redis1

What it does:

- Deletes the image budarajumadhurika/redis1 from your system.

```
MINGW64:/c/Users/kmit/Desktop/Docker
5f67cb1 size: 856

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ ^[[200~docker rm 856ae4657fa9
bash: $'\E[200~docker': command not found

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker rm 856ae4657fa9
856ae4657fa9

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker rmi snehal1729/redis1
Untagged: snehal1729/redis1:latest

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker rmi snehal/redis1
Untagged: snehal/redis1:latest

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker rmi snehal1729/redis1
Error response from daemon: No such image: snehal1729/redis1:latest

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$
```

12. docker ps -a

What it does:

Shows all containers again to confirm changes.

```
MINGW64:/c/Users/kmit/Desktop/Docker
$ docker rmi snehal1729/redis1
Untagged: snehal1729/redis1:latest

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker rmi snehal/redis1
Untagged: snehal/redis1:latest

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker rmi snehal1729/redis1
Error response from daemon: No such image: snehal1729/redis1:latest

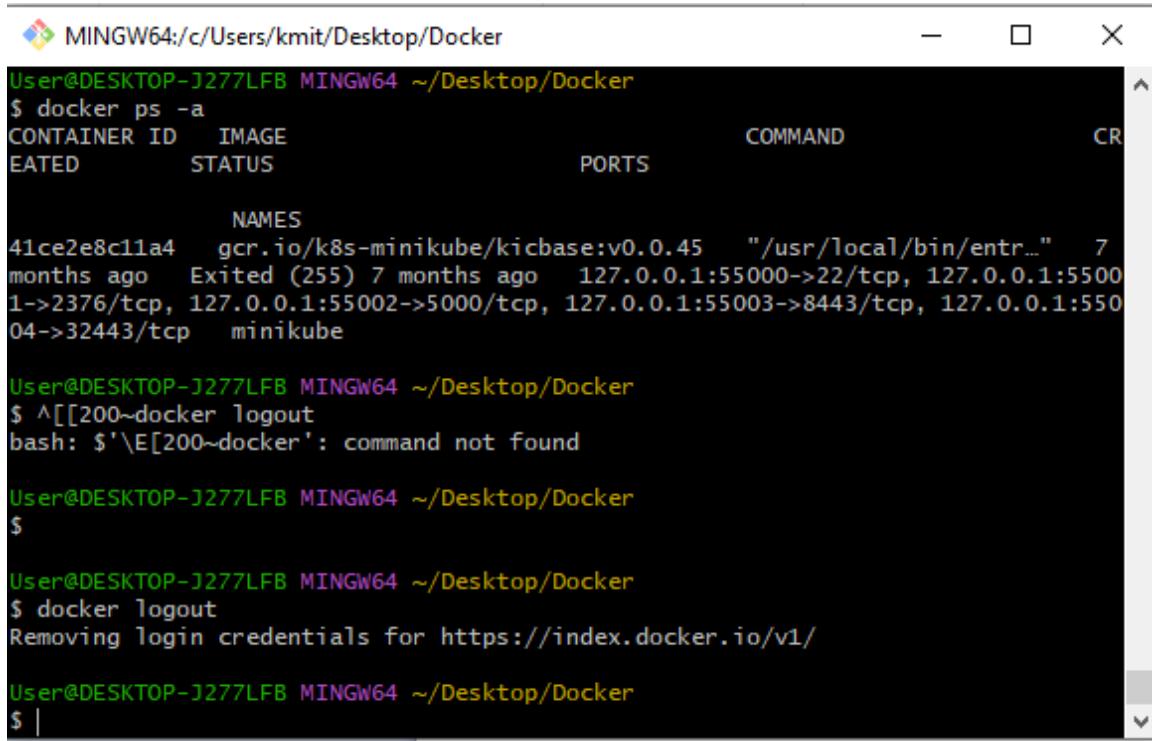
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker ps -a
CONTAINER ID        IMAGE               COMMAND                  CR
EATED             STATUS              PORTS
NAMES
41ce2e8c11a4      gcr.io/k8s-minikube/kicbase:v0.0.45   "/usr/local/bin/entr..."  7
months ago         Exited (255) 7 months ago   127.0.0.1:55000->22/tcp, 127.0.0.1:5500
1->2376/tcp, 127.0.0.1:55002->5000/tcp, 127.0.0.1:55003->8443/tcp, 127.0.0.1:550
04->32443/tcp    minikube

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$
```

13. docker logout

What it does:

- Logs you out of Docker Hub.



```
MINGW64:/c/Users/kmit/Desktop/Docker
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker ps -a
CONTAINER ID   IMAGE          COMMAND
EATED          STATUS          PORTS
NAMES
41ce2e8c11a4   gcr.io/k8s-minikube/kicbase:v0.0.45   "/usr/local/bin/entr..."  7
months ago     Exited (255) 7 months ago   127.0.0.1:55000->22/tcp, 127.0.0.1:5500
1->2376/tcp, 127.0.0.1:55002->5000/tcp, 127.0.0.1:55003->8443/tcp, 127.0.0.1:550
04->32443/tcp minikube

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ ^[[200~docker logout
bash: $'\E[200~docker': command not found

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker logout
Removing login credentials for https://index.docker.io/v1/

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ |
```

14. docker pull snehal1729/redis1

What it does:

- Downloads the image budarajumadhurika/redis1 from Docker Hub.

```
MINGW64:/c/Users/kmit/Desktop/Docker
1->2376/tcp, 127.0.0.1:55002->5000/tcp, 127.0.0.1:55003->8443/tcp, 127.0.0.1:550
04->32443/tcp minikube

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ ^[[200~docker logout
bash: $'\E[200~docker': command not found

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker logout
Removing login credentials for https://index.docker.io/v1/

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker pull snehal1729/redis1
Using default tag: latest
latest: Pulling from snehal1729/redis1
Digest: sha256:479c2a2f716056dcf47d41956210f52ecded3f0dc8e3e8cfad989e0295f67cb1
Status: Downloaded newer image for snehal1729/redis1:latest
docker.io/snehal1729/redis1:latest

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ |
```

15. docker run --name myredis -d snehal1729/redis1

What it does:

- Starts a new container using the image budarajumadhurika/redis1.

```
$ ^[[200~docker logout
bash: $'\E[200~docker': command not found

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ 

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker logout
Removing login credentials for https://index.docker.io/v1/

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker pull snehal1729/redis1
Using default tag: latest
latest: Pulling from snehal1729/redis1
Digest: sha256:479c2a2f716056dcf47d41956210f52ecded3f0dc8e3e8cfad989e0295f67cb1
Status: Downloaded newer image for snehal1729/redis1:latest
docker.io/snehal1729/redis1:latest

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker run --name myredis -d snehal1729/redis1
9369876d40486fe9671b7a5e9c0db0affa921683fdd7b993b751c8479a8e73ea

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ |
```

16. docker exec -it myredis redis-cli

What it does:

- Opens the Redis command-line interface (like a terminal) inside the running container myredis.

The screenshot shows a terminal window titled "MINGW64:/c/Users/kmit/Desktop/Docker". The terminal output is as follows:

```
bash: $'\E[200~docker': command not found
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ 
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker logout
Removing login credentials for https://index.docker.io/v1/
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker pull snehal1729/redis1
Using default tag: latest
latest: Pulling from snehal1729/redis1
Digest: sha256:479c2a2f716056dcf47d41956210f52ecded3f0dc8e3e8cfad989e0295f67cb1
Status: Downloaded newer image for snehal1729/redis1:latest
docker.io/snehal1729/redis1:latest

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker run --name myredis -d snehal1729/redis1
9369876d40486fe9671b7a5e9c0db0affa921683fdd7b993b751c8479a8e73ea

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker exec -it myredis redis-cli
127.0.0.1:6379>
```

17. SET name "Abcdef"

What it does:

- Saves a key-value pair in Redis (key = name, value = Abcdef).

```
MINGW64:/c/Users/kmit/Desktop/Docker
ewer' 'image' 'for' 'snehal1729/redis1:latest'
127.0.0.1:6379> docker.io/snehal1729/redis1:latest
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker(error) ERR unknown command 'docker
.io/snehal1729/redis1:latest', with args beginning with:
127.0.0.1:6379>
127.0.0.1:6379> User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker

$ docker exec -it myredis redis-cli
127.0.0.1:6379>
(error) ERR unknown command 'User@DESKTOP-J277LFB', with args beginning with: 'M
INGW64' '~/Desktop/Docker'
127.0.0.1:6379> $ docker run --name myredis -d snehal1729/redis1
(error) ERR unknown command '$', with args beginning with: 'docker' 'run' '--nam
e' 'myredis' '-d' 'snehal1729/redis1'
127.0.0.1:6379> 9369876d40486fe9671b7a5e9c0db0affa921683fdd7b993b751c8479a8e73ea
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker$ docker exec -it myredis redis-cl
127.0.0.1:6379>SET name "snehal"
(error) ERR unknown command '9369876d40486fe9671b7a5e9c0db0affa921683fdd7b993b75
1c8479a8e73eaUser@DESKTOP-J277LFB', with args beginning with: 'MINGW64' '~/Deskt
op/Docker$' 'docker' 'exec' '-it' 'myredis' 'redis-cl127.0.0.1:6379>SET' 'name'
'snehal'
127.0.0.1:6379> SET name "Snehal"
OK
127.0.0.1:6379> |
```

18. GET name

What it does:

Retrieves the value of the key name from Redis (it will return "Abcdef").

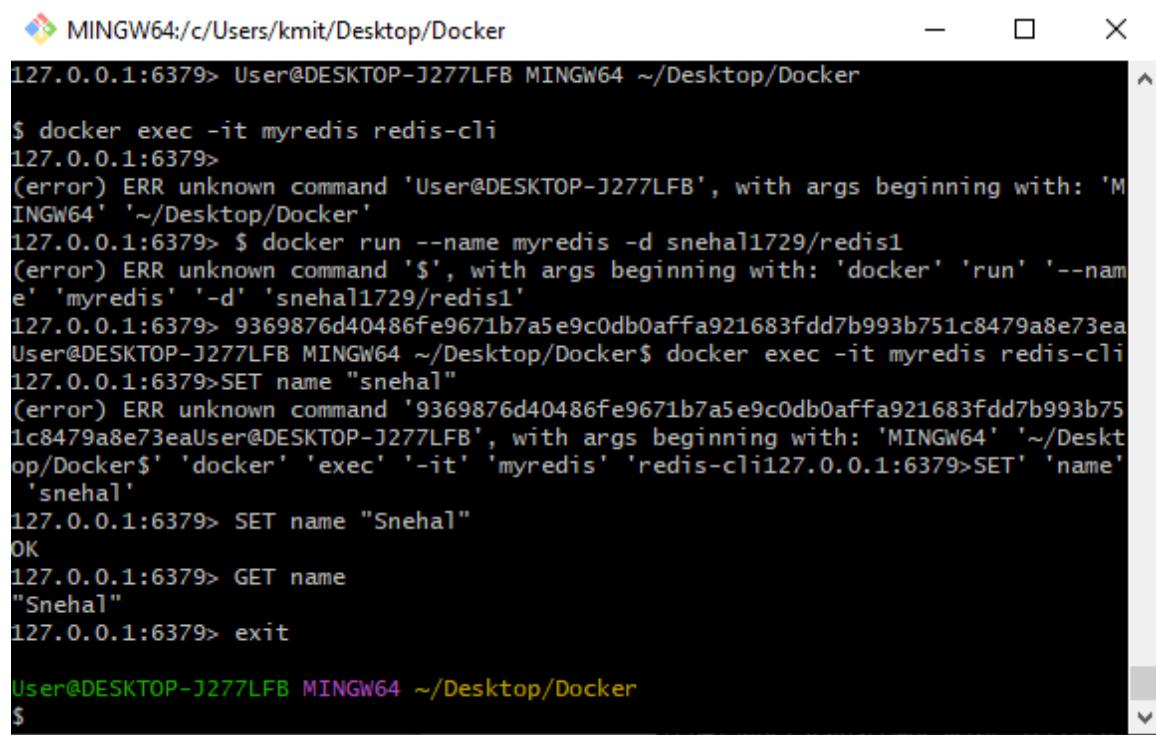
```
MINGW64:/c/Users/kmit/Desktop/Docker
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker(error) ERR unknown command 'docker
.io/snehal1729/redis1:latest', with args beginning with:
127.0.0.1:6379>
127.0.0.1:6379> User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker

$ docker exec -it myredis redis-cli
127.0.0.1:6379>
(error) ERR unknown command 'User@DESKTOP-J277LFB', with args beginning with: 'M
INGW64' '~/Desktop/Docker'
127.0.0.1:6379> $ docker run --name myredis -d snehal1729/redis1
(error) ERR unknown command '$', with args beginning with: 'docker' 'run' '--nam
e' 'myredis' '-d' 'snehal1729/redis1'
127.0.0.1:6379> 9369876d40486fe9671b7a5e9c0db0affa921683fdd7b993b751c8479a8e73ea
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker$ docker exec -it myredis redis-cl
127.0.0.1:6379>SET name "snehal"
(error) ERR unknown command '9369876d40486fe9671b7a5e9c0db0affa921683fdd7b993b75
1c8479a8e73eaUser@DESKTOP-J277LFB', with args beginning with: 'MINGW64' '~/Deskt
op/Docker$' 'docker' 'exec' '-it' 'myredis' 'redis-cl127.0.0.1:6379>SET' 'name'
'snehal'
127.0.0.1:6379> SET name "Snehal"
OK
127.0.0.1:6379> GET name
"Snehal"
127.0.0.1:6379> |
```

19. exit

What it does:

- Exits the Redis CLI.



The screenshot shows a terminal window titled "MINGW64:/c/Users/kmit/Desktop/Docker". The user is connected to a Redis instance at 127.0.0.1:6379. They attempt to run a Docker command inside the Redis container, which fails due to syntax errors. Then, they use the Redis `SET` command to set a key-value pair. Finally, they use the Redis `GET` command to retrieve the value, and finally exit the Redis CLI.

```
127.0.0.1:6379> User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker exec -it myredis redis-cli
127.0.0.1:6379>
(error) ERR unknown command 'User@DESKTOP-J277LFB', with args beginning with: 'MINGW64' '~/Desktop/Docker'
127.0.0.1:6379> $ docker run --name myredis -d snehal1729/redis1
(error) ERR unknown command '$', with args beginning with: 'docker' 'run' '--name' 'myredis' '-d' 'snehal1729/redis1'
127.0.0.1:6379> 9369876d40486fe9671b7a5e9c0db0affa921683fdd7b993b751c8479a8e73ea
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker$ docker exec -it myredis redis-cli
127.0.0.1:6379>SET name "snehal"
(error) ERR unknown command '9369876d40486fe9671b7a5e9c0db0affa921683fdd7b993b751c8479a8e73ea
User@DESKTOP-J277LFB', with args beginning with: 'MINGW64' '~/Desktop/Docker$ docker exec -it myredis redis-cli
127.0.0.1:6379>SET name 'snehal'
127.0.0.1:6379> SET name "Snehal"
OK
127.0.0.1:6379> GET name
"Snehal"
127.0.0.1:6379> exit
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$
```

20. docker ps -a

What it does:

- Shows all containers again to check their status.

```
MINGW64:/c/Users/kmit/Desktop/Docker
'snehal'
127.0.0.1:6379> SET name "Snehal"
OK
127.0.0.1:6379> GET name
"Snehal"
127.0.0.1:6379> exit

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker ps -a
CONTAINER ID   IMAGE          COMMAND
EATED          STATUS          PORTS
NAMES
9369876d4048  snehal1729/redis1   "docker-entrypoint.s..."  4
minutes ago    Up 3 minutes      6379/tcp

myredis
41ce2e8c11a4  gcr.io/k8s-minikube/kicbase:v0.0.45  "/usr/local/bin/entr..."  7
months ago     Exited (255) 7 months ago  127.0.0.1:55000->22/tcp, 127.0.0.1:550
01->2376/tcp, 127.0.0.1:55002->5000/tcp, 127.0.0.1:55003->8443/tcp, 127.0.0.1:55
004->32443/tcp  minikube

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ |
```

21. docker stop myredis

What it does:

stops the container myredis.

```
MINGW64:/c/Users/kmit/Desktop/Docker
"Snehal"
127.0.0.1:6379> exit

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker ps -a
CONTAINER ID   IMAGE          COMMAND
EATED          STATUS          PORTS
NAMES
9369876d4048  snehal1729/redis1   "docker-entrypoint.s..."  4
minutes ago    Up 3 minutes      6379/tcp

myredis
41ce2e8c11a4  gcr.io/k8s-minikube/kicbase:v0.0.45  "/usr/local/bin/entr..."  7
months ago     Exited (255) 7 months ago  127.0.0.1:55000->22/tcp, 127.0.0.1:550
01->2376/tcp, 127.0.0.1:55002->5000/tcp, 127.0.0.1:55003->8443/tcp, 127.0.0.1:55
004->32443/tcp  minikube

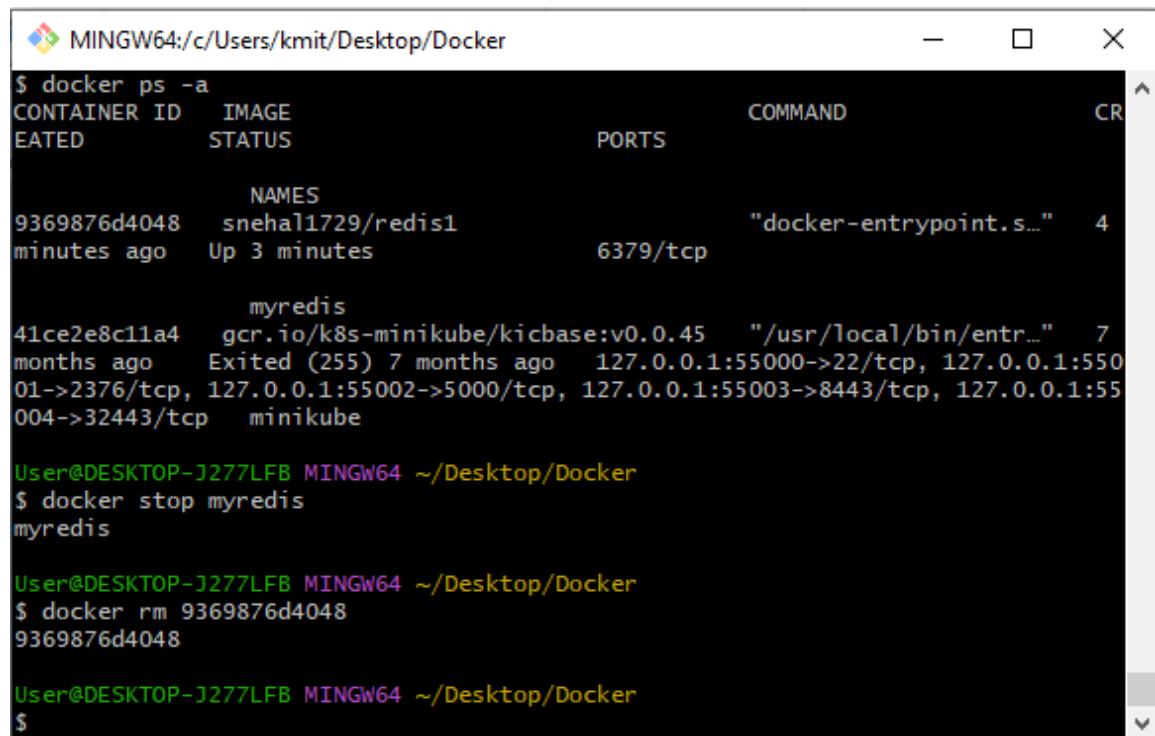
User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker stop myredis
myredis

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ |
```

22. docker rm 50a6e4a9c326

What it does:

- Deletes the container with ID 50a6e4a9c326.



The screenshot shows a terminal window titled "MINGW64:/c/Users/kmit/Desktop/Docker". It displays the output of several Docker commands:

```
$ docker ps -a
CONTAINER ID   IMAGE          COMMAND
9369876d4048   snehal1729/redis1   "docker-entrypoint.s..."  4
               NAMES
               myredis
41ce2e8c11a4   gcr.io/k8s-minikube/kicbase:v0.0.45   "/usr/local/bin/entr..."  7
months ago     Exited (255) 7 months ago  127.0.0.1:55000->22/tcp, 127.0.0.1:550
01->2376/tcp, 127.0.0.1:55002->5000/tcp, 127.0.0.1:55003->8443/tcp, 127.0.0.1:5504->32443/tcp  minikube

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker stop myredis
myredis

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker rm 9369876d4048
9369876d4048

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$
```

23. docker images

What it does:

- Lists all images again to confirm which ones remain.

```
MINGW64:/c/Users/kmit/Desktop/Docker
9369876d4048

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker images
REPOSITORY          TAG      IMAGE ID      CREATED       SIZE
<none>            <none>    c74d230cbf13  About an hour ago  200MB
redisnew           latest   479c2a2f7160  14 hours ago   200MB
snehal1729/redis1  latest   479c2a2f7160  14 hours ago   200MB
docker/desktop-kubernetes
-v1.29.0-cri-dockerd-v0.3.11-1-debian
registry.k8s.io/kube-apiserver        v1.30.5
                                         7a7b02256c8d  10 months ago  625MB
registry.k8s.io/kube-scheduler        v1.30.5
                                         7746ea55ad74  11 months ago  153MB
                                         62c91756a3c9  11 months ago  84.6M
B
registry.k8s.io/kube-controller-manager v1.30.5
                                         bbd15d267294  11 months ago  146MB
registry.k8s.io/kube-proxy             v1.30.5
                                         fa20f91153b9  11 months ago  118MB
gcr.io/k8s-minikube/kicbase          v0.0.45
```

24. docker rmi snehal1729/redis1

What it does:

Deletes the image budarajumadhurika/redis1 again.

```
MINGW64:/c/Users/kmit/Desktop/Docker
B
registry.k8s.io/coredns/coredns        v1.11.3
                                         9caabbf6238b  12 months ago  85.1M
B
registry.k8s.io/etcfd                 3.5.12-0
                                         44a8e24dcba   18 months ago  211MB
docker/desktop-vpnkit-controller
6e
registry.k8s.io/pause                  3.9
                                         7ecf567ea070  2 years ago   47MB
                                         7031c1b28338  2 years ago   1.07M
B
docker/desktop-storage-provisioner    v2.0
                                         115d77efe6e2  4 years ago   59.2M
B
registry.k8s.io/etcfd                 3.5.15-0
                                         a6dc63e6e8cf  55 years ago  56.9M
B

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ docker rmi snehal1729/redis1
Untagged: snehal1729/redis1:latest

User@DESKTOP-J277LFB MINGW64 ~/Desktop/Docker
$ |
```

Step 3: Remove Login Credentials (Optional)

If you no longer need to be logged in, you can log out:

`docker logout`

What It Does:

- ② Logs you out from Docker Hub and removes your stored credentials.

Scenario based Questions:

1. Your application is running inside Docker, but you're not sure if the container is active. Which command helps you check running containers?

ANS. `docker ps`

2. You notice one of your containers is consuming high CPU and want to stop it. How do you stop a running container named `web_app`?

ANS. `docker stop web_app`

3. You've written a Dockerfile in your project directory and want to build an image named `myapi`.

What Docker CLI command do you use?

ANS. `docker build -t myapi .`

4. You want to run your `web_app` container and expose its internal port 5000 on host port 8080. Which command should you use?

ANS. `docker run -p 8080:5000 web_app`

5. Your container is running, but you want to go inside and debug it using a shell. What CLI command helps you enter the container's shell?

ANS. `docker exec -it <container_name_or_id> bash`

6. You've built a few test images and want to delete one called `old_api`. How would you do this?

ANS. `docker rmi old_api`

7. Your app crashed in the container. You need to check the logs. How do you check this?

ANS. `docker logs <container_name_or_id>`

8. You want to start a container and let it run in the background. Which command should you use?

ANS. `docker run -d <image_name>`

9. You're not sure which container is using port 3000. How do you check?

ANS. `docker ps --filter "publish=3000"`

10. You built an image but forgot to tag it. How do you tag it now?

ANS. `docker tag <existing_image_id_or_name> <new_repo>:<new_tag>`

11. You want to export an image to a .tar file for offline transfer. How do you export it?

ANS. docker save -o image.tar <image_name_or_id>

12. You want to restart a container automatically if it crashes. What Docker CLI command do you use?

ANS. docker run --restart unless-stopped <image_name>

13. You want to see how much CPU/RAM each container uses. What Docker CLI command do you use? How to Control or Limit RAM in Docker? How to Measure Actual Memory Usage?

ANS. docker stats

 docker run -m 512m <image>

14. You want your container to run a shell command before starting the app. What do you include in docker file?

ANS. CMD /bin/sh -c "your-command && exec your-app"

15. You want to create a custom Nginx image with your static files. What do you do?

ANS. FROM nginx:latest

 COPY ./static /usr/share/nginx/html

Conclusion:

Through this hands-on exploration of Docker commands and Dockerfile usage, students will get a clear picture about the end-to-end container workflow—from pulling images, creating and running containers, to customizing them with Dockerfiles. By practicing key commands such as docker build, docker run, docker ps, and docker exec, students will gain practical skills in managing containerized environments. The Dockerfile exercise reinforced the importance of layered builds, efficient image creation, and environment reproducibility. Overall, these activities strengthened the understanding of containerization principles, enabling building, testing and deploying applications in isolated, consistent environments.

```
C:\Windows\System32\cmd.exe
Microsoft Windows [Version 10.0.19045.6216]
(c) Microsoft Corporation. All rights reserved.

C:\Users\kmit\Desktop\Dockers-compose>docker-compose up -d
[+] Running 1/2
  Container docker-compose-db-1      Running          0.0s
  - Container docker-compose-wordpress-1 Starting          0.0s
Error response from daemon: ports are not available: exposing port TCP 0.0.0.0:8080 -> 127.0.0.1:0: listen tcp 0.0.0.0:8080: bind: Only one usage of each socket address (protocol/network address/port) is normally permitted.

C:\Users\kmit\Desktop\Dockers-compose>
```

```
C:\Windows\System32\cmd.exe
Microsoft Windows [Version 10.0.19045.6216]
(c) Microsoft Corporation. All rights reserved.

C:\Users\kmit\Desktop\Dockers-compose>docker-compose up -d
[+] Running 1/2
  Container docker-compose-db-1      Running          0.0s
  - Container docker-compose-wordpress-1 Starting          0.0s
Error response from daemon: ports are not available: exposing port TCP 0.0.0.0:8080 -> 127.0.0.1:0: listen tcp 0.0.0.0:8080: bind: Only one usage of each socket address (protocol/network address/port) is normally permitted.

C:\Users\kmit\Desktop\Dockers-compose>docker-compose down
[+] Running 3/3
  Container docker-compose-wordpress-1 Removed          0.2s
  Container docker-compose-db-1     Removed          1.9s
  Network docker-compose_default   Removed          0.2s

C:\Users\kmit\Desktop\Dockers-compose>
```

```
C:\Windows\System32\cmd.exe
C:\Users\kmit\Desktop\Dockers-compose>docker-compose up --scale wordpress=2 -d
[+] Running 3/4
  Network docker-compose_default    Created          0.1s
  Container docker-compose-db-1     Started          0.7s
  Container docker-compose-wordpress-2 Created          0.2s
  - Container docker-compose-wordpress-1 Starting          0.6s
Error response from daemon: ports are not available: exposing port TCP 0.0.0.0:8080 -> 127.0.0.1:0: listen tcp 0.0.0.0:8080: bind: Only one usage of each socket address (protocol/network address/port) is normally permitted.

C:\Users\kmit\Desktop\Dockers-compose>
```

```

● PS C:\Users\kmit\Desktop\Compose-lab> docker compose up -d
time="2025-08-26T10:42:20+05:30" level=warning msg="C:\\\\Users\\\\kmit\\\\Desktop\\\\Compose-lab\\\\docker-compose.yml: the attribute `version` is obsolete, it will be ignored, please remove it to avoid potential confusion"
[+] Running 2/2
  ✓ Container compose-lab-db-1  Running          0.0s
  ✓ Container compose-lab-web-1 Started          1.1s
○ PS C:\Users\kmit\Desktop\Compose-lab>

```

localhost:9080

Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
 Commercial support is available at nginx.com.

Thank you for using nginx.

Transfer it to another machine with Docker Compose installed.

```

PS C:\Users\kmit\Desktop\Compose-lab> docker compose up -d
time="2025-08-26T10:46:27+05:30" level=warning msg="C:\\\\Users\\\\kmit\\\\Desktop\\\\Compose-lab\\\\docker-compose.yml: the attribute `version` is obsolete, it will be ignored, please remove it to avoid potential confusion"
[+] Running 3/3
  ✓ Container compose-lab-db-1  Running          0.0s
  ✓ Container compose-lab-redis-1 Running          0.0s
  ✓ Container compose-lab-web-1  Running          0.0s
PS C:\Users\kmit\Desktop\Compose-lab>

```

```

C:\Users\kmit\Desktop\Compose-lab> docker compose ps
time="2025-08-26T10:47:39+05:30" level=warning msg="C:\\\\Users\\\\kmit\\\\Desktop\\\\Compose-lab\\\\docker-compose.yml: the attribute `version` is obsolete, please remove it to avoid potential confusion"
NAME           IMAGE        COMMAND       SERVICE      CREATED     STATUS      PORTS
compose-lab-db-1  postgres:15  "docker-entrypoint.s..."  db      11 minutes ago  Up 11 minutes  5432/tcp
compose-lab-redis-1  redis:alpine  "docker-entrypoint.s..."  redis    2 minutes ago  Up 2 minutes  6379/tcp
compose-lab-web-1   nginx:latest  "/docke...r-entrypoint.s..."  web      5 minutes ago  Up 5 minutes  0.0.0.0:9080->80/tcp, [::]:9080->80/tcp

```

Check that Nginx and Postgres work there as well.

```
10      ▷ Run Service
11  db:
12    image: postgres:15
13    environment:
14      POSTGRES_USER: demo
15      POSTGRES_PASSWORD: demo
16      POSTGRES_DB: demo_db
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS powershell + × ⌂ ⌂ ... [] ×

```
● PS C:\Users\kmit\Downloads\compose-lab> docker compose up -d
time="2025-08-26T10:55:15+05:30" level=warning msg="C:\\Users\\kmit\\Downloads\\compose-la
b\\docker-compose.yml: the attribute 'version' is obsolete, it will be ignored, please rem
ove it to avoid potential confusion"
[+] Running 3/3
  ✓ Container compose-lab-db-1    Running          0.0s
  ✓ Container compose-lab-redis-1  Running          0.0s
  ✓ Container compose-lab-web-1   Started          2.1s
○ PS C:\Users\kmit\Downloads\compose-lab> 
```

← → ⌂ ⌂ localhost:8081

Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.

All services run correctly on the new host.

```
PS C:\Users\kmit\Downloads\compose-lab> docker compose up -d
[+] Running 5/5
  ✓ Network compose-lab_app-net    Created          0.1s
  ✓ Volume "compose-lab_db-data"   Created          0.0s
  ✓ Container compose-lab-db-1    Started          2.5s
  ✓ Container compose-lab-redis-1  Started          2.4s
  ✓ Container compose-lab-web-1   Started          1.9s
PS C:\Users\kmit\Downloads\compose-lab> 
```

```
✓ docker-compose.yml X
  ↴ docker-compose.yml
  7   services:
 17     db:
 19       environment:
 20         POSTGRES_USER: demo
 21         POSTGRES_PASSWORD: demo
 22         POSTGRES_DB: demo_db
 23       volumes:
 24         - db-data:/var/lib/postgresql/data
 25       networks:
 26         - app-net
 27       ports:
 28         - "9090:4060"
 29
 30   ▶ Run Service
 31   redis:
 32     image: redis:alpine
 33     networks:
 34       - app-net
```

. Insert some data into Postgres

Then

```
✓ Container compose-lab-web-1    Started          2.4s
PS C:\Users\kmit\Downloads\compose-lab> docker compose down
[+] Running 4/4
✓ Container compose-lab-redis-1  Removed         0.4s
✓ Container compose-lab-web-1   Removed         0.4s
✓ Container compose-lab-db-1   Removed         0.4s
✓ Network compose-lab_app-net  Removed         0.2s
PS C:\Users\kmit\Downloads\compose-lab>
```

```
PS C:\Users\kmit\Downloads\compose-lab> docker compose up -d
[+] Running 3/3
  ✓ Container compose-lab-redis-1  Running          0.05s
  ✓ Container compose-lab-db-1    Running          0.05s
  ✓ Container compose-lab-web-1  Running          0.05s
PS C:\Users\kmit\Downloads\compose-lab> docker compose ps
NAME           IMAGE      COMMAND      SERVICE      CREATED      STATUS
S              PORTS
PS C:\Users\kmit\Downloads\compose-lab> docker compose up -d
time="2025-08-26T11:05:38+05:30" level=warning msg="Found orphan containers ([compose-lab-redis-1]) for this project. If you removed or renamed this service in your compose file, you can run this command with the --remove-orphans flag to clean it up."
[+] Running 2/2
  ✓ Container compose-lab-db-1  Running          0.0s
  ✓ Container compose-lab-web-1 Started         1.1s
PS C:\Users\kmit\Downloads\compose-lab>
```

```

PS C:\Users\kmit\Downloads\compose-lab> docker compose ps
NAME           IMAGE          COMMAND                  SERVICE    CREATED        STATUS
5              PORTS
compose-lab-db-1  postgres:15  "docker-entrypoint.s..."  db         50 seconds ago  Up 49s
seconds          0.0.0.0:9090->5432/tcp, [::]:9090->5432/tcp
compose-lab-redis-1  redis:alpine  "docker-entrypoint.s..."  redis      7 minutes ago   Up 7m
minutes          6379/tcp
compose-lab-web-1   nginx:latest  "/docker-entrypoint..."  web       About a minute ago  Up Ab
out a minute     0.0.0.0:8081->80/tcp, [::]:8081->80/tcp
PS C:\Users\kmit\Downloads\compose-lab>

#10 naming to docker.io/library/compose-lab-web:latest
#10 naming to docker.io/library/compose-lab-web:latest done
#10 unpacking to docker.io/library/compose-lab-web:latest
#10 unpacking to docker.io/library/compose-lab-web:latest 0.4s done
#10 DONE 1.8s

#11 resolving provenance for metadata file
#11 DONE 0.0s
[+] Running 4/4
  ✓ compose-lab-web          Built
  ✓ Container compose-lab-db-1  Running
  ✓ Container compose-lab-redis-1  Running
  ✓ Container compose-lab-web-1  Started
PS C:\Users\kmit\Downloads\compose-lab>

PS C:\Users\kmit\Downloads\compose-lab> docker compose ps
>>
NAME           IMAGE          COMMAND                  SERVICE    CREATED        STATUS        PORTS
compose-lab-db-1  postgres:15  "docker-entrypoint.s..."  db         6 minutes ago  Up 6 minutes  0.0.0.0:9090->5432/tcp, [::]:9090->5432
/tcp
compose-lab-redis-1  redis:alpine  "docker-entrypoint.s..."  redis      12 minutes ago  Up 12 minutes  6379/tcp
compose-lab-web-1   compose-lab-web  "python app.py"      web       56 seconds ago  Up 54 seconds  0.0.0.0:5000->5000/tcp, [::]:5000->5000
/tcp
PS C:\Users\kmit\Downloads\compose-lab>

```

localhost:5000

Hello from Flask + Docker!

```

PROBLEMS    OUTPUT    DEBUG CONSOLE    TERMINAL    PORTS
docker + ×  ⌂ ... | ⌂ ×

✓ compose-lab-web          Built
✓ Container compose-lab-redis-1  Running
✓ Container compose-lab-db-1  Running
✓ Container compose-lab-web-1  Recreated
Attaching to db-1, redis-1, web-1
web-1  | * Serving Flask app 'app'
web-1  | * Serving Flask app 'app'
web-1  | * Debug mode: off
web-1  | WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
web-1  | * Running on all addresses (0.0.0.0)
web-1  | * Running on http://127.0.0.1:5000
web-1  | * Running on http://172.23.0.4:5000
web-1  | Press CTRL+C to quit
web-1  | * Debug mode: off
web-1  | WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
web-1  | * Running on all addresses (0.0.0.0)
web-1  | * Running on http://127.0.0.1:5000
web-1  | * Running on http://172.23.0.4:5000
web-1  | Press CTRL+C to quit
  | 
View in Docker Desktop  View Config  Enable Watch

```

← → ⌂ ⌂ localhost:5000

Hello Docker Compose!

LAB ACTION PLAN FOR WEEK 7

Objective:

To enable students to:

- Understand the structure and lifecycle of a Maven project.
- Build and package Java and Web applications using Maven.
- Add dependencies using **pom.xml**, compile and test using plugins.
- Resolve errors and conflicts arising from dependency mismatches.
- Work with parent and multi-module Maven projects.
- Generate executable JARs and deployable WARs using Maven.

Students must **document observations**, include **screenshots of executions**, and answer **scenario-based questions** after completing the tasks.

1. Understanding Maven Project Structure

Maven standardizes the project structure for both Java and web-based applications. `src/`

```
└── main/
    ├── java/      → Application source code
    └── resources/ → Configuration files like config.properties
└── test/
    ├── java/      → Unit test source code
    └── resources/ → Test resources
```

The compiled files and reports are generated in the `target/` directory after a successful build.

2. Steps to Perform Maven Build and Testing

----mvn clean install

- **clean**: Deletes the previous build (`target/` folder)
- **install**: Builds, tests, and installs the package to local `.m2` repository

After execution:

- `target/` folder contains compiled `.class` files and the final `.jar` or `.war`

- Artifact is stored in:
`~/.m2/repository/groupId/artifactId/version/`

Add a Dependency (e.g., Gson):

In `pom.xml`:

```
<dependency>
  <groupId>com.google.code.gson</groupId>
  <artifactId>gson</artifactId>
  <version>2.10</version>
</dependency>
```

After running:

```
mvn clean install
```

Check:

- Gson JAR is downloaded to `.m2/repository/com/google/code/gson/gson/`
- If version is wrong or not found → **BUILD FAILURE** with dependency resolution error.

Creation of Maven Java Project

Step 1. Open Eclipse IDE

 └— 1.1. Launch Eclipse workspace

Step 2. Install Maven Plugin (if not installed)

 └— 2.1. Go to "Help" in the top menu

 └— 2.1.1. Click "Eclipse Marketplace"

 └— 2.1.2. Search for "Maven Integration for Eclipse"

 └— 2.1.3. Install the plugin if not already installed

Step 3. Create a New Maven Project

 └— 3.1. File -> New -> Project...

 └— 3.1.1. Expand "Maven"

└— 3.1.2. Select "Maven Project" and click "Next"

Step 4. Set Project Configuration

└— 4.1. Select workspace location (default or custom)

└— 4.2. Click "Next"

Step 5. Choose Maven Archetype

└— 5.1. Select an archetype(e.g "org.apache.maven.archetypes -> maven-archetype-quickstart 1.4 ")

└— 5.2. Click "Next"

Step 6. Define Project Metadata

└— 6.1. Group ID: (e.g., com.example)

└— 6.2. Artifact ID: (e.g., my-maven-project)

└— 6.3. Version: (default is usually fine)

└— 6.4. Click "Finish"

In Console, artifacts are grouped. When prompted with Y/N, type 'Y'.

Step 7. Maven Project Created

└— 7.1. Project structure is generated with a standard Maven layout

└— 7.2. Includes:

 └— src/main/java (for Java source code)

 └— src/test/java (for test code)

 └— pom.xml (Maven configuration file)

Step 8. Update Project Settings (if needed)

└— 8.1. Right-click on the project -> Maven -> Update Project...

└— 8.2. Ensure dependencies are up to date

Step 9. Build and Run Maven Project

└— 9.1. Right-click on App.java -> Run As -> Maven Clean

 └— 9.1.1. Right-click on App.java -> Run As -> Maven Install

└─ 9.1.2. Right-click on App.java -> Run As -> Maven Test

└─ 9.1.3. Right-click on App.java -> Run As -> Maven Build

Step 10. In the Maven Build dialog:

└─ Enter Goals: clean install test

└─ Click on Apply -> Click on Run

Step 11. Check console for BUILD SUCCESS message.

Step 12. Run the application:

└─ Right-click on App.java -> Run As -> Java Application

└─ Output: "Hello World" displayed.

Creation of Maven web Java Project

Step 1: Open Eclipse

└─ 1.1 Launch Eclipse IDE.

└─ 1.2 Select or create a workspace.

Step 2: Create a New Maven Project

└─ 2.1. File -> New -> Project...

└─ 2.1.1. Expand "Maven"

└─ 2.1.2. Select "Maven Project" and click "Next"

Step 3: Choose Maven Archetype

└─ 3.1. Select an archetype(e.g "'org.apache.maven.archetypes' -> 'maven-archetype-webapp' 1.4 ")

└─ 3.2. Click "Next"

Step 4: Configure the Maven Project

└─ 4.1 Group Id: Enter a group ID (e.g., com.example).

└─ 4.2 Artifact Id: Enter an artifact ID (e.g., my-web-app).

└─ 4.3 Click **Finish** to create the project.

Step 5: Add Maven Dependencies

└— 5.1 Open the **pom.xml** file in the Maven project.

└— 5.2 Add the necessary dependencies for your web project (e.g., Servlet, JSP):

Go to browser -> Open mvnrepository.com

Search for 'Java Servlet API' -> Select the latest version.

Copy the dependency code -> Paste it in MavenWeb's pom.xml under the target folder

└— Example:

```
```xml
<dependency>
 <groupId>javax.servlet</groupId>
 <artifactId>javax.servlet-api</artifactId>
 <version>4.0.1</version>
 <scope>provided</scope>
</dependency>
...
```

```

Step 6:- Configure server:

└— Window -> Show View -> Servers

└— Add server -> Select Tomcat v9.0 server -> Click Next

└— Configure server options (e.g., ports, server location).

Step 7:- Modify 'tomcat-users.xml':

└— Add role and user details under <tomcat-users> tag.

Step 8:- Build the project:

└— Right-click on index.jsp -> Run As -> Maven Clean

└— Right-click on index.jsp -> Run As -> Maven Install

└— Right-click on index.jsp -> Run As -> Maven Test

└— Right-click on index.jsp -> Run As -> Maven Build

Step 9. In the Maven Build dialog:

└─ Enter Goals: clean install test

└─ Click on Apply -> Click on Run

Step 10. Check console for BUILD SUCCESS message.

Step 11. Run the application:

└─ Right-click on index.jsp -> Run As -> Run on Server

└─ Select the Tomcat server -> Click on Finish

Step 12. Output: "Hello World" webpage displayed.

Note:-Now push yours Maven java project and Maven Web Project into your github

3. Configure Java Version via Compiler Plugin

In `pom.xml`:

```
<plugin>
<groupId>org.apache.maven.plugins</groupId>
<artifactId>maven-compiler-plugin</artifactId>
<version>3.11.0</version>
<configuration>
<source>17</source>
<target>17</target>
</configuration> </plugin>
```

Check:

- Run `mvn package`
- Inspect generated JAR: `jar tf target/myapp.jar`

4. JUnit Testing and Reports

Place test files in `src/test/java`. Example:

```

public class AppTest {

    @Test

    public void testSum() {

        assertEquals(5, 2 + 3);

    }
}

```

Run:

`mvn test`

Check:

- target/test-classes/ → compiled test .class files
- target/surefire-reports/ → .txt or .xml test results
If test fails → corresponding log and failure trace shown

5. Handling Errors in pom.xml

- Typos in version numbers or missing repositories cause BUILD FAILURE
- Maven shows precise error in console
- Fix the dependency tag → re-run `mvn clean install`

6. Adding Resource Files

Put config.properties inside:

`src/main/resources/config.properties`

To read:

```
InputStream input =
getClass().getClassLoader().getResourceAsStream("config.properties");
```

After build, check target/classes/ → file should exist there.

7. Multi-Module and Parent Projects

Structure:

```

parent/
└── pom.xml (packaging: pom)
    ├── core/
    │   └── pom.xml
    └── web/
        └── pom.xml

```

- Parent pom.xml defines <modules> and common dependencies
- Each submodule builds independently into its target/ directory

8. Executable JARs

Add to pom.xml:

```
<build>
  <plugins>
    <plugin>
      <groupId>org.apache.maven.plugins</groupId>
      <artifactId>maven-jar-plugin</artifactId>
      <configuration>
        <archive>
          <manifest>
            <mainClass>com.example.Main</mainClass>
          </manifest>
        </archive>
      </configuration>
    </plugin>
  </plugins>
</build>
```

Run:

```
mvn package
java -jar target/myapp.jar
```

9. Building a WAR File

Create a Maven web project with structure:

```
src/main/webapp/
  └── WEB-INF/web.xml
```

Add:

```
<packaging>war</packaging>
```

Command:

```
mvn package
```

Generates target/mywebapp.war → deploy on Tomcat server.

10. Scenario-Based Questions:

1. If my error is about:

*Failed to execute goal org.apache.maven.plugins:maven-compiler-plugin:3.8.1:compile
[ERROR] -> [Help 1]
[ERROR] To see the full stack trace of the errors, re-run Maven with the -X switch.*

How can I resolve it using maven terminal?

ANS: mvn clean compile -X

2. A dependency you added is not recognized by the compiler. What steps would you take to confirm it is available in .m2 and listed in dependency tree?

ANS: Check if dependency JAR exists in local repo:

`~/.m2/repository/<groupId-path>/<artifactId>/<version>/`

3. A teammate sends a .patch file for a bug fix. How would you apply it and include it in your Maven build?

Apply the patch in your project root using git

`git apply path/to/patchfile.patch`

You have multiple JUnit test failures and want to rerun only failed tests.

How would you approach this?

ANS: Use Maven's -Dtest

4. `mvn test -Dsurefire.runOrder=failed`

5. Your Maven build fails due to “Unsupported class version error.” What plugin and configuration would you review?

ANS: Check the maven-compiler-plugin Java version settings in pom.xml

6. You need to change your Java application from a WAR to a standalone JAR. What pom.xml changes are needed?

ANS: Modify pom.xml packaging:

`<packaging>jar</packaging>`

7. You are required to change the default build output directory from target/ to build_output/. How would you configure it?

ANS: In pom.xml:

```
<build>
  <directory>build_output</directory>
</build>
```

- 8. You want to skip tests during the Maven build. What command would you use?**

ANS: mvn clean install -DskipTests

- 9. How would you generate a site report (with test coverage, dependency analysis) for a Maven project?**

ANS: mvn site

- 10. How do you build a Java project using Maven, and what files are generated in the target/ folder after running mvn clean install?**

ANS: mvn clean install

- 11. How does Maven resolve dependency conflicts when two libraries use different versions of the same dependency, and how can you view and manage the dependency tree?**

Maven uses nearest-wins strategy.

`mvn dependency:tree`

Use `<dependencyManagement>` in parent pom.xml to control versions explicitly.

- 12. How do you write and run a JUnit test in a Maven project, and where are the compiled test classes and reports stored after running mvn test?**

ANS: Write test in src/test/java with proper JUnit annotations.

Run tests with:

`mvn test`

Compiled tests: target/test-classes/

- 13. How can you create an executable JAR with a main method using Maven, and which plugin helps configure this behavior?**

ANS: Use the maven-jar-plugin or maven-shade-plugin.

- 14. How do you install and use a custom third-party JAR file in your Maven project, and how can you confirm it's included in the build and classpath?**

ANS: mvn install:install-file -Dfile=path/to/jar -DgroupId=com.custom - DartifactId=custom-lib -Dversion=1.0 -Dpackaging=jar

- 15. How do you create a Maven web project that packages into a WAR file, and what is the standard folder structure for such a project?**

ANS: In pom.xml

<packaging>war</packaging>

16. What command do you use to build a WAR file in Maven, where is it generated, and how can you deploy it to a server like Apache Tomcat?

Build WAR:

mvn clean package

WAR generated at:

target/yourproject.war

Deploy by copying WAR to Tomcat's webapps/ folder or via Tomcat Manager.

17. How do you add JSTL and servlet-api dependencies in a Maven web project, and why should the servlet API use provided scope instead of compile?

Add JSTL and servlet-api dependencies; why servlet-api is

<dependency>

<groupId>javax.servlet</groupId>

<artifactId>javax.servlet-api</artifactId>

<version>4.0.1</version>

<scope>provided</scope>

</dependency>

<dependency>

<groupId>javax.servlet.jsp.jstl</groupId>

<artifactId>javax.servlet.jsp.jstl-api</artifactId>

<version>1.2.1</version>

</dependency>

18. How do you set up a multi-module Maven web project with separate modules for core logic and web interface, and how are these modules built and connected?

ANS: Parent pom.xml with <modules> listing core and web modules.

Core module: logic, packaged as jar.

Web module: depends on core, packaged as war.

mvn clean install

19. How do you configure a Maven web project, and how does its packaging and execution differ from a traditional WAR-based application?

ans: Uses packaging WAR.

src/main/webapp folder contains web resources.

Built WAR deployed to servlet container.

Can use plugins like tomcat7-maven-plugin for execution.

Conclusion

Mastering Maven empowers students to structure projects efficiently, manage complex dependencies, and automate testing and packaging. By practicing real-world scenarios—such as resolving build errors, handling resource files, applying patches, and deploying to servers—students gain valuable hands-on experience aligned with professional software development workflows. Maven not only streamlines builds but also enforces standardization and reusability across projects. Through Maven, students learn efficient project management, dependency control, multi-module integration, and reproducible builds. Understanding Maven's lifecycle, plugin system, and error handling prepares students for professional DevOps and CI/CD workflows. This lab equips students with hands-on experience in packaging, testing, resolving conflicts, and deploying real-world Java and Web applications.

Multi-Model exercise:

New Maven Project

New Maven project

Select project name and location

Create a simple project (skip archetype selection)

Use default Workspace location

Location:

Add project(s) to working set

Working set:

Advanced

New Maven Project

New Maven project

Configure project

Artifact

Group Id: KMIT-WEEK-7

Artifact Id: MultiModule_exercise

Version: 0.0.1-SNAPSHOT

Packaging: pom

Name:

Description:

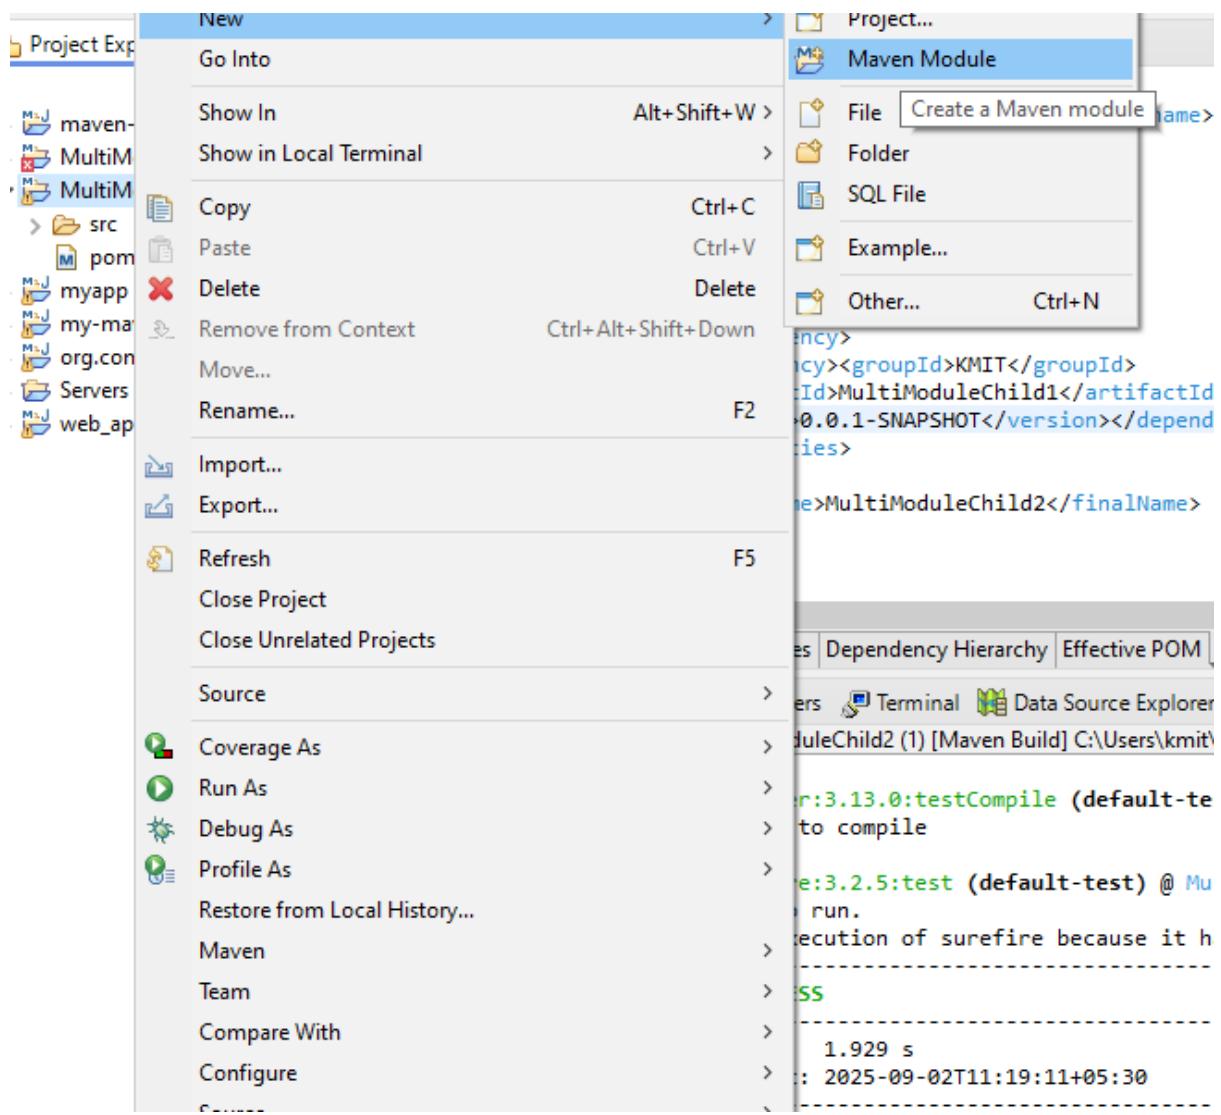
Parent Project

Group Id:

Artifact Id:

Version:

Advanced





New Maven Module

Select the module name and parent



Create a simple project (skip archetype selection)

Module Name:

Parent Project:

Add project(s) to working set

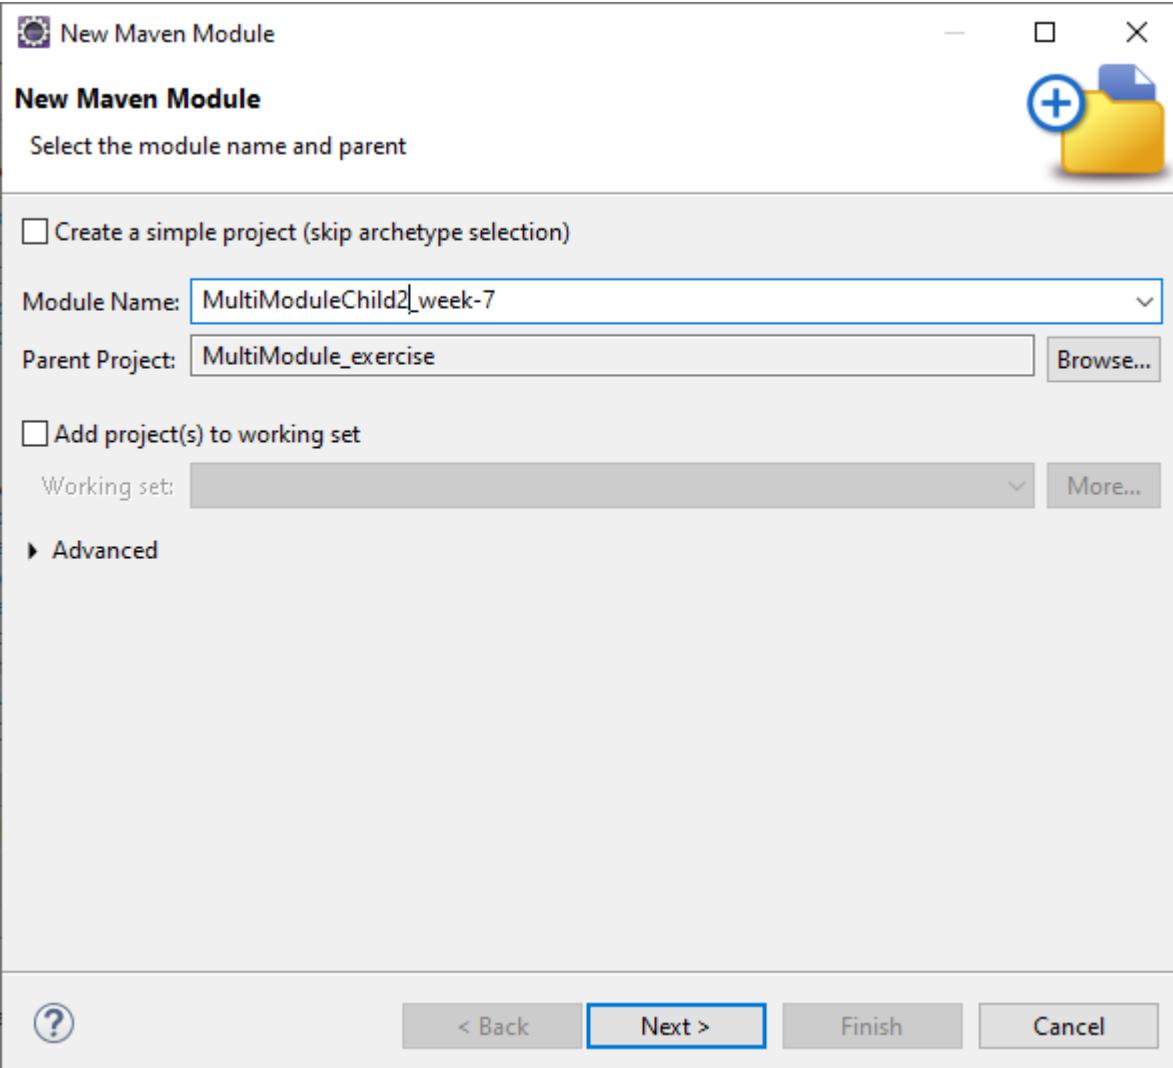
Working set:

► Advanced



< Back

Next >



New Maven Module

New Maven Module

Select an Archetype

Catalog: All Catalogs | Filter: maven-archetype-webapp

Group Id	Artifact Id	Version
com.haoxuer.maven.archetype	maven-archetype-webapp	1.01
com.lodsve	lodsve-maven-archetype-webapp	1.0.2-RELEASE
org.apache.maven.archetypes	maven-archetype-webapp	1.5
org.bytesizebook.com.guide.boot	maven-archetype-webapp	1.0

An archetype to create the starting code for the first three chapters of Guide to Web Development with Java, 2nd edition.
<https://repo1.maven.org/maven2>

Show the last version of Archetype only Include snapshot archetypes [Add Archetype...](#)

[Advanced](#)

[?](#) < Back [Next >](#) [Finish](#) [Cancel](#)

New Maven Module

New Maven Module

Specify Archetype parameters

Group Id: KMIT-WEEK-7
 Artifact Id: MultiModelChild2example
 Version: 0.0.1-SNAPSHOT
 Package: org.MultiModelChild2example
 run archetype generation interactively

Properties available from archetype:

Name	Value

[Add...](#) [Remove](#)

[Advanced](#)

[?](#) < Back [Next >](#) [Finish](#) [Cancel](#)

```
<terminated> C:\Users\kmit\Documents\eclipse\plugins\org.eclipse.jdt.core\hotspot\jre\full\win32\x86_64_21.0.3.v20240426-1530\jre\bin\javaw.exe (02-Sept-2025, 12:26:42)
[INFO] Parameter: package, Value: org.MultiModelChild2example
[INFO] Parameter: groupId, Value: KMIT-WEEK-7
[INFO] Parameter: artifactId, Value: MultiModelChild2example
[INFO] Parameter: version, Value: 0.0.1-SNAPSHOT
[WARNING] CP Don't override file C:\Users\kmit\workspace\MultiModule_exercise\MultiModelChild2example\src\main\webapp
[WARNING] CP Don't override file C:\Users\kmit\workspace\MultiModule_exercise\MultiModelChild2example\.mvn
[INFO] Project created from Archetype in dir: C:\Users\kmit\workspace\MultiModule_exercise\MultiModelChild2example
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 40.094 s
[INFO] Finished at: 2025-09-02T12:27:28+05:30
[INFO] -----
```

Name: MultiModule_exercise (1)

Main JRE Refresh Source Environment Common

Base directory: \${project_loc:MultiModule_exercise}

Goals: clean install test

Profiles:

User settings: C:\Users\kmit\.m2\settings.xml

Offline Update Snapshots
 Debug Output Skip Tests Non-recursive
 Resolve Workspace artifacts

Threads: 1 Color Output: Auto

Parameter Name	Value

Add... Edit... Remove

Revert Apply

Run Close

```
[...] Skipping execution of surefire because it has already been run for this configuration
[INFO] -----
[INFO] Reactor Summary for MultiModule_exercise 0.0.1-SNAPSHOT:
[INFO]
[INFO] MultiModule_exercise ..... SUCCESS [ 0.498 s]
[INFO] MultiModuleChild1_week-7 ..... SUCCESS [ 4.735 s]
[INFO] MultiModelChild2example Maven Webapp ..... SUCCESS [ 3.408 s]
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 8.800 s
[INFO] Finished at: 2025-09-02T12:29:30+05:30
[INFO] -----
```

```
[...] Copying 0 resource from src\test\resources to target\test-classes
[INFO]
[INFO] --- compiler:3.13.0:testCompile (default-testCompile) @ MultiModuleChild1_week-7 ---
[INFO] Nothing to compile - all classes are up to date.
[INFO]
[INFO] --- surefire:3.2.5:test (default-test) @ MultiModuleChild1_week-7 ---
[INFO] Skipping execution of surefire because it has already been run for this configuration
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 1.807 s
[INFO] Finished at: 2025-09-02T12:29:59+05:30
[INFO] -----
```

Name: MultiModelChild2example

Main JRE Refresh Source Environment Common

Base directory: \${project_loc:MultiModelChild2example}

Goals: clean install test

Profiles:

User settings: C:\Users\kmit\.m2\settings.xml

Threads: 1 Color Output: Auto

Parameter Name Value Add... Edit... Remove

Revert Apply

Run Close

```
[INFO] --- compiler:3.13.0:testCompile (default-testCompile) @ MultiModelChild2example ---
[INFO] No sources to compile
[INFO]
[INFO] --- surefire:3.3.0:test (default-test) @ MultiModelChild2example ---
[INFO] No tests to run.
[INFO] Skipping execution of surefire because it has already been run for this configuration
[INFO]
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 2.088 s
[INFO] Finished at: 2025-09-02T12:30:30+05:30
[INFO] -----
```

SBQs(week 8)

1. Q: What is Jenkins primarily used for?

A: Jenkins is used for Continuous Integration and Continuous Delivery (CI/CD) automation.

2. Q: What is a feature of Jenkins?

A: Automates building, testing, and deploying code using pipelines and plugins.

3. Q: What is the default port on which Jenkins runs?

A: Port **8080**.

4. Q: What can be integrated with Jenkins for version control?

A: **Git, GitHub, Bitbucket, GitLab, and SVN.**

5. Q: What is the purpose of Jenkins plugins?

A: To extend Jenkins functionality and integrate external tools.

6. Q: Which type of Jenkins job is best for one-off tasks or small scripts?

A: Freestyle Project.

7. Q: How can you manage sensitive information such as API keys in Jenkins?

A: Use Credentials Manager to store and access them securely.

8. Q: What does the "Blue Ocean" feature in Jenkins refer to?

A: A modern, user-friendly UI for visualizing and managing pipelines.

9. Q: What does the "Blue Ocean" feature in Jenkins refer to?

A: It provides a graphical view of pipeline stages and builds.

10. Q: Which Jenkins component allows distributed builds across multiple machines?

A: Jenkins agents (nodes)

11. Q: List five important Jenkins plugins for microservices CI/CD and their purposes.

A: Git Plugin – Integrates Git repositories.

Pipeline Plugin – Enables CI/CD pipeline scripting.

Docker Plugin – Builds and deploys Docker containers.

Kubernetes Plugin – Runs builds in Kubernetes pods.

JUnit Plugin – Publishes and visualizes test results.

12. Q: Steps to install a plugin via Jenkins UI and key considerations?

A: Go to Manage Jenkins → Manage Plugins → Available.

Search and install the plugin.

Restart Jenkins.

Check compatibility with Jenkins version and dependencies before installing.

13. **Q:** Steps to install a plugin via Jenkins UI and key considerations?

A: Navigate to Manage Jenkins → Manage Plugins → Available tab.

Install selected plugin and restart Jenkins.

Verify update logs and ensure version compatibility.

14. **Q:** After installing the Git Plugin, how to configure it?

A: Go to Manage Jenkins → Global Tool Configuration.

Add Git installation path.

In pipeline, use:

```
git url: 'https://github.com/user/repo.git', credentialsId: 'git-creds'.
```

15. **Q:** Common plugin issues and how to troubleshoot them?

A: Issues: Dependency conflicts, version mismatches, plugin crashes.

Fixes: Check Jenkins logs, update plugins, verify compatibility matrix, or rollback versions.

Week 9: Pipeline Creation using script

1. Evaluation of Jenkins pipeline.
2. **WORKING ON BUILD TRIGGERS FOR LAST JENKINS PIPILINE**
3. Hands-on practice on creation of scripted Jenkins pipeline.
4. Take the screenshots for above task

Procedure

General :

Description :- provide the description of the project

Build triggers: here we can provide with build triggers of your choice

Advance project option:

Definition:

Choose pipeline script

Here code for pipeline -script

Copy the code there

```
pipeline {  
    agent any  
    tools{  
        maven 'MAVEN-HOME'  
    }  
    stages {  
        stage('git repo & clean') {  
            steps {  
                //bat "rmdir /s /q mavenjava"  
                bat "git clone provide your github link"  
                bat "mvn clean -f mavenjava"  
            }  
        }  
        stage('install') {  
            steps {  
                bat "mvn install -f mavenjava" #project name#  
            }  
        }  
        stage('test') {  
            steps {  
                bat "mvn test -f mavenjava"  
            }  
        }  
    }  
}
```

```
        }
    }
stage('package') {
    steps {
        bat "mvn package -f mavenjava"
    }
}
}
```

Apply and save

Jenkins

Enabled

Description
Scripted Jenkins pipeline for Maven build

Plain text [Preview](#)

Discard old builds [?](#)
 Do not allow concurrent builds
 Do not allow the pipeline to resume if the controller restarts
 GitHub project
 Permission to Copy Artifact
 Pipeline speed/durability override [?](#)
 Preserve stashes from completed builds [?](#)
 This project is parameterized [?](#)
 Throttle builds [?](#)

Build Triggers

Build after other projects are built [?](#)
 Build periodically [?](#)
 Build whenever a SNAPSHOT dependency is built [?](#)
 GitHub hook trigger for GITScm polling [?](#)
 Poll SCM [?](#)
 Quiet period [?](#)
 Trigger builds remotely (e.g., from scripts) [?](#)

Advanced Project Options

Advanced [▼](#)

Pipeline

Dashboard > All > Scripted Jenkins > Configuration

Configure

Script

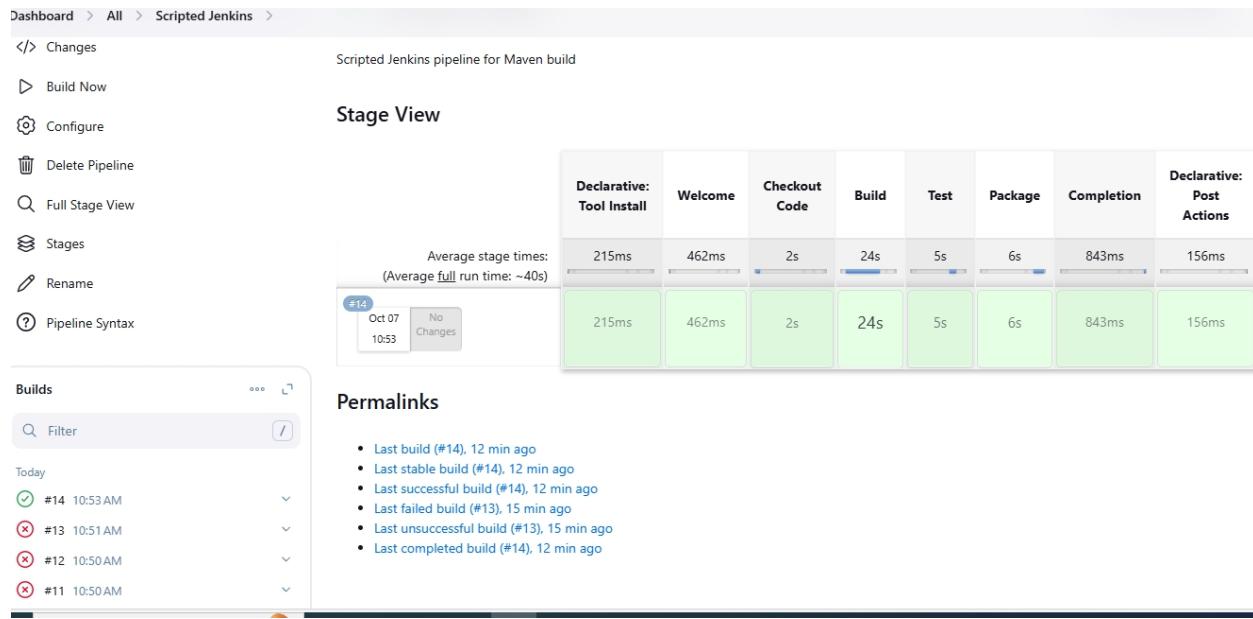
```
1 pipeline {  
2     agent any  
3     tools {  
4         maven "MAVEN_HOME" // Make sure this matches your Jenkins Maven installation name  
5     }  
6     stages {  
7         stage("Welcome") {  
8             steps {  
9                 echo "Starting Jenkins build for $JOB_NAME"  
10                echo "Build Number: $env.BUILD_NUMBER"  
11            }  
12        }  
13        stage("Checkout Code") {  
14            steps {  
15                checkout scm  
16            }  
17        }  
18    }  
19}
```

Use Groovy Sandbox [?](#)

[Pipeline Syntax](#)

[Save](#) [Apply](#)

REST API Jenkins 2.489



SBQ's:

1. Your manager asks you to clean the workspace before building — which stage in this pipeline takes care of it?

Ans: The Build stage (inside the dir('selab-internal') block) runs mvn clean install, and additionally, the Checkout Code stage removes any existing folder with rmdir /s /q selab-internal, which ensures the workspace is cleaned before the build.

2. The GitHub repository link is missing in the script — where exactly do you provide it?

Ans: The GitHub link is provided in the Checkout Code stage in the line:

```
git clone https://github.com/NaveenCK-10/selab-internal.git
```

3. If Maven is not configured globally in Jenkins, which section of the pipeline will fail first?

Ans: The pipeline will fail at the Build stage where Maven is invoked using:

```
mvn clean install
```

because Jenkins won't be able to locate Maven without a global tool configuration.

4. A teammate complains the pipeline is not creating .war files — which stage is responsible?

Ans: The Package stage is responsible for creating the final .jar or .war artifacts using:

```
mvn package
```

5. Your test cases failed, but the pipeline still continued — how will Jenkins behave in the test stage?

Ans: By default, if a Maven test fails, it throws an error and the pipeline stops. In our current script, the pipeline stops at the Test stage on failure, and subsequent stages (Package) are skipped.

6. Instead of running nightly builds, you want this pipeline to trigger only when GitHub changes occur — where will you configure it?

Ans: You configure it in Build Triggers → GitHub hook trigger for GITScm polling in Jenkins. Alternatively, a GitHub webhook can notify Jenkins to trigger the pipeline on every push.

7. If you replace mvn clean with mvn compile, what difference will it make to the project build?

Ans: mvn clean deletes previous build artifacts to avoid conflicts, while mvn compile only compiles the source code without cleaning previous outputs. Replacing clean with compile may lead to using old artifacts in the build.

8. The project folder name is not mavenjava but studentapp — which parts of the script must you edit?

Ans: Update all references of dir('selab-internal') to dir('studentapp') and also the rmdir /s /q and git clone folder name accordingly.

9. If the install stage fails, will the test and package stages still run in this pipeline?

Ans: No. In a declarative Jenkins pipeline, if a stage fails (here Build or Install), the following stages (Test and Package) are skipped automatically.

10. A student asks where to add deployment steps to Tomcat after packaging — which is the best place in this script?

Ans: Deployment steps should be added after the Package stage, ideally in a new stage called Deploy that runs commands like copying .war to Tomcat's webapps directory.

11. Why is tools { maven 'MAVEN-HOME' } used in this pipeline?

Ans: It tells Jenkins which Maven installation to use so that Maven commands like mvn clean install can run properly.

12. If GitHub credentials are private, how can you secure them in the git repo & clean stage?

Ans: We can store GitHub credentials in Jenkins Credentials and reference them using credentialsId in the git command:

```
git branch: 'main', credentialsId: 'github-credentials-id', url: 'https://github.com/...'
```

13. Which Jenkins plugin is needed to recognize the pipeline { ... } structure in this script?

Ans: We need the Pipeline plugin (also called Workflow plugin) to understand the pipeline { ... } structure.

14. In Windows, the script uses bat commands — what change would you make if Jenkins runs on Linux?

Ans: If Jenkins runs on Linux, we should replace all `bat` commands with `sh` commands.

For example: `'mvn clean install'`

15. How will you modify the pipeline to stop execution if the GitHub clone command fails?

Your project lead asks you to print a welcome message in Jenkins to confirm the pipeline is working — how would you script it?

Ans: We can add `|| exit 1` at the end of the git clone command so Jenkins stops immediately if cloning fails. Or we can use checkout step which fails the stage automatically on error.

16. A teammate forgot to pull the latest GitHub code before building; how can you fix this in your pipeline?

Ans: We can add an initial stage like this:

```
stage('Welcome') {  
    steps {  
        echo "Hello! Jenkins pipeline is running for selab-internal"  
    }  
}
```

It shows a friendly message in the console when the pipeline starts.

17. Your Java file Hello.java must be compiled every time code is committed — how will you add this step?

Ans: We can include `mvn compile` or `mvn install` in the Build stage. Since Maven automatically compiles Java files, every commit will trigger compilation when the pipeline runs.

18. Tests should stop the pipeline if they fail — where will you put the mvn test command?

Ans: Keep `mvn test` in the Test stage. By default, if any test fails, Maven returns an error, and Jenkins stops the pipeline automatically.

19. Every evening at 6 PM your pipeline should run automatically — how can you set this in Jenkins?

Ans: We can add a cron schedule in Jenkins Build Triggers. For example:

```
0 18 * * *
```

This runs the pipeline at 6:00 PM daily.

20. You only want to package the project if compilation succeeds — how would you connect the stages?

Ans: Jenkins declarative pipelines automatically stop stages on failure. So if Build (compile/install) fails, the Package stage will not run. This ensures packaging happens only after successful compilation.

21. Your professor wants a .jar file generated for submission — what pipeline stage will you add?

Ans: The .jar file is created in the Package stage using:

```
mvn package
```

It will be located in selab-internal/target/ after a successful build

22. A Git clone step is failing due to wrong credentials — how would you secure and use them in your script?

Ans: We should store GitHub credentials in Jenkins Credentials and use them in the pipeline:

```
git branch: 'main', credentialsId: 'github-credentials-id', url: 'https://github.com/NaveenCK-10/selab-internal.git'
```

This keeps credentials secure and allows Jenkins to clone private repos.

23. A teammate wants to see the build number inside console logs — how do you print it in the pipeline?

Ans: We can print the build number using the environment variable \${env.BUILD_NUMBER}. For example:

```
echo "Build Number: ${env.BUILD_NUMBER}"
```

It will appear in the console every time the pipeline runs.

LAB ACTION PLAN FOR WEEK 10

Working with minikube and Nagios

1. Hands-on practice of creating, running and scaling pods in minikube.
2. Running Nginx server on specified port number by explaining the Nginx monitoring tool
3. Running Nagios server and Understanding the Monitoring tool using Docker.
4. AWS-free Trier account Creation steps
5. Upload the screenshots for the tasks

Kubernetes:

Kubernetes is a tool that automates how we run and manage applications inside the container.

Dockers will only run containers, if in any case the container fails/stopped/killed, the docker will not help us, here is where Kubernetes plays an important role, Kubernetes cluster will be responsible in creating a new container and managing various containers.

POD: In Kubernetes, a Pod is the smallest deployable unit that you can create and manage.

Minikube:

Minikube creates a VM on your local machine and deploys a simple Kubernetes cluster with one node. It's a lightweight implementation. Minikube *is* a version of Kubernetes.

Nagios:

Nagios is an **open-source IT infrastructure monitoring tool**. It monitors

- Servers
- Network devices
- Applications and services

It **alerts administrators** when issues occur and notifies when they are resolved.

Step 1: Install Prerequisites

Before installing Minikube, ensure the following are installed:

1. **Virtualization Support:**
 - Verify virtualization is enabled:

2. Hypervisor:

- Minikube supports multiple hypervisors (e.g., **Hyper-V**, **VirtualBox**, or **Docker** as a driver).
- Install one of the following:
 - **Hyper-V** (pre-installed on Windows 10/11 Pro or Enterprise).
 - **Docker Desktop** (if you want to use Docker as the driver).

```
MICROSOFT WINDOWS [version 10.0.19045.0332]
(c) Microsoft Corporation. All rights reserved.

C:\Users\kmit>system info
'system' is not recognized as an internal or external command,
operable program or batch file.

C:\Users\kmit>systeminfo

Host Name:           DESKTOP-J277LFB
OS Name:            Microsoft Windows 10 Pro
OS Version:         10.0.19045 N/A Build 19045
OS Manufacturer:   Microsoft Corporation
OS Configuration:  Standalone Workstation
OS Build Type:     Multiprocessor Free
Registered Owner:  kmit
Registered Organization:
Product ID:        00331-10000-00001-AA061
Original Install Date: 17-12-2024, 04:01:40
System Boot Time:  14-10-2025, 09:48:50
System Manufacturer: LENOVO
System Model:      10MQS20500
System Type:       x64-based PC
Processor(s):      1 Processor(s) Installed.
                    [01]: Intel64 Family 6 Model 158 Stepping 9 GenuineIntel ~2712 M
BIOS Version:      LENOVO M1AKT3EA, 21-01-2019
Windows Directory: C:\Windows
System Directory:  C:\Windows\system32
Boot Device:       \Device\HarddiskVolume1
System Locale:    en-us;English (United States)
Input Locale:     00004009
Time Zone:        (UTC+05:30) Chennai, Kolkata, Mumbai, New Delhi
Total Physical Memory: 8,080 MB
Available Physical Memory: 1,233 MB
Virtual Memory: Max Size: 11,345 MB
Virtual Memory: Available: 2,227 MB
Virtual Memory: In Use: 9,118 MB
Page File Location(s): C:\pagefile.sys
Domain:           WORKGROUP
Logon Server:     \\DESKTOP-J277LFB
Hotfix(s):        17 Hotfix(s) Installed.
                    [01]: KB5064400
                    [02]: KB5045936
                    [03]: KB5011049
```

Step 2: Download Minikube

1. Open a PowerShell or Command Prompt with administrator privileges.
2. Download the latest Minikube executable using this command:
 3. curl -LO
<https://storage.googleapis.com/minikube/releases/latest/minikube-installer.exe>
4. Install Minikube by running the installer:
 5. .\minikube-installer.exe

Step 3: Add Minikube to PATH

If Minikube is not automatically added to your PATH during installation:

1. Open **System Properties → Environment Variables**.
 2. Add the directory where Minikube is installed (e.g., C:\Program Files\Minikube) to your PATH variable.
-

Step 4: Start Minikube

1. Open a terminal (PowerShell or CMD). Do the following commands
2. Start Minikube with a specified driver (e.g., Hyper-V, Docker, or VirtualBox). For example:
 - **Hyper-V:**
○ minikube start --driver=hyperv
 - **Docker:**
○ minikube start --driver=docker
3. Verify Minikube is running:
4. minikube status

```
C:\Users\kmit>
C:\Users\kmit>minikube start --driver=docker
* minikube v1.34.0 on Microsoft Windows 10 Pro 10.0.19045.6332 Build 19045.6332
* minikube 1.37.0 is available! Download it: https://github.com/kubernetes/minikube/releases/tag/v1.37.0
* To disable this notice, run: 'minikube config set WantUpdateNotification false'

* Using the docker driver based on existing profile
* Starting "minikube" primary control-plane node in "minikube" cluster
* Pulling base image v0.0.45 ...
* Restarting existing docker container for "minikube" ...
! Failing to connect to https://registry.k8s.io/ from inside the minikube container
* To pull new external images, you may need to configure a proxy: https://minikube.sigs.k8s.io/docs/reference/networking/proxy/
* Preparing Kubernetes v1.31.0 on Docker 27.2.0 ...
* Verifying Kubernetes components...
  - Using image gcr.io/k8s-minikube/storage-provisioner:v5
* Enabled addons: storage-provisioner, default-storageclass
* Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default

C:\Users\kmit>
C:\Users\kmit>kubectl get nodes
NAME      STATUS    ROLES     AGE   VERSION
minikube  Ready     control-plane  293d  v1.31.0

C:\Users\kmit>
```

Step 5: Interact with Minikube

kubectl is a command-line tool used in Kubernetes to interact with and manage Kubernetes clusters.

Once Minikube is running:

1. Use kubectl to interact with the cluster.

- o Install `kubectl` if not already installed:
 - o `minikube kubectl -- get pods -A`
 - o Or download it separately from the [official Kubernetes site](#).
2. Open the Minikube dashboard (optional):
3. `minikube dashboard`

```
C:\Users\kmit>kubectl get pods -A
NAMESPACE      NAME                               READY   STATUS    RESTARTS   AGE
kube-system    coredns-6f6b679f8f-8wfqf          1/1     Running   1 (5m12s ago) 293d
kube-system    coredns-6f6b679f8f-ptm4v          1/1     Running   1 (5m12s ago) 293d
kube-system    etcd-minikube                      1/1     Running   1 (5m12s ago) 293d
kube-system    kube-apiserver-minikube           1/1     Running   1 (5m12s ago) 293d
kube-system    kube-controller-manager-minikube  1/1     Running   1 (5m12s ago) 293d
kube-system    kube-proxy-wxt5w                  1/1     Running   1 (5m12s ago) 293d
kube-system    kube-scheduler-minikube          1/1     Running   1 (5m12s ago) 293d
kube-system    storage-provisioner              1/1     Running   3 (4m38s ago) 293d

C:\Users\kmit>
```

Optional: Check Your Installation

Run the following to verify the installation:

Minikube version

```
kubectl version --client
C:\Users\kmit>kubectl version --client
Client Version: v1.30.5
Kustomize Version: v5.0.4-0.20230601165947-6ce0bf390ce3

C:\Users\kmit>
```

Troubleshooting

1. **If Minikube fails to start:**
 - o Ensure your hypervisor (Hyper-V/Docker/VirtualBox) is installed and running.
 - o Check the Minikube logs:
 - o `minikube logs`
2. **Updating Minikube:**
 3. `minikube update-check`
 4. `minikube update`

Minikube Automation Steps

Step 1: Start Minikube Cluster

- Open your terminal and run the command:

```
minikube start
C:\Users\kmit>minikube start
* minikube v1.34.0 on Microsoft Windows 10 Pro 10.0.19045.6332 Build 19045.6332
* Using the docker driver based on existing profile
* Starting "minikube" primary control-plane node in "minikube" cluster
* Pulling base image v0.0.45 ...
* Updating the running docker "minikube" container ...
! Failing to connect to https://registry.k8s.io/ from inside the minikube container
* To pull new external images, you may need to configure a proxy: https://minikube.sigs.k8s.io/docs/reference/networking/proxy/
* Preparing Kubernetes v1.31.0 on Docker 27.2.0 ...
* Verifying Kubernetes components...
  - Using image gcr.io/k8s-minikube/storage-provisioner:v5
* Enabled addons: storage-provisioner, default-storageclass
* Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default

C:\Users\kmit>
```

Step 2: Create and Manage Deployment

1. Create an application in Kubernetes:

- Command:

```
kubectl create deployment mynginx --image=nginx
```

```
C:\Users\kmit>kubectl create deployment mynginx --image=nginx
deployment.apps/mynginx created

C:\Users\kmit>
```

if already created then

```
kubectl set image deployment/myngnix nginx=nginx:latest
```

- Verify the deployment using: Kubernetes responds by showing you a list that includes the names of your deployment groups

```
kubectl get deployments
```

```
C:\Users\kmit>kubectl get deployments
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
mynginx  1/1     1           1           57s

C:\Users\kmit>
```

- Ensure `mynginx` appears in the output.

Check the following commands:

- `kubectl get pods`

```
C:\Users\kmit>kubectl get pods
NAME                  READY   STATUS    RESTARTS   AGE
mynginx-79bb8756c7-xm6g5   1/1     Running   0          92s

C:\Users\kmit>
```

- `kubectl describe pods`

```
C:\Users\kmit>kubectl describe pods
Name:           mynginx-79bb8756c7-xm6g5
Namespace:      default
Priority:       0
Service Account: default
Node:           minikube/192.168.49.2
Start Time:     Tue, 14 Oct 2025 10:23:08 +0530
Labels:         app=mynginx
                pod-template-hash=79bb8756c7
Annotations:   <none>
Status:        Running
IP:            10.244.0.8
IPs:
  IP:          10.244.0.8
Controlled By: ReplicaSet/mynginx-79bb8756c7
Containers:
  nginx:
    Container ID:  docker://c8378173f47af0611077a51b228cf13893130cbe4ffa08accd06a08799a7455
    Image:         nginx
    Image ID:     docker-pullable://nginx@sha256:3b7732505933ca591ce4a6d860cb713ad96a3176b82f7979a8dfa9973486a0d6
    Port:          <none>
    Host Port:    <none>
    State:        Running
      Started:   Tue, 14 Oct 2025 10:23:36 +0530
    Ready:        True
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-h6lhm (ro)
Conditions:
  Type          Status
  PodReadyToStartContainers  True
  Initialized   True
  Ready         True
  ContainersReady  True
  PodScheduled  True
Volumes:
  kube-api-access-h6lhm:
    Type:           Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:   kube-root-ca.crt
```

2. Expose Deployment as a Service:

- Command:

```
kubectl expose deployment mynginx --type=NodePort --port=80 --target-port=80  
C:\Users\kmit>kubectl expose deployment mynginx --type=NodePort --port=80 --target-port=80  
service/mynginx exposed  
C:\Users\kmit>
```

Step 3: Scale the Deployment

Command:Scales the Nginx deployment to 4 replicas (pods).

```
kubectl scale deployment mynginx --replicas=4  
C:\Users\kmit>kubectl scale deployment mynginx --replicas=4  
deployment.apps/mynginx scaled  
C:\Users\kmit>
```

```
kubectl get service mynginx
```

```
C:\Users\kmit>kubectl get service mynginx  
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE  
mynginx   NodePort  10.102.91.131  <none>        80:32151/TCP  3m28s  
C:\Users\kmit>
```

Step 4: Access the Nginx App

1. Using Port Forwarding:

- Command:

```
kubectl port-forward svc/mynginx 8081:80
```

```
C:\Users\kmit>kubectl port-forward svc/mynginx 8081:80
Forwarding from 127.0.0.1:8081 -> 80
Forwarding from [::1]:8081 -> 80
```

- Access the app via <http://localhost:8081>.



If Error, use this option, **Using Minikube Tunnel**:

- Start the tunnel:

```
minikube tunnel
```

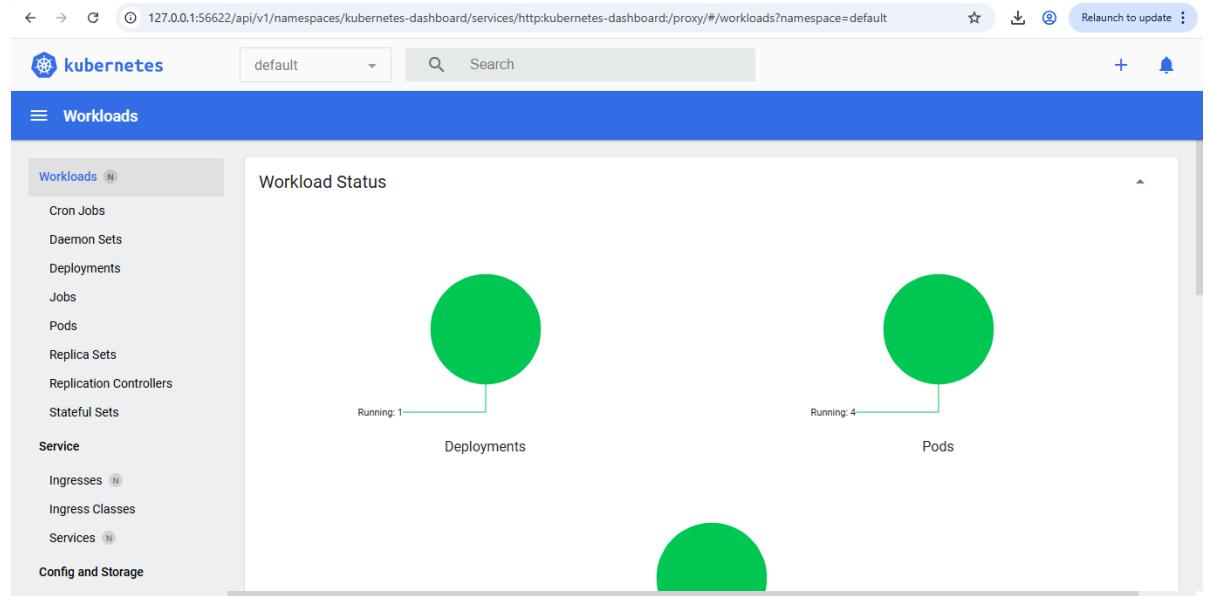
- Retrieve the service URL:

```
minikube service mynginx --url
```

- Open the provided URL in your browser.
- Open the kubernets dashboard

- Open the minikube dashboard

Minikube dashboard



The screenshot shows the Kubernetes Pods dashboard for the default namespace. The sidebar on the left is identical to the previous screenshot. The main area is titled 'Pods' and lists four running pods: 'mynginx-79bb8756c7-gz7bc', 'mynginx-79bb8756c7-kmmrh', 'mynginx-79bb8756c7-stjp', and 'mynginx-79bb8756c7-xm6g5'. Each pod row includes columns for Name, Images, Labels, Node, Status, Restarts, CPU Usage (cores), Memory Usage (bytes), and Created time. All pods are running on the minikube node and were created approximately 13 minutes ago.

Name	Images	Labels	Node	Status	Restarts	CPU Usage (cores)	Memory Usage (bytes)	Created
mynginx-79bb8756c7-gz7bc	nginx	app: mynginx pod-template-hash: 79bb8756c7	minikube	Running	0	-	-	13 minutes ago
mynginx-79bb8756c7-kmmrh	nginx	app: mynginx pod-template-hash: 79bb8756c7	minikube	Running	0	-	-	13 minutes ago
mynginx-79bb8756c7-stjp	nginx	app: mynginx pod-template-hash: 79bb8756c7	minikube	Running	0	-	-	13 minutes ago
mynginx-79bb8756c7-xm6g5	nginx	app: mynginx pod-template-hash: 79bb8756c7	minikube	Running	0	-	-	16 minutes ago

Step 5: Stop and Clean Up

1. Stop Nginx Deployment:

- o Commands:

```
kubectl delete deployment mynginx  
C:\Users\kmit>kubectl delete deployment mynginx  
deployment.apps "mynginx" deleted  
C:\Users\kmit>
```

```
kubectl delete service mynginx
```

```
C:\Users\kmit>kubectl delete service mynginx  
service "mynginx" deleted  
C:\Users\kmit>
```

2. Stop Minikube (Optional):

- o Command:

```
minikube stop
```

3. Delete Minikube Cluster (Optional):

- o Command:
minikube delete

Nagios Automation Steps

Step 1: Pull the Nagios Docker Image

- Open a terminal and run:

```
docker pull jasonrivers/nagios:latest
```

```
C:\Users\kmit>docker pull jasonrivers/nagios:latest
latest: Pulling from jasonrivers/nagios
9ffe54c5c139: Download complete
279b28aefaf10: Download complete
a900dfccceb38: Download complete
9a90645e352c: Download complete
8e911c59da28: Download complete
8c389e58e867: Download complete
3d5785144815: Download complete
b0e280e9aa8c: Download complete
8fb30af17153: Download complete
785b9873bdf4: Download complete
ff65ddf9395b: Download complete
a2fc4187e3b4: Download complete
c219d58cc3f9: Download complete
566cdc02555d: Download complete
0ef9446ba5cc: Download complete
53aff88babbc4: Download complete
d72f92e29533: Download complete
706ed7d4ce0a: Download complete
d3245570f968: Download complete
b69c76bd2b6b: Download complete
e58e184b986a: Download complete
c700be87d617: Download complete
eeb77e6dde3e: Download complete
4f4fb700ef54: Already exists
0bd0f5795eeb: Download complete
71bfb306f8cb: Download complete
738fc7520889: Download complete
fe8a6b2cf4e3: Download complete
e6f8fab512d1: Download complete
15f36d0b0439: Download complete
d5aa2a3a6539: Download complete
Digest: sha256:2a7c2b20d118baf92b47b69a3901e68dd7664617801b94e560bc4d6564d6ae54
Status: Downloaded newer image for jasonrivers/nagios:latest
docker.io/jasonrivers/nagios:latest
```

Step 2: Run Nagios

- Command:

```
docker run --name nagiosdemo -p 8888:80 jasonrivers/nagios:latest
```

```
C:\Users\kmit>docker run --name nagiosdemo -p 8881:80 jasonrivers/nagios:latest
Adding password for user nagiosadmin
chown: warning: '.' should be '::': 'nagios.nagios'
Started runsvdir, PID is 13
checking permissions for nagios & nagiosgraph
rsyslogd: [origin software="rsyslogd" swVersion="8.2312.0" x-pid="25" x-info="https://www.rsyslog.com"] start

Nagios Core 4.5.7
Copyright (c) 2009-present Nagios Core Development Team and Community Contributors
Copyright (c) 1999-2009 Ethan Galstad
Last Modified: 2024-10-24
License: GPL

Website: https://www.nagios.org
Nagios 4.5.7 starting... (PID=26)
Local time is Tue Oct 14 05:23:37 UTC 2025
nagios: Nagios 4.5.7 starting... (PID=26)
nagios: Local time is Tue Oct 14 05:23:37 UTC 2025
nagios: LOG VERSION: 2.0
wproc: Successfully registered manager as @wproc with query handler
nagios: qh: Socket '/opt/nagios/var/rw/nagios.qh' successfully initialized
nagios: qh: core query handler registered
nagios: qh: echo service query handler registered
nagios: qh: help for the query handler registered
nagios: wproc: Successfully registered manager as @wproc with query handler
wproc: Registry request: name=Core Worker 41;pid=41
nagios: wproc: Registry request: name=Core Worker 41;pid=41
wproc: Registry request: name=Core Worker 45;pid=45
nagios: wproc: Registry request: name=Core Worker 45;pid=45
wproc: Registry request: name=Core Worker 43;pid=43
nagios: wproc: Registry request: name=Core Worker 43;pid=43
wproc: Registry request: name=Core Worker 44;pid=44
nagios: wproc: Registry request: name=Core Worker 44;pid=44
wproc: Registry request: name=Core Worker 42;pid=42
nagios: wproc: Registry request: name=Core Worker 42;pid=42
wproc: Registry request: name=Core Worker 46;pid=46
nagios: wproc: Registry request: name=Core Worker 46;pid=46
postfix/master[28]: daemon started -- version 3.8.6, configuration /etc/postfix
Successfully launched command file worker with pid 60
nagios: Successfully launched command file worker with pid 60
-
```

Step 3: Access Nagios Dashboard

- Open your browser and navigate to:

<http://localhost:8888>

The screenshot shows the Nagios Core 4.5.7 dashboard. At the top, a message box通知 "A new version of Nagios Core is available! Visit nagios.org to download Nagios 4.5.9." On the left, a sidebar menu includes sections for General, Current Status (with Hosts, Services, Host Groups, and Problems), and Reports (Availability, Trends, Alerts, History, Summary). The main content area displays a "System Status" dashboard with a grid of host icons, server statistics (Memory, CPU, Disk), a service status grid, and a bar chart for CPU usage. A "Latest Alerts" section shows a single alert from "nagiosadmin" about memory usage. On the right, there's an "Administrative Tools" sidebar with links for initial setup tasks like configuring system settings and hosts.

- **Login Credentials:**
 - Username: nagiosadmin
 - Password: nagios
- Once logged in, explore:
 - Hosts: View systems being monitored.

The screenshot shows the "Host Status Details For All Host Groups" page. The left sidebar is identical to the previous dashboard. The main area displays a table of host status details. The table has columns for Host, Status, Last Check, Duration, and Status Information. One row is shown for "localhost" which is UP and has been checked at 10-14-2025 05:25:45. The "Status Information" column indicates "PING OK - Packet loss = 0%, RTA = 0.09 ms". Above the table, summary boxes show "Host Status Totals" (1 Up, 0 Down, 0 Unreachable, 0 Pending) and "Service Status Totals" (1 Ok, 1 Warning, 0 Unknown, 1 Critical, 4 Pending).

- Services: Check tasks being monitored (e.g., CPU usage).

Current Network Status

Last Updated: Tue Oct 14 05:26:58 UTC 2025
Updated every 90 seconds
Nagios® Core™ 4.5.7 - www.nagios.org
Logged in as nagiosadmin

Host Status Totals

Up	Down	Unreachable	Pending
1	0	0	0

All Problems All Types

Service Status Totals

Ok	Warning	Unknown	Critical	Pending
2	1	0	1	3

All Problems All Types

Service Status Details For All Hosts

Host	Service	Status	Last Check	Duration	Attempt	Status Information
localhost	Current Load	CRITICAL	10-14-2025 05:26:20	0d 0h 2m 39s	3/4	CRITICAL - load average: 6.41, 743.634
	Current Users	OK	10-14-2025 05:26:02	0d 0h 3m 21s+	1/4	USERS OK - 0 users currently logged in
	HTTP	WARNING	10-14-2025 05:26:45	0d 0h 1m 13s	2/4	HTTP WARNING: HTTP/11 401 Unauthorized - 695 bytes in 0.013 second response time
	PING	OK	10-14-2025 05:26:28	0d 0h 3m 21s+	1/4	PING OK - Packet loss = 0%, RTA = 0.07 ms
	Root Partition	PENDING	N/A	0d 0h 3m 21s+	1/4	Service check scheduled for Tue Oct 14 05:27:11 UTC 2025
	Swap Usage	PENDING	N/A	0d 0h 3m 21s+	1/4	Service check scheduled for Tue Oct 14 05:27:54 UTC 2025
	Total Processes	PENDING	N/A	0d 0h 3m 21s+	1/4	Service check scheduled for Tue Oct 14 05:28:37 UTC 2025

Results 1 - 7 of 7 Matching Services

- Alerts: Access recent notifications.

Alert History

Last Updated: Tue Oct 14 05:27:24 UTC 2025
Nagios® Core™ 4.5.7 - www.nagios.org
Logged in as nagiosadmin

All Hosts and Services

Latest Archive ← Tue Oct 14 00:00:00 UTC 2025 to Present..

File: /opt/nagios/var/nagios.log

Log File Navigation

October 14, 2025 05:00

[10-14-2025 05:27:20] SERVICE ALERT: localhost;Current Load;CRITICAL;HARD;4;CRITICAL - load average: 5.41, 6.94, 6.24
[10-14-2025 05:26:43] SERVICE ALERT: localhost;HTTP;WARNING;SOFT;2;HTTP WARNING: HTTP/11 401 Unauthorized - 695 bytes in 0.013 second response time
[10-14-2025 05:26:20] SERVICE ALERT: localhost;Current Load;CRITICAL;SOFT;3;CRITICAL - load average: 6.41, 743, 6.34
[10-14-2025 05:25:45] SERVICE ALERT: localhost;HTTP;WARNING;SOFT;1;HTTP WARNING: HTTP/11 401 Unauthorized - 695 bytes in 0.230 second response time
[10-14-2025 05:25:19] SERVICE ALERT: localhost;Current Load;CRITICAL;SOFT;2;CRITICAL - load average: 7.83, 777, 6.37
[10-14-2025 05:24:19] SERVICE ALERT: localhost;Current Load;CRITICAL;SOFT;1;CRITICAL - load average: 10.36, 812, 6.38
[10-14-2025 05:23:37] Nagios 4.5.7 starting.. (PID=26)

Step 4: Monitoring Host Details

1. **Navigate to the Host Information Page:**
 - Select a host from the **Hosts** menu.
2. **Key Details:**
 - Host Status: Indicates if the system is UP or DOWN.

- Metrics: View CPU usage, memory status, and network activity.
- Actions: Reschedule checks, disable notifications, or schedule downtime.

The screenshot shows the Nagios web interface at localhost:8881. The left sidebar has links for General, Home, Documentation, Current Status, Tactical Overview, Map, Hosts, Services, Host Groups, Problems, Reports, Availability, Trends, Alerts, and History. The main content area shows 'Host Information' for 'localhost'. It includes details like last update (Tue Oct 14 05:27:55 UTC 2025), status (UP), and performance data (rtt=0.087000ms;3000.000000;5000.000000;0.000000 pl=0%;80;100;0). Below this is the 'Host State Information' section, which lists various states and their details. To the right is the 'Host Commands' section, which contains a list of actions with icons and descriptions. At the bottom is a 'Host Comments' section.

Step 5: Stop and Remove Nagios

1. Stop the Container:

- Command:

```
docker stop nagiosdemo
```

```
C:\Users\kmit>
C:\Users\kmit>docker stop nagiosdemo
nagiosdemo

C:\Users\kmit>
```

2. Delete the Container:

- Command:

```
docker rm nagiosdemo
```

```
C:\Users\kmit>docker rm nagiosdemo  
nagiosdemo  
C:\Users\kmit>
```

3. Remove the Image (Optional):

- o List images:

```
docker images
```

- o Delete the Nagios image:

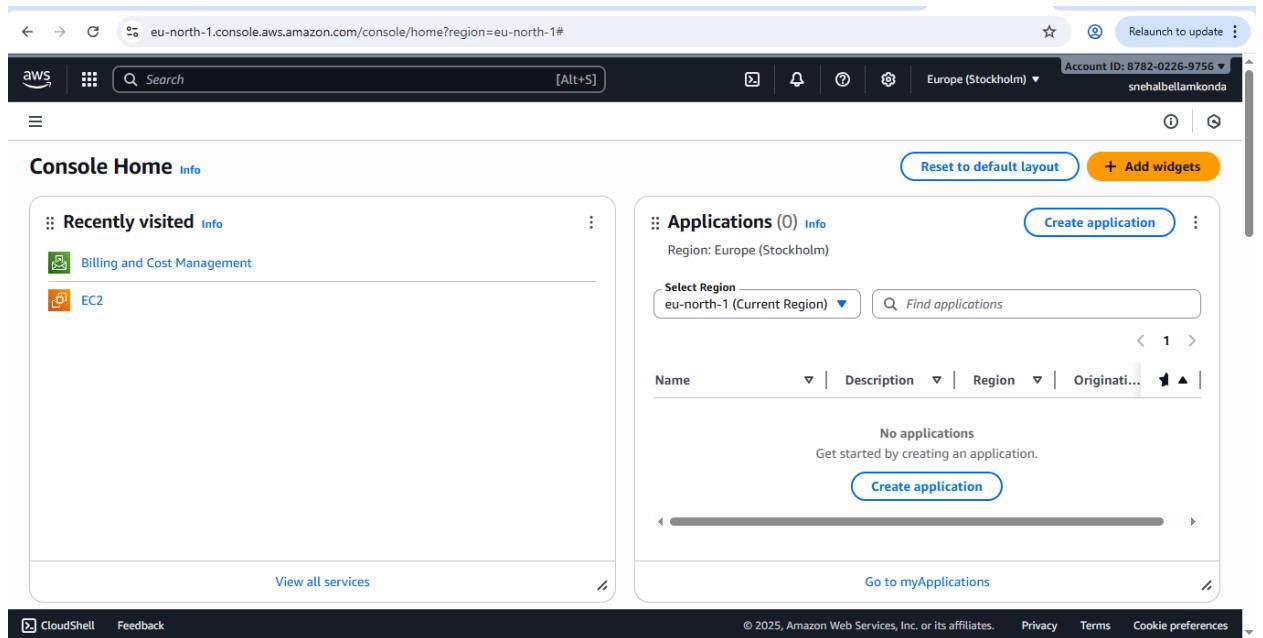
```
docker rmijasonrivers/nagios:latest
```

4. Observe the docker containers in DockerHub, we can see the latest Nagios Installed running on port:8888

Steps for AWS-free trier account creation

- open <https://aws.amazon.com>
- click on create AWS account
- Provide email id and name details for AWS account creation then a screen will appear as below
- Enter verification code received in your mail
- Provide your contact details and agree Terms and conditions, then click on create
- Provide billing details, click on verify and continue
- Amazon charges 2 rs and it will be credit back into your account once verification is over
- Once payment success a screen appear as below, provide your working contact number after selecting country, and click on send SMS after entering the details
- Provide verification code received in your mobile
- Now select basic plan and click on complete sign up
- Now a screen appear as follows then click on AWS management console

- You will navigating to the aws page ..where you need to Select role as Academic/Researcher and interest as DevOps and click on submit
- Sign into AWS console as Root user
- A console screen appears aws will appear



SBQ's

1. Your Pod keeps restarting repeatedly. What will you do?

ANS: Check logs: `kubectl logs <pod-name>`

Describe pod: `kubectl describe pod <pod-name>`

Look for crash loop or configuration issues.

2. A Kubernetes pod is stuck in a "Pending" state. What could be the possible reasons, and how would you troubleshoot it?

ANS : • Causes: No resources, nodeSelector mismatch, or PVC issues.

• Troubleshoot: `kubectl describe pod <pod-name>`

3. How would you debug a failed deployment in Kubernetes?

ANS : Check rollout status: `kubectl rollout status deployment <deployment-name>`

Check events: kubectl describe deployment <deployment-name>

4. You have a Kubernetes Deployment with multiple replicas, and some pods are failing health checks. How would you identify the root cause and fix it?

ANS : Run: kubectl describe pod <pod-name>

Inspect readiness/liveness probe errors.

Check container logs.

5. How do you roll back a faulty deployment?

ANS : kubectl rollout undo deployment <deployment-name>

6. How do you debug a running Pod?

ANS : kubectl exec -it <pod-name> -- /bin/bash

Or check logs: kubectl logs <pod-name>

7. You need to expose a local service externally.How to do?

ANS : kubectl expose deployment myapp --type=NodePort --port=80

minikube service myapp --url

8. How to Start and stop Nagios

ANS : docker start nagiosdemo

docker stop nagiosdemo

9. You installed Nagios but the web interface shows “*Unable to connect to Nagios process*”.how to resolve this?

ANS : Restart container.

Check configuration errors inside container logs.

10. You added a new host in Nagios, but it's not appearing on the web interface.how to check?

ANS : Check nagios.cfg and included configs.

Ensure config file syntax is correct.

Restart Nagios container.

11. How can you check whether Nagios is running properly?

ANS : docker ps

12. How do you view Nagios logs in real-time

ANS : docker logs -f nagiosdemo

13. What are the advantages of using Nagios?

ANS : Real-time monitoring

Custom alerting system

Supports plugins for extensibility

Easy visualization through Web UI

Conclusion: In this week we learnt how to deploy pods in minikube and how to monitor our local system using Nagios tool and will know how to create the aws free tier account.

LAB REPORT – WEEK 11: Jenkins CI/CD Integration using GitHub Webhooks and Email Notifications

```
ngrok
* Decouple policy and sensitive data with vaults: https://ngrok.com/r/secrets

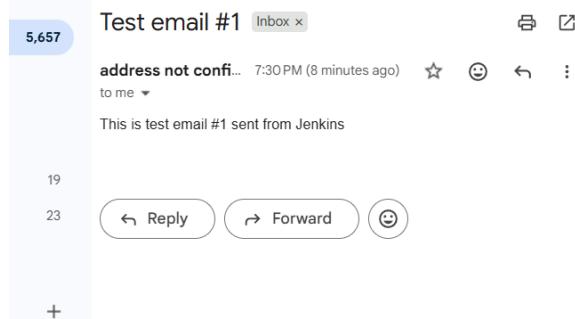
Session Status      online
Account            mad674 (Plan: Free)
Update             update available (version 3.32.0, Ctrl-U to update)
Version            3.24.0-msix
Region             India (in)
Web Interface     http://127.0.0.1:4040
Forwarding        https://9f27a2046653.ngrok-free.app -> http://localhost:80

Connections        ttl     opn     rti     rt5     p50     p90
                   0       0       0.00   0.00   0.00   0.00
```

A screenshot of a GitHub repository page. At the top, there are navigation links: issues, Pull requests, Actions, Projects, Wiki, Security, Insights, and Settings. The Settings tab is currently selected. Below the navigation, a message says "s successfully created. We sent a ping payload to test it out! Read more about it at https://docs.github.com/webhooks/#ping-event." Under the "Webhooks" section, there is a single webhook entry: "https://9f27a2046653.ngrok-free.ap... (push)". A note below it says "This hook has never been triggered." There are "Edit" and "Delete" buttons next to the URL.

A screenshot of the Jenkins global configuration page. On the left, there is a sidebar with sections: General, Access, Collaborators, Moderation options, Code and automation (Rules, Actions, Models), Webhooks (selected), Copilot, Environments, Codespaces, and Pages. Below the sidebar, there is a "Git" section with a "Triggers" button. Under "Triggers", there are several checkboxes: "Build after other projects are built", "Build periodically", "GitHub hook trigger for GITScm polling" (which is checked), and "Poll SCM". In the main content area, there is a "Webhooks" section with a note: "Webhooks allow external services to be notified when certain events happen. When the specified events happen, we'll send a POST request to each of the URLs you provide. Learn more in our [Webhooks Guide](#)". Below this, there is a list of webhooks: "https://59c8389d8820.ngrok-free.ap... (push)" (with a checkmark) and "Last delivery was successful". There are "Edit" and "Delete" buttons next to the URL.

```
Started by GitHub push by mad674
Running as SYSTEM
Building in workspace C:\ProgramData\Jenkins\.jenkins\workspace\ webhook
The recommended git tool is: NONE
> credentials specified
> git.exe rev-parse --resolve-git-dir
:C:\ProgramData\Jenkins\.jenkins\workspace\ webhook\.git # timeout=10
Getting changes from the remote Git repository
> git.exe config remote.origin.url https://github.com/mad674/ngrok.git # timeout=10
Getting upstream changes from https://github.com/mad674/ngrok.git
> git.exe --version # timeout=10
> git --version # 'git version 2.45.1.windows.1'
> git.exe fetch --tags --force --progress -- https://github.com/mad674/ngrok.git
refs/heads/*:refs/remotes/origin/* # timeout=10
> git.exe rev-parse "refs/remotes/origin/main^{commit}" # timeout=10
Checking out Revision 63c2e22b0d237ad5b69eb259aba959016e148bcb
refs/remotes/origin/main)
> git.exe config core.sparsecheckout # timeout=10
> git.exe checkout -f 63c2e22b0d237ad5b69eb259aba959016e148bcb # timeout=10
Commit message: "fix"
> git.exe rev-list --no-walk 2145a7540661aa3a018e1eaf050240059072eb31 # timeout=10
Finished: SUCCESS
```



Viva Questions & Answers

1. Q: What is Continuous Integration (CI)?

A: CI is the practice of automatically integrating code changes from multiple developers into a shared repository several times a day.

2. Q: What is Continuous Deployment (CD)?

A: CD automates the release of validated code to production after successful testing and integration.

3. Q: What is the role of Jenkins in a CI/CD pipeline?

A: Jenkins automates building, testing, and deploying code, ensuring a continuous integration and delivery workflow.

4. Q: What is a webhook in GitHub?

A: A webhook is an HTTP callback that notifies Jenkins when an event (like a push) occurs in a GitHub repository.

5. Q: Why are webhooks used in Jenkins integration?

A: They remove the need for periodic polling and allow immediate triggering of builds upon a GitHub event.

6. Q: What are the different types of build triggers available in Jenkins?

A: Build triggers include: Poll SCM, Build periodically, Trigger builds remotely, and GitHub webhook triggers.

7. Q: What is the difference between polling and webhook triggers?

A: Polling checks for changes at fixed intervals, while webhooks notify Jenkins instantly when changes occur.

8. Q: What is ngrok and why is it used in Jenkins-GitHub integration?

A: ngrok is a tunneling tool that exposes local servers to the internet, allowing GitHub to send webhooks to a local Jenkins instance.

9. Q: How does ngrok help in setting up webhooks for Jenkins running locally?

A: ngrok creates a public HTTPS URL that forwards requests to the local Jenkins, enabling GitHub to reach it.

10. Q: Why do we configure email notifications in Jenkins?

A: To receive automated build success or failure updates, helping developers monitor build status efficiently.

The screenshot shows the AWS Academy Learner Lab interface. On the left, there's a sidebar with icons for Account, Dashboard, Courses (which is selected), Calendar, Inbox, History, and Help. The main area has a navigation bar with Home, Modules (selected), Discussions, Grades, and Lucid (Whiteboard). A central panel displays a blue V-shaped icon with a red curved arrow at the top. At the bottom right of the main panel, there's a "Learner Lab" section with a list of links: Environment Overview, Environment Navigation, Access the AWS Management Console, Region restriction, Service usage and other restrictions, Using the terminal in the browser, Running AWS CLI commands, Using the AWS SDK for Python, Preserving your budget, Accessing EC2 Instances, SSH Access to EC2 Instances, SSH Access from Windows, and SSH Access from a Mac. Below this section are "Previous" and "Next" buttons. The status bar at the bottom shows weather information (23°C Sunny) and system details (ENG IN 10:05 11-11-2025).

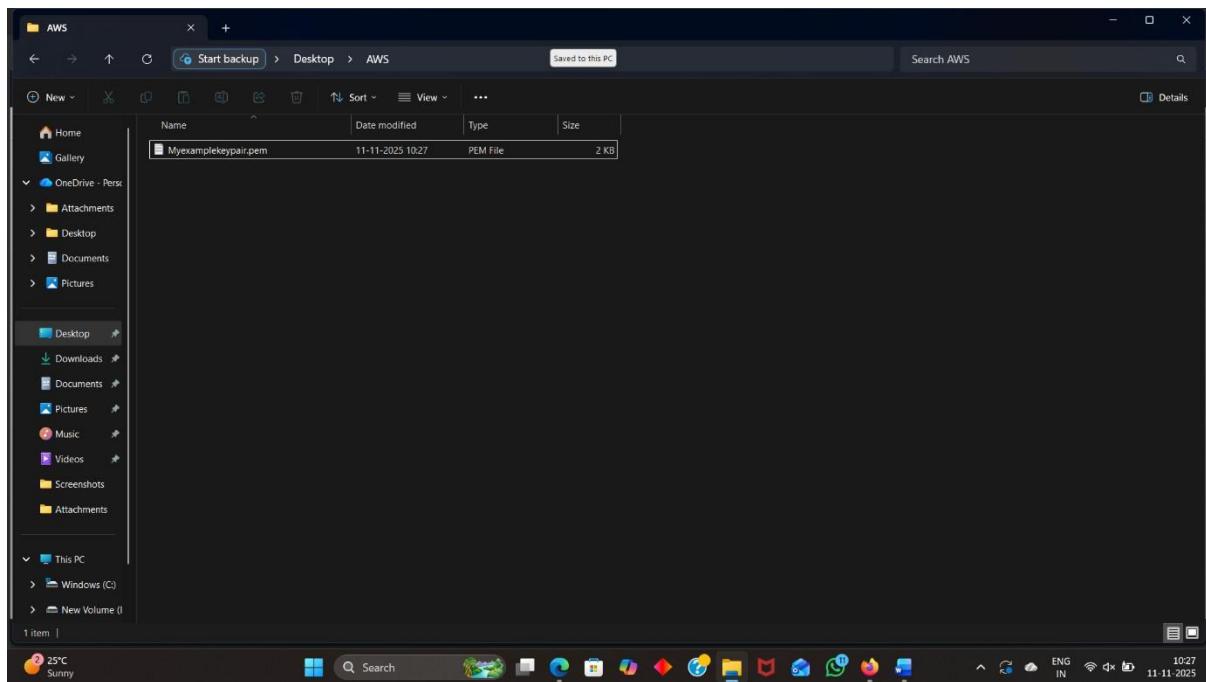
This screenshot shows the AWS Academy Learner Lab interface after starting a lab session. The main panel now features a terminal window with the command "eee_N_535353@runweb195508:~\$". The rest of the interface remains the same as the first screenshot, including the sidebar, navigation bar, and the "Learner Lab" section with its list of links. The status bar at the bottom is identical to the first screenshot.

The screenshot shows the AWS EC2 Home page. On the left, a sidebar navigation includes EC2 Global View, Events, Instances (with sub-options like Instances, Instance Types, Launch Templates, Spot Requests, Savings Plans, Reserved Instances, Dedicated Hosts, Capacity Reservations, Capacity Manager), Images (AMIs, AMI Catalog), and Elastic Block Store (Volumes, Snapshots). The main content area displays a summary of resources: 0 Instances (running), 0 Auto Scaling Groups, 0 Capacity Reservations, 0 Dedicated Hosts, 0 Elastic IPs, 0 Instances, 1 Key pairs, 0 Load balancers, 0 Placement groups, 1 Security groups, 1 Snapshots, and 0 Volumes. Below this is a 'Launch instance' section with a 'Launch instance' button and a 'Migrate a server' link. To the right is a 'Service health' section showing the United States (N. Virginia) region is operating normally. Further right is an 'Account attributes' section listing the Default VPC (vpc-0b0751dc4785a7253), Settings (Data protection and security, Allowed AMIs, Zones, EC2 Serial Console, Default credit specification, EC2 console preferences), and an 'Explore AWS' section with tips for cost optimization and a '10 Things You Can Do Today to Reduce AWS Costs' list.

The screenshot shows the 'Launch an instance' wizard. Step 1: Name and tags. It asks for a name (input: myexamplewebserver) and allows adding additional tags. Step 2: Application and OS Images (Amazon Machine Image). It lists various AMI categories: Amazon Linux, macOS, Ubuntu, Windows, Red Hat, SUSE Linux, and Debian. A search bar is available to find more images. Step 3: Summary. It shows 1 instance selected, the software image (Canonical, Ubuntu, 24.04, amd64), virtual server type (t3.micro), firewall (New security group), storage (1 volume(s) - 8 GiB), and provides 'Launch instance' and 'Preview code' buttons.

The screenshot shows the AWS EC2 Instances Launch page. At the top, it displays the instance details: Architecture (64-bit (x86)), AMI ID (ami-0ecb62995f68bb549), Publish Date (2025-10-22), and Username (ubuntu). A green "Verified provider" badge is present. On the left, the "Instance type" section shows a dropdown set to t2.micro, with detailed pricing information for various On-Demand and On-Demand SUSE base options. Below this, a note states "Additional costs apply for AMIs with pre-installed software". The "Key pair (login)" section requires a key pair name, with a dropdown menu showing "Select" and a "Create new key pair" button. On the right, the "Summary" panel shows 1 instance, the software image (Canonical, Ubuntu, 24.04, amd64), the virtual server type (t2.micro), the firewall (New security group), and storage (1 volume(s) - 8 GiB). Buttons for "Cancel", "Launch instance", and "Preview code" are at the bottom.

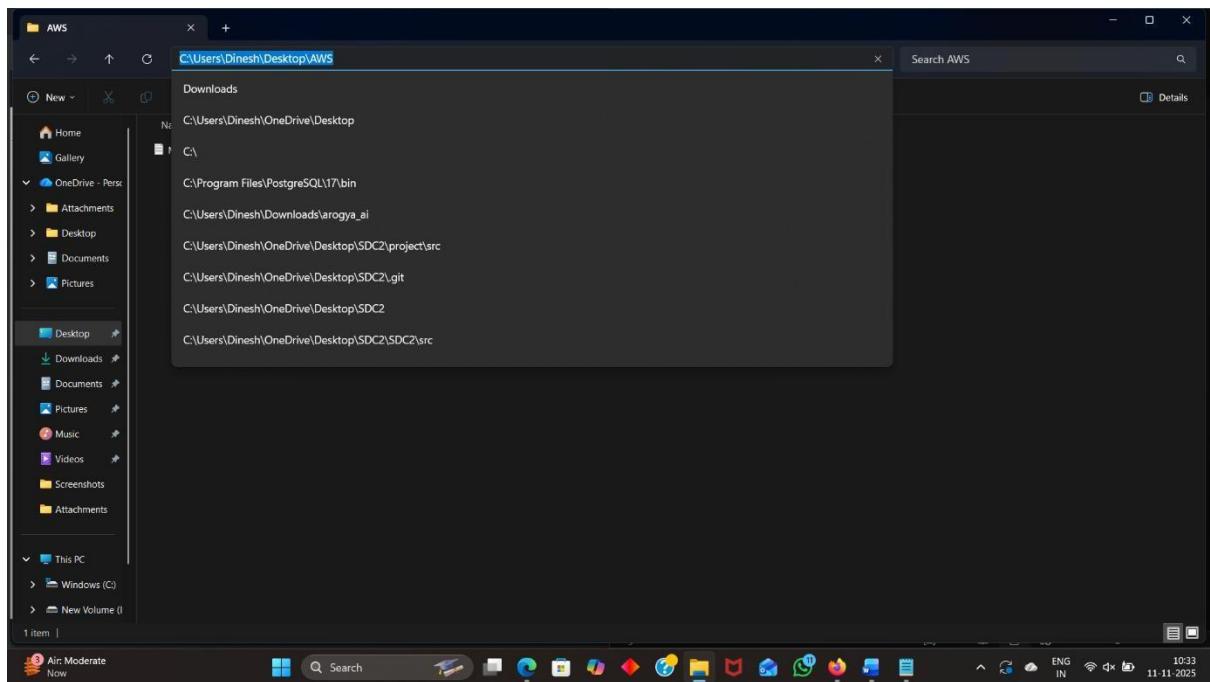
This screenshot shows the "Create key pair" dialog box overlaid on the main EC2 launch page. The dialog has a title "Create key pair" and a "Key pair name" field containing "Enter key pair name". It also includes a note about key pairs allowing secure connection. The "Key pair type" section shows two options: "RSA" (selected) and "ED25519" (disabled). The "Private key file format" section shows two options: ".pem" (selected) and ".ppk" (disabled). A warning message in a yellow box states: "⚠️ When prompted, store the private key in a secure and accessible location on your computer. You will need it later to connect to your instance. [Learn more](#)". The background of the main page remains visible, showing the same instance configuration and summary as the first screenshot.

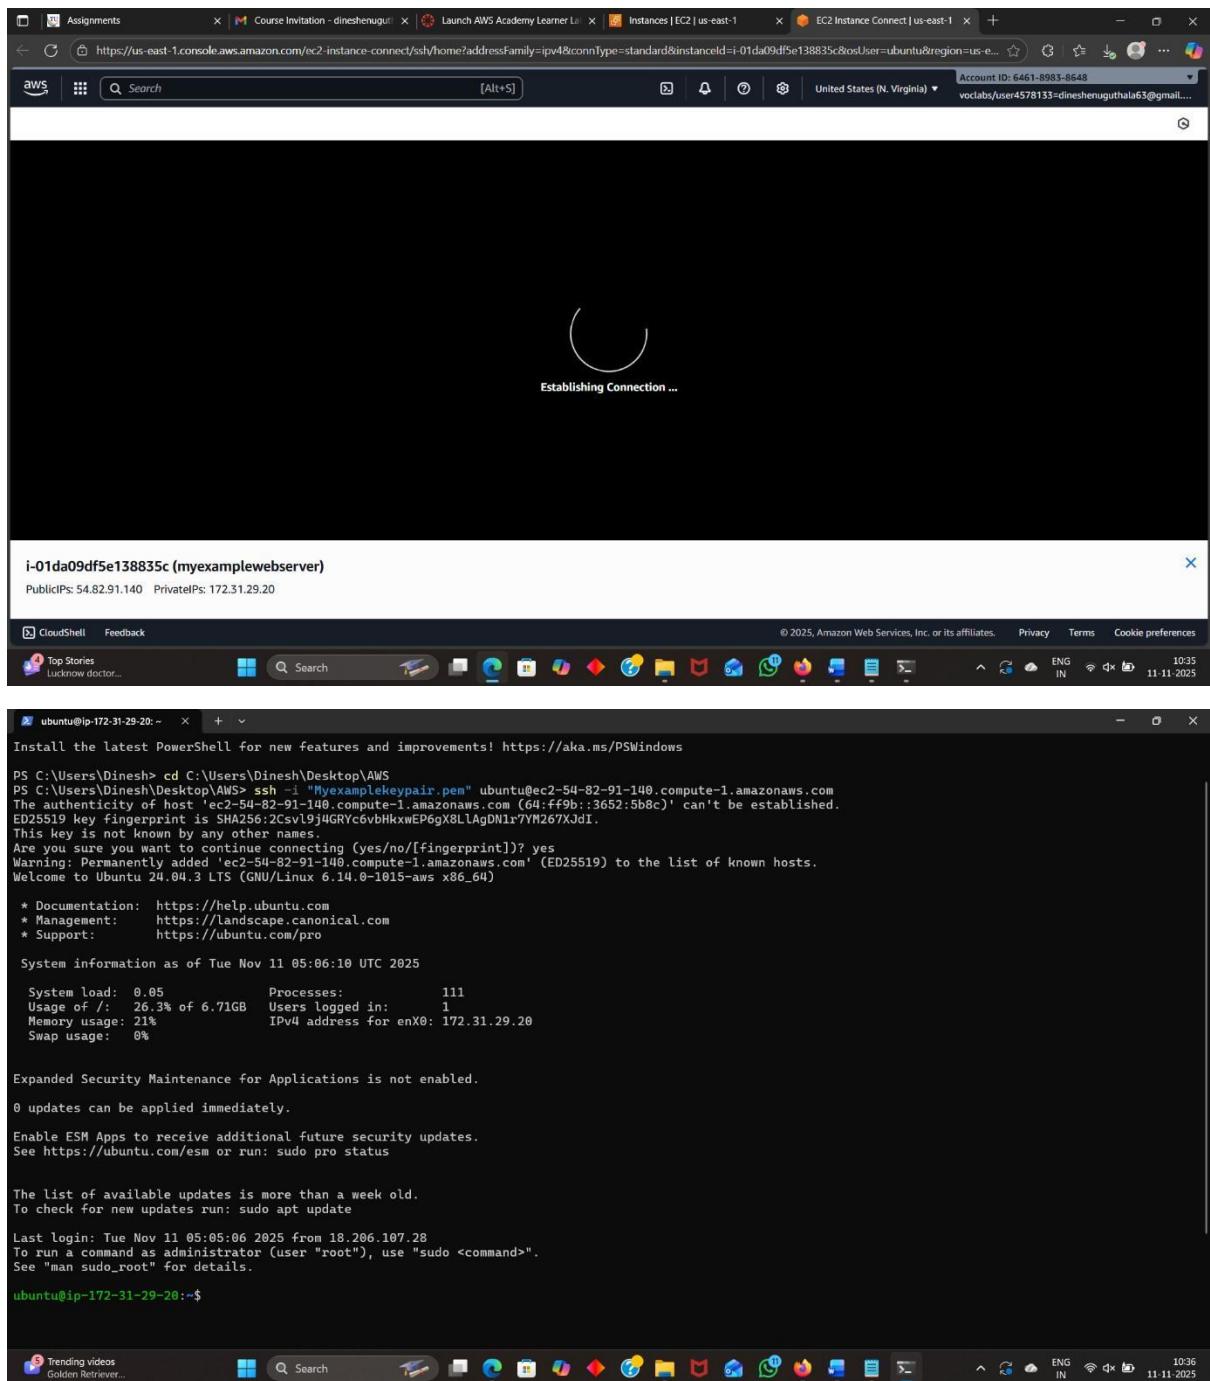


A screenshot of a web browser displaying the AWS EC2 Instances Launch page. The URL is https://us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#LaunchInstances. The page shows a green success message: "Success: Successfully initiated launch of instance (i-01da09df5e138835c)". Below this, there's a "Launch log" section. A "Next Steps" section contains several cards: "Create billing usage alerts", "Connect to your instance", "Connect an RDS database", "Create EBS snapshot policy", "Manage detailed monitoring", "Create Load Balancer", "Create AWS budget", and "Manage CloudWatch alarms". The status bar at the bottom shows "CloudShell Feedback" and the date and time as "11-11-2025 10:28".

The screenshot shows the AWS EC2 Instances page. On the left, a navigation sidebar is open under the 'EC2' section, showing options like Dashboard, EC2 Global View, Events, Instances (selected), Instance Types, Launch Templates, Spot Requests, Savings Plans, Reserved Instances, Dedicated Hosts, Capacity Reservations, and Capacity Manager. The main content area displays a table titled 'Instances (1) Info'. The table has columns for Name, Instance ID, Instance state, Instance type, Status check, Alarm status, Availability Zone, and Public IPv4. One row is visible, showing 'myexamplewe...' as the name, 'i-01da09df5e138835c' as the Instance ID, 'Running' as the Instance state, 't2.micro' as the Instance type, 'Initializing' as the Status check, 'us-east-1d' as the Availability Zone, and 'ec2-54-82...' as the Public IPv4. Below the table, a section titled 'Select an instance' is visible.

The screenshot shows the 'Connect to instance' page for the instance 'i-01da09df5e138835c'. The top navigation bar includes links for Assignments, Course Invitation, Launch AWS Academy Learner Lab, and Connect to Instance. The main content area is titled 'Connect info' and contains instructions to connect using a browser-based client. It shows the 'EC2 Instance Connect' tab is active, along with Session Manager, SSH client, and EC2 serial console tabs. Under 'Instance ID', it shows 'i-01da09df5e138835c (myexamplewebserver)'. The 'Connection type' section has two options: 'Connect using a Public IP' (selected) and 'Connect using a Private IP'. Under 'Public IPv4 address', it lists '54.82.91.140'. The 'Username' field is set to 'ubuntu'. A note at the bottom states: 'Note: In most cases, the default username, ubuntu, is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.' At the bottom right are 'Cancel' and 'Connect' buttons. The bottom navigation bar includes CloudShell, Feedback, and standard system icons.





```
ubuntu@ip-172-31-29-20:~ % + ~
Reading state information... Done
10 packages can be upgraded. Run 'apt list --upgradable' to see them.
ubuntu@ip-172-31-29-20:~ $ sudo apt-get install docker.io
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
bridge-utils containerd dns-root-data dnsmasq-base pigz runc ubuntu-fan
Suggested packages:
ifupdown aufs-tools cgroupfs-mount | cgroup-lite debootstrap
docker-buildx docker-compose-v2 docker-doc rinse zfs-fuse | zfsutils
The following NEW packages will be installed:
bridge-utils containerd dns-root-data dnsmasq-base docker.io pigz runc
ubuntu-fan
0 upgraded, 8 newly installed, 0 to remove and 10 not upgraded.
Need to get 76.0 MB of archives.
After this operation, 288 MB of additional disk space will be used.
Do you want to continue? [Y/n] y
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 pigz amd64 2.8-1 [65.6 kB]
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 bridge-utils amd64 1.7.1-1ubuntu2 [33.9 kB]
Get:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main amd64 runc amd64 1.3.3-0ubuntu1~24.04.2 [8815 kB]
Get:4 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main amd64 containerd amd64 1.7.28-0ubuntu1~24.04.1 [38.4 MB]
Get:5 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main amd64 dns-root-data all 2824071801~ubuntu0.24.04.1 [5918 B]
Get:6 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main amd64 dnsmasq-base amd64 2.90-2ubuntu0.1 [376 kB]
Get:7 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates/universe amd64 docker.io amd64 28.2.2-0ubuntu1~24.04.1 [28.3 MB]
Get:8 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates/universe amd64 ubuntu-fan all 0.12.16+24.04.1 [34.2 kB]
Fetched 76.0 MB in 7(71.0 MB/s)
Preconfiguring packages ...
Selecting previously unselected package pigz.
(Reading database ... 71735 files and directories currently installed.)
Preparing to unpack .../0-pigz_2.8-1_amd64.deb ...
Unpacking pigz (2.8-1) ...
Selecting previously unselected package bridge-utils.
Preparing to unpack .../1-bridge-utils_1.7.1-1ubuntu2_amd64.deb ...
Unpacking bridge-utils (1.7.1-1ubuntu2) ...
Selecting previously unselected package runc.
Preparing to unpack .../2-runc_1.3.3-0ubuntu1~24.04.2_amd64.deb ...
Unpacking runc (1.3.3-0ubuntu1~24.04.2) ...
Selecting previously unselected package containerd.
Selecting previously unselected package containerd.
Preparing to unpack .../3-containerd_1.7.28-0ubuntu1~24.04.1_amd64.deb ...

```

S 25°C Sunny Search ENG IN 10:37 11.11.2025

```
ubuntu@ip-172-31-29-20:~ % + ~
Created symlink /etc/systemd/system/multi-user.target.wants/containerd.service → /usr/lib/systemd/system/containerd.service.
Setting up ubuntu-fan (0.12.16+24.04.1) ...
Created symlink /etc/systemd/system/multi-user.target.wants/ubuntu-fan.service → /usr/lib/systemd/system/ubuntu-fan.service.
Setting up docker.io (28.2.2-0ubuntu1~24.04.1) ...
info: Selecting GID from range 100 to 999 ...
info: Adding group 'docker' (GID 113) ...
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
Created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /usr/lib/systemd/system/docker.socket.
Processing triggers for dbus (1.14.10-4ubuntu4.1) ...
Processing triggers for man-db (2.12.0-4build2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

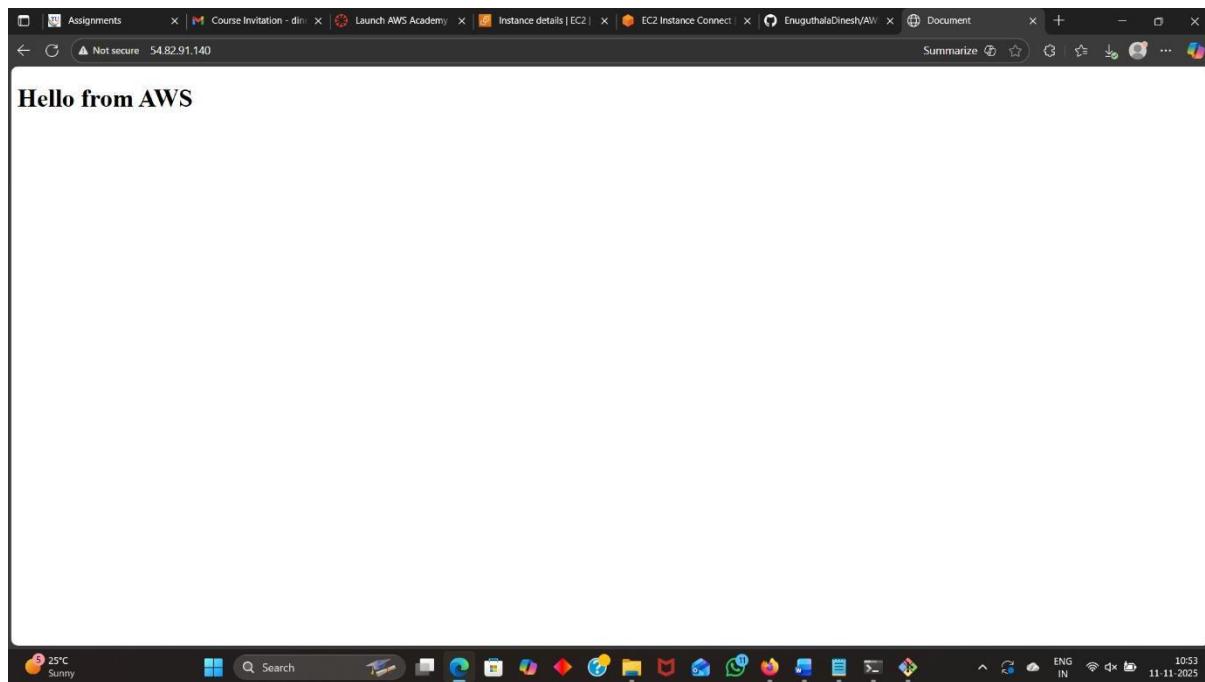
No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
ubuntu@ip-172-31-29-20:~ $ sudo apt install git
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
git is already the newest version (1:2.43.0-1ubuntu7.3).
git set to manually installed.
0 upgraded, 0 newly installed, 0 to remove and 10 not upgraded.
ubuntu@ip-172-31-29-20:~ $ sudo apt install nano
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
nano is already the newest version (7.2-2ubuntu0.1).
nano set to manually installed.
0 upgraded, 0 newly installed, 0 to remove and 10 not upgraded.
ubuntu@ip-172-31-29-20:~ |
```

S 25°C Sunny Search ENG IN 10:38 11.11.2025



A screenshot of the AWS Management Console EC2 Instances page. The URL in the address bar is 'https://us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#InstanceDetails:instanceId=i-01da09df5e138835c'. The left sidebar shows navigation options like Dashboard, EC2 Global View, Events, Instances (selected), Instance Types, Launch Templates, Spot Requests, Savings Plans, Reserved Instances, Dedicated Hosts, Capacity Reservations, Images, AMIs, and Elastic Block Store. The main content area is titled 'Instance summary for i-01da09df5e138835c (myexamplewebserver)'. It displays detailed information about the instance, including its ID (i-01da09df5e138835c), Public IP (54.82.91.140), Instance State (Running), and VPC ID (vpc-0b0751dc4785a7253). A context menu is open over the 'Actions' button, with the 'Terminate (delete) instance' option highlighted. The bottom of the screen shows the AWS footer and the Windows taskbar.

Screenshot of the AWS EC2 Instance Details page showing the termination process.

Instance summary for i-01da09df5e138835c (myexamplewebserver)

Terminate (delete) instance

On an EBS-backed instance, the default action is for the root EBS volume to be deleted when the instance is terminated. Storage on any local drives will be lost.

Are you sure you want to terminate these instances?

Instance ID: i-01da09df5e138835c (myexamplewebserver) | Termination protection: Disabled

To confirm that you want to delete the instances, choose the terminate button below. Instances with termination protection enabled will not be terminated. Terminating the instance cannot be undone.

Skip OS shutdown: This option skips the graceful OS shutdown process. Use only when your instance must be stopped immediately, such as during an emergency or failure.

Skip OS shutdown

Cancel **Terminate (delete)**

CloudShell Feedback

Screenshot of the AWS EC2 Instance Details page showing the successful termination of the instance.

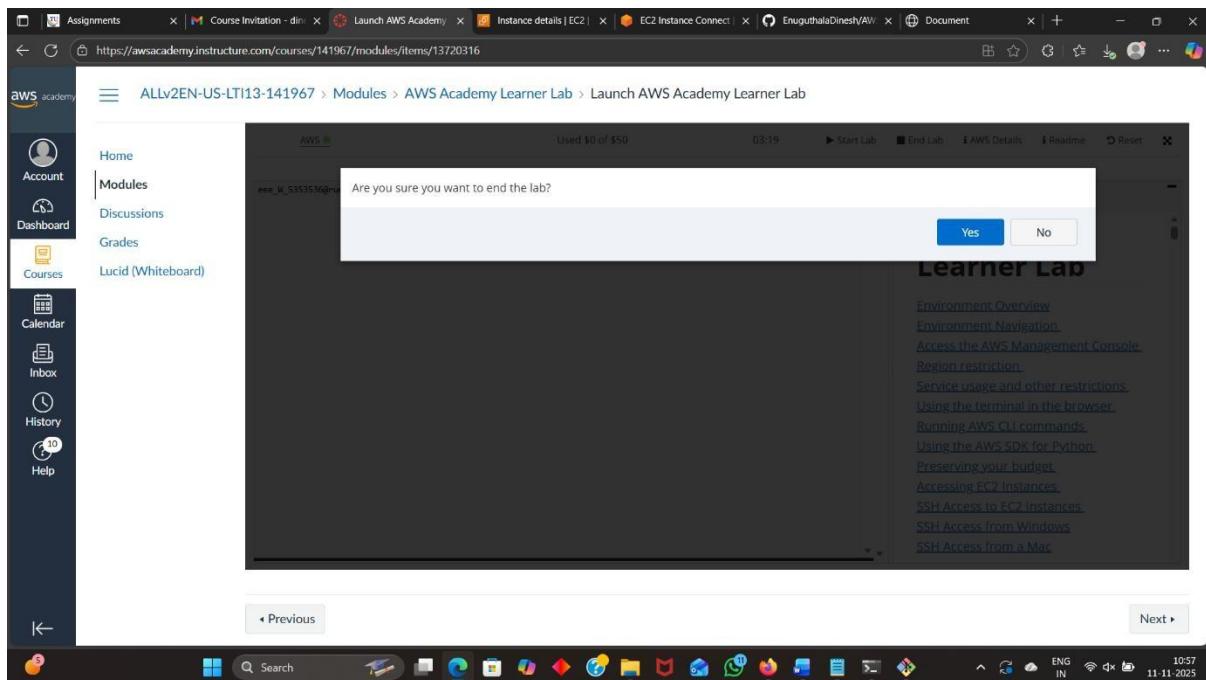
Successfully Initiated termination (deletion) of i-01da09df5e138835c

Instance summary for i-01da09df5e138835c (myexamplewebserver)

Updated less than a minute ago

Attribute	Value
Public IPv4 address	54.82.91.140 open address
Private IP4 address	172.31.29.20
Instance state	Shutting-down
Private IP DNS name (IPv4 only)	ip-172-31-29-20.ec2.internal
Instance type	t2.micro
VPC ID	vpc-0b0751dc4785a7253
Subnet ID	subnet-087b86d65ddff9481e
Public DNS	ec2-54-82-91-140.compute-1.amazonaws.com open address
Elastic IP addresses	-
AWS Compute Optimizer finding	Opt-in to AWS Compute Optimizer for recommendation
Auto Scaling Group name	-

CloudShell Feedback



Scenario-Based Questions and Answers

1. You have a simple index.html file on your laptop and you launched an EC2 instance with Amazon Linux 2. The instance is running but when you open the public IP in browser, the page doesn't load.

What steps will you take to host the index.html?

Answer: Install Apache → Start service → Copy index.html to /var/www/html/ → Allow port 80 in Security Group.

2. You deployed your index.html to /var/www/html/ directory on EC2, but the web page still isn't loading. What are two possible issues you would check?

Answer: Check if Apache is running and if port 80 is open in Security Group.

3. You installed Apache HTTP server on EC2 to host index.html, but the service stops after instance reboot. What command should you run to ensure it auto-starts on boot?

Answer: sudo systemctl enable httpd

4. You are deploying a Maven web application onto an EC2 instance. Maven is not installed on the instance. What commands or steps will you follow to install Maven on Amazon Linux/Ubuntu EC2?

Answer:

- Amazon Linux: sudo yum install maven -y
- Ubuntu: sudo apt update && sudo apt install maven -y

5. You built a Maven project on EC2, and a .war file generated inside target/. You want to deploy it using Tomcat. Where will you place the .war file and why?

Answer: Place it in /usr/share/tomcat/webapps/ — Tomcat auto-deploys .war files from there.

6. Your Maven web app is deployed to Tomcat on EC2, but accessing it via browser gives 404 error. What configuration or path issues will you check?

Answer: Check if .war deployed correctly, Tomcat logs, context path, and service status.

7. You can access your web application locally on the EC2 instance using curl localhost:8080 but not from your browser. What AWS setting is likely missing or misconfigured?

Answer: Security Group missing inbound rule for port 8080.

8. You have deployed your web app on EC2 successfully, but after shutting down your local Wi-Fi and reconnecting, the public IP changed and the app is not opening. What AWS feature helps avoid this issue?

Answer: Use an Elastic IP.

9. You want to automate deployment of your index.html whenever you restart the EC2 instance. Which EC2 feature can be used to run commands automatically during instance setup?

Answer: Use User Data script.

10. Your Maven application needs external dependencies during build but EC2 has no internet. What AWS service or change can allow the EC2 instance to download dependencies securely?

Answer: Use a NAT Gateway or VPC Endpoint for internet access.